

THE CODE AFTER SERIES

Code After

A Translation Project for the AI Era

Richard Yan

Series Announcement Paper

Open-Access Release

May 8, 2026

codeafter.ai

Author: Richard Yan

Email: richard.yan@codeafter.ai

ORCID: <https://orcid.org/0009-0000-7611-6323>

Affiliation: Code After AI

Document type: Series Announcement Paper — Open-Access Release

Release date: 8 May 2026

Repository: Zenodo

DOI: 10.5281/zenodo.20079145

Additional dissemination: ResearchGate; SSRN; codeafter.ai

Length: Approx. 11,000 words

Language: English (Chinese edition in parallel to follow)

© 2026 Richard Yan · codeafter.ai

Licensed under Creative Commons Attribution — NonCommercial — NoDerivatives 4.0 International (CC BY-NC-ND 4.0). Attribution required. Commercial use and derivative works prohibited. For inquiries about translation, local-language editions, and publishing partnership, see Section VIII of this paper or contact partners@codeafter.ai.

ABSTRACT

This paper announces the Code After Series: six papers over twenty-four months. Each extends the framework of *Code After: Law, Accounting, and the Governance of Artificial Intelligence* (v0.9, April 2026) into a distinct institutional domain being reshaped by AI. The paper opens by naming the structural break the project is built to address — the move from a Pre-Code condition, in which deterministic code was executed by transparent tools, to a Post-Code condition, in which probabilistic intelligence operates as a quasi-agent inside institutions built for the earlier regime. It sets out the project's core thesis — that AI is a general-purpose decoupling force acting on institutions built for stable representation — and specifies the methodological discipline, publication architecture, and distribution model through which the series will be produced and disseminated. It also serves as a formal invitation to local publishers, translators, and institutional partners positioned to produce language editions outside the G2.

SUGGESTED CITATION

Yan, R. (2026). *Code After: A Translation Project for the AI Era*. Zenodo. DOI: 10.5281/zenodo.20079145. Also deposited on SSRN and ResearchGate.

COMPANION WORK

Yan, R. (2026). *Code After: Law, Accounting, and the Governance of Artificial Intelligence* (v0.9, April 2026). Zenodo. DOI: 10.5281/zenodo.19537473. Also deposited on SSRN (ID: 6563959) and ResearchGate (Richard-Yan-10).

Table of Contents

I. Why Code After	4
II. What Came Before.....	6
III. The Core Thesis.....	8
IV. The Widening Gap	10
Capital	10
Compute and Inference	11
The Application Layer and Agents	12
World Models and Physical Ai	13
V. The Series	16
VI. Layered Outputs	18
VII. Publication Architecture and Bilingual Commitment.....	19
VIII. The Partworks Model.....	20
A. Incremental Entry	20
B. The Asymmetry to Acknowledge.....	20
C. Licensing and Terms.....	21
D. Contact.....	21
IX. The Voice	22
X. Why This, Why Now	23
Appendix A. Terms from V0.9 Used in This Paper	24
Appendix B. Terms Developed in This Paper	25

I. Why Code After

Code After names a break.¹

For most of modern history, code meant deterministic instructions executed by tools that did what they were told. Law, accounting, regulation, the architectures of authority that hold a society together — this is the inherited apparatus of modern governance. All of it was built on the assumptions that condition made available. Humans act. Tools execute. Responsibility is traceable. The instrument disappears into the result, and the result can be audited back to the person who caused it.

Stable couplings hold throughout — between action and attribution, between rule and execution, between representation and reality. Call this the Pre-Code condition. The distinction is structural rather than chronological — Pre-Code names not a period before software but a regime in which code, however sophisticated, behaved as deterministic instrumentation subordinate to human intention. Most of the institutions a citizen, a regulator, a court, an auditor, or a policymaker uses to make sense of the world today were designed for it.

The condition no longer holds. AI systems operate as probabilistic, adaptive, and partially opaque agents. They generate knowledge, shape decisions, and participate in economic and institutional processes in ways the inherited languages of evidence, liability, compliance, and sovereignty were not designed to describe. The break is not technical. It is conceptual. A Pre-Code ontology is being used to govern Post-Code actors, and the mismatch produces structural failure at every point where the two are required to meet. The term actor here is functional rather than metaphysical: a system whose outputs participate materially in institutional processes, regardless of whether the system itself possesses agency in the human sense.

Code After: Law, Accounting, and the Governance of Artificial Intelligence (v0.9, April 2026) named the break and developed it inside one area of institutional life — governance — through four structural gaps. The Code After Series extends the operation. If probabilistic intelligence is becoming a quasi-agent — capable of shaping outcomes without being a legal subject in the inherited sense — inside education, work, evidence, measurement, jurisdiction, and language itself, then each one requires the same procedure: a diagnosis of what the inherited instruments no longer reach, and a rewritten operating logic for what now has to be governed.

The project supplies the vocabulary and analytical structure for that transition. Not from one technology to another. From a civilisation in which human intelligence was the only intelligence participating in institutional decisions to one in which it is not — and from one organised around deterministic tools to one co-built with probabilistic intelligences. This is the why of the project. What follows is the what — the people the work is built for, the time available to do it, and the publication architecture through which it will be done.

AI is being built by a very small number of people, working in a very small number of labs, concentrated in a very small number of places, for a world that includes almost everyone else. This is not a complaint. It is a description of the present moment. The capital, talent, compute, data, and institutional concentration required to build frontier AI systems exist at scale in the United States and China, and at partial scale in the

¹ Code is retained as the term throughout this project because the institutional infrastructures of the AI era remain computationally mediated even after deterministic execution ceases to be their defining characteristic. The project name marks a change in what code is, not a renunciation of the term.

United Kingdom and a short list of adjacent economies. The consequences of what is being built will be carried by the roughly eight billion people who live everywhere else.

The gap between who is building AI and who will live with what it builds is the defining inequality of the period. It is wider than any comparable gap in the industrial transition, which unfolded over a century while the AI transition is unfolding over a decade. It is also less visible. Earlier technologies let a user build a working picture of the machine through interaction. One could watch the machine, follow the process, work out the rule. AI systems break that learning. The output arrives without the reasoning that produced it, and often without reasoning in the human sense at all.

A person can use AI fluently without any reliable account of how it works. This is the structural reason the gap widens faster than understanding can catch up. It is also the everyday face of the Pre-Code / Post-Code break — the reason an ordinary user, in an ordinary interaction, can no longer learn the system through use.

The Code After project exists to narrow this gap through translation. Not technical translation — the work of explaining how large language models work, which others do well — but institutional translation. What is changing in the laws, the economies, the credentials, the languages, the systems of evidence, and the structures of authority under which ordinary lives are lived. What the changes mean for the people affected by them. What action is still possible, in the time before the new institutions settle into place and the old ones lose the capacity to reach them.

I write for the majority of the world that does not currently have a voice in the discussion and cannot afford not to have one. Citizens, professionals, policymakers, and the institutions that serve them, in every part of the world outside the frontier economies and in most parts of the frontier economies themselves. Writing for this audience requires a specific publication objective and process, which Section VI describes, and a specific discipline in how the writing is done, which Section IX sets out in full. Together, the publication process and the writing discipline establish the series as the continuation of the work v0.9 began.

II. What Came Before

Code After: Law, Accounting, and the Governance of Artificial Intelligence (v0.9, April 2026) is the foundation of the project. It is an open-access manuscript of roughly 55,000 words that diagnoses four structural gaps through which the governance of AI fails and proposes a constitutional framework designed to close them. It was released through Zenodo with a persistent DOI and is also shared through SSRN and ResearchGate. It has begun to draw responses from scholars and practitioners across law, accounting, and AI policy.

The reception of v0.9, even at this early stage, has clarified four things.

First, the framework travels. The four gaps — Visibility, Rule-Execution, Categorical, and Measurement — describe recurring and recognisable failures of governance, and the analytical machinery developed around them extends beyond the jurisdictions the manuscript addressed. Readers in countries the manuscript did not cover have found the diagnostics applicable to their own. The break appears to be of one kind, even where the institutional starting points differ.

Second, the density is a major barrier. The manuscript was written for readers with substantial prior exposure to at least one of the multiple disciplines it synthesises, and readers without that exposure have found it difficult to hold the whole argument at once. This is a structural property of interdisciplinary work written at frontier density, and it is not correctable within the form. A second approach is required.

Third, the pace does not fit. Code After began in late 2024 as a single academic paper on the fusion of law and accounting in AI governance. By early 2025, it had expanded into a fifteen-chapter monograph of around 150,000 words intended for submission to a major university press. As the first draft came into shape, it became clear that the monograph path would take years — review, revision, production, print — during which AI would move so fast that much of the analysis might be overtaken before the book reached a reader. The decision was made in early 2026 to release a shorter open-access version at 55,000 words, structured to establish a dated scholarly record on Zenodo and to reach readers at a pace the subject matter required. That became v0.9, released in April 2026.

Even that decision has proved insufficient. In the eighteen months since the project began, AI has moved faster than any reasonable projection available at the start — more capital, more compute, more energy, more models, more capability, more competition between the U.S. and China and within each of their ecosystems for talent, capital, and market penetration. The ground the analysis stands on has kept shifting. The monograph path was too slow for the original 2025 conditions.

The open-access release was faster, and AI has still outpaced it. A publication rhythm sized to the subject matter has to assume continuous adjustment, because the subject matter does not stop moving while the analysis is being written. This is why the series is structured the way it is. Six papers at four-month intervals, each capable of standing alone, each updatable as the substrate evolves, each released at a cadence that has at least a chance of matching the ground it is describing. The work cannot outrun AI. It can, with discipline, avoid falling so far behind that it stops being useful.

Fourth, the framework reaches further than v0.9 can hold. As the work continued, it became clear that the arguments developed for governance also explain what is happening in other areas the manuscript did not address — language, education, work, evidence, measurement, jurisdiction.

The break v0.9 named in governance recurs across every domain the inherited apparatus of modern life was built to service. Going back to add another gap and another section to v0.9 would distort a manuscript already built for a different purpose. The framework's reach is real, but v0.9 is not the vehicle for it.

A separate paper on language was considered in response, and from that process the recognition followed: the framework would need a series of its own, each extension sized to its domain and written on its own terms.

These lessons did not arrive as settled conclusions. They emerged in the doing. The series proceeds on that basis, and the discipline of the work is partly the discipline of staying honest about what is still being found out.

III. The Core Thesis

The four gaps of v0.9 are instances of a more general pattern. They are what happens when Pre-Code instruments are required to reach Post-Code actors. The mechanism is decoupling.

Three terms now carry the weight of the project. The Pre-Code condition names institutional systems built on the assumption that code operates as deterministic instrumentation, subordinate to human intention and traceable through stable chains of execution and accountability. The Post-Code condition names institutional systems operating in the presence of probabilistic, adaptive, partially opaque computational actors whose outputs cannot be fully reduced to deterministic instruction or audited through inherited governance frameworks. Decoupling names the structural separation between an institution's inherited operating assumptions and the actual behaviour of the systems it attempts to govern. These definitions are stipulative, in force throughout the series, and intended to be cited together.²

Within the Post-Code condition, three further terms work at different levels and should not be used interchangeably. *Probabilistic intelligence* names the AI itself — the underlying technology. *Post-Code actor* names any AI system whose outputs materially participate in institutional processes. *Quasi-agent* names a Post-Code actor capable of shaping outcomes without being a legal subject in the inherited sense. Three layers: the AI, the AI in operation, the AI that shapes outcomes.

AI systems are opaque where inherited frameworks assumed transparency, adaptive where they assumed stability, distributed where they assumed locality, and running at machine speed where they assumed human time. These properties break the couplings on which industrial institutions rested — between representation and reality, between rule and execution, between measurement and value, between instruction and outcome, between credential and competence, between evidence and authorship.

Where a coupling breaks, the same morphology follows. A gap opens between the framework's operating logic and the subject matter's actual behaviour. Outputs continue to issue but no longer reach into the systems they are meant to govern, teach, measure, certify, or record. Intermediary actors and technical layers arise to restore function in practice, usually invisibly and usually without democratic sanction. Authority migrates from the institutions that retain the right to decide to the operators who possess the capacity to act.

Legitimacy lags behind the migration. The distance between who is accountable and who is in control widens, and the world continues to operate under a new and unacknowledged constitution. Call this the decoupling morphology: a recurring sequence in which a stable coupling breaks; legacy frameworks continue to issue outputs that no longer reach; technical intermediaries arise to restore function; legitimacy lags operational reality; and a new operational constitution settles into place without explicit ratification.

The project's central claim can be stated in one sentence. AI is a general-purpose decoupling force acting on institutional systems, and the institutions built on the assumption of stable coupling are failing in ways that follow a consistent morphology across domains. Decoupling is what the Pre-Code / Post-Code break does once it reaches an institution.

² *Institution is used broadly throughout the series to include legal, administrative, economic, educational, linguistic, and epistemic systems that organise collective social coordination.*

The series is the demonstration of that claim across six distinct domains. If the morphology holds across all six, the thesis holds. If it does not, the thesis fails, and the project will say so plainly.

The framework has limits. Decoupling describes what happens when AI meets an institutional foundation built on stable representation of the world. It does not fully describe domains that are mostly physical and immediate, that operate with little representational mediation, or that were already abstracted from underlying reality before AI arrived. Manual labour, direct sensory experience, and forms of financial engineering that were already detached from real economic activity sit outside the framework's reach, or inside it only partially. The boundary is a range rather than a wall, and the series will name it where it matters. Stating the limits is part of the discipline the project accepts. A theory that explains everything explains nothing.

The relevance of Lawrence Lessig is partial and specific. Lessig showed that in digital environments, behaviour is regulated by four interacting forces — law, norms, markets, and architecture, with architecture itself functioning as a regulatory instrument. That insight expanded what counts as a regulatory instrument and remains foundational in cyberlaw and digital-governance scholarship.³

Code After addresses a different problem. Lessig described how regulation works when the architecture is stable, legible, and authored — when "code" is the kind of thing the term meant before this project's break. The series describes what happens after that condition no longer holds. AI systems are adaptive, probabilistic, and often opaque to the institutions that were built to regulate them. Under those conditions, the question is no longer how architecture regulates. It is that the architecture of regulation fails to reach the systems it claims to govern.⁴

³ L. Lessig, *Code and Other Laws of Cyberspace* (1999); *Code: Version 2.0* (2006).

⁴ Code After is the project name; Pre-Code condition and Post-Code condition are the paired analytical terms. The first is narrative and identifies the work; the second is stipulative and does the analysis. The two registers are kept separate by design — the project name does not need to do analytical work, and the analytical terms do not need to brand.

IV. The Widening Gap

Every paper in the series will include an update. The interval between papers is four months. That is long enough for the AI field to move materially, and the analysis of any domain has to account for where the field is rather than where it stood when the previous paper appeared. The update is not a postscript. It is part of the method. Each paper will revisit the framework's earlier arguments, test them against intervening developments, and state openly where the analysis has held, where it requires revision, and where events have overtaken it.

This is the first such update. It covers the roughly eighteen months during which Code After was developed from a single academic paper, to a fifteen-chapter manuscript, to v0.9, and the period since v0.9's release on 12 April 2026.

The G2 and G3 framing has organised the manuscript's analysis of jurisdictional position from the project's beginning. G2 names the two countries, the U.S. and China, that possess the full stack of frontier AI capacity at scale; G3 names the broader set including the EU, whose regulatory and market-access leverage shapes how AI systems are deployed globally.

It is the framework's most-tested analytical claim. Over the eighteen months during which the project has been in development, the framing has not weakened. It has hardened. The concentration of frontier AI capacity within a small number of G2 labs continues, and the advantages of those labs compound while the distance between them and every other actor widens. All three are more clearly observable now than they were when the project began. The series will therefore continue to use the G2 and G3 framing, because the evidence of the period has provided no reason to abandon it and several reasons to hold it more firmly.

The gap is not only one of capability. It is one of institutional position relative to the frontier — the structural feature Section III describes as decoupling, visible now in four dimensions.

Capital

The scale of funding required to operate at the frontier continues to increase, and at exceptional speed. Valuations and capital requirements for leading foundation-model firms have reached levels that only the U.S. and China can accommodate at scale. The effective unit of analysis is no longer the country; it is the small number of labs within those two countries large enough to be capitalised at all.

The comparison to Europe's most prominent frontier firm makes the scale concrete. Mistral's September 2025 Series C closed at a post-money valuation of approximately €11.7 billion — a serious figure, reflecting both real European ambition and real European capability.

OpenAI's \$122 billion funding round in March 2026 valued the firm at approximately \$852 billion post-money. Anthropic's \$30 billion Series G in February 2026 valued it at approximately \$380 billion post-money.⁵

The gap between Mistral and the leading U.S. labs is roughly 28x relative to Anthropic and 62x relative to OpenAI. It is a gap in capital, not in technical ambition or engineering quality. If OpenAI and Anthropic

⁵ Mistral AI, *Series C announcement* (September 2025); OpenAI, *funding round announcement* (March 2026); Anthropic, *Series G announcement* (February 2026). Valuations as reported by primary company communications and corroborated in standard technology-press coverage.

complete the public offerings widely reported across mainstream financial press for late 2026, that capital gap is likely to widen further.

Capital at this scale concentrates almost entirely in the U.S. and China, mobilised by different institutional routes — primarily private capital markets in the U.S., a combination of state-aligned investment and major commercial actors in China. Actors elsewhere, including private firms in either country operating without those resources behind them, compete under structurally different conditions.

The capital concentration is reinforced by infrastructure. Frontier model firms sit on top of cloud and data-centre buildouts whose scale now rivals major public works. Q1 2026 earnings, reported on 29 April, confirmed the trajectory. The four U.S. hyperscalers — Microsoft, Alphabet, Amazon, Meta — spent on the order of \$130 billion on capital expenditure during the quarter, roughly 80 percent above the figure for Q1 2025.

Combined 2026 capex guidance from the four firms now sits in the range of \$650 to \$700 billion, nearly double the approximately \$410 billion they spent across 2025. Around three-quarters of the 2026 figure is directed at AI infrastructure — chips, servers, networking, data centres, and the power required to run them. Combined 2026 capex is on track to approach the level of the four firms' combined operating cash flows, against a long-run average closer to 40 percent.⁶

These are not marginal adjustments to an established capital cycle. They are the outline of an infrastructure regime that only a handful of U.S. and Chinese actors can enter at scale. This is the decoupling morphology operating at the level of capital. The capacity to build the systems that increasingly govern modern life is concentrated in a number of operators that can be counted on two hands. Authority migrates, in the precise sense the framework names, to the actors who possess that capacity. The institutions formally responsible for governing the resulting systems lack the resources to observe them at the scale at which they are being built, let alone to fund or replicate them.

Compute and Inference

The binding constraint is shifting from training to inference. Earlier phases of frontier development were limited primarily by training compute; current constraints are increasingly defined by the cost and availability of serving models at scale. Usage limits on the consumer products of the frontier labs have tightened over the past year, and plan tiers have been restructured upward, some substantially.

OpenAI has withdrawn Sora's consumer access over the course of 2026, discontinuing the web and app experience in April and the API later in the year. The withdrawal is consistent with the broader inference-capacity pressures visible across frontier deployment, though OpenAI has not publicly framed it in those terms. The feedback loop v0.9 described is now visible in operation. Usage produces revenue and data;

⁶ Capital expenditure figures drawn from the Q1 2026 earnings releases of Microsoft, Alphabet, Amazon, and Meta (29 April 2026) and the corresponding SEC Form 8-K filings, with mainstream financial-press aggregation used for full-year guidance synthesis. Q1 2026 combined capex on the order of \$130 billion, against approximately \$72 billion in Q1 2025, derives from per-company filings (cash payments for property, plant, and equipment, with finance leases included where reported). The 2026 full-year guidance range of roughly \$650 to \$700 billion synthesises updated investor guidance issued the same week. The 2025 baseline of approximately \$410 billion is consistent with mainstream financial-press aggregation. The AI-infrastructure share (approximately three-quarters of 2026 spend) and the capex-to-operating-cash-flow ratio (approaching 100 percent in 2026, against a long-run average closer to 40 percent) reflect standard analyst aggregations, including UBS, Bank of America, Epoch AI, and other major sell-side and independent research houses.

revenue and data fund compute; compute enables more capable models; more capable models attract further usage. The frontier is extending its lead through the same mechanism that created it.⁷

The Application Layer and Agents

Agentic systems have moved from frontier novelties into broader consumer and enterprise integration. The pattern is visible at two levels.

At the consumer level, Manus was among the first widely visible general-agent products from the Chinese ecosystem. In 2025 it drew intense attention in China — domestic coverage framed it as a "GPT moment for agents" and as the second major Chinese AI breakthrough of the year after DeepSeek — establishing it as a focal example of how rapidly agent capability had crossed into ordinary consumer use.

The case then took on a second, geopolitical, dimension. Meta announced its acquisition of Manus in late December 2025 in a deal reported at approximately \$2 billion. On April 27, 2026, China's National Development and Reform Commission (NDRC) announced that it would prohibit foreign investment in Manus, citing foreign-investment and national-security grounds, and ordered the parties to unwind the transaction.

The case carries a dual lesson. The first half is the one already visible: agent capability is reaching ordinary users at speed. The second is newer. Sovereign action is now drawing, in real time, the cross-border boundaries of access to frontier AI capability. Within the broader process of U.S.–China decoupling in AI, the Manus case marks a clear inflection point: the cross-border flow of frontier AI assets is no longer determined by market logic alone, but is increasingly subject to the hard-edged constraints of national-security review.

At the infrastructure level, the pattern presents differently. Open-source agentic frameworks have emerged as a distinct category. OpenClaw is the most visible example. Uptake within China's developer community and beyond has been rapid. Reports describe nearly a thousand people queuing outside Tencent's Shenzhen headquarters for installations. Local governments have subsidised deployment events. Cloud providers and regional actors have forked and repackaged the codebase at scale. Audited user metrics remain limited, but the pace of replication and the intensity of public and developer discourse around the framework suggest that market reception has far exceeded initial industry expectations.

OpenClaw functions as the *Orchestration Layer* between foundation models and user-facing applications, managing state and memory, tool and API access, permissions, and routing across multiple model backends. It does not compete with frontier models. It makes them work. The orchestration layer is the path by which the concentrated capability sitting at the top of the stack reaches the users, applications, and workflows at the bottom.⁸

⁷ OpenAI, *Sora consumer access withdrawal* (April 2026); *API discontinuation announced for September 2026*. Company communications and standard technology-press coverage at the time of release.

⁸ On Manus's March 2025 launch and immediate public reception in China, see China Daily (6 March 2025) describing the product as "the GPT moment for AI agents" and as the second major Chinese AI breakthrough after DeepSeek, and Tech Xplore (March 2025) on the invitation-only debut and ensuing public reaction. On Meta's late-2025 acquisition announcement, the subsequent regulatory review by Chinese authorities, and the NDRC's formal prohibition of the transaction, see Bloomberg (27 April 2026) on the NDRC's decision to prohibit foreign investment in Manus and to order cancellation; CNBC (27 April 2026) on the NDRC's role as state planner and on the framing of the action under foreign-investment and national-security grounds; and parallel reporting across major Western financial and technology press, including Wall Street Journal, The Washington Post, Reuters, and

The emergence of this middle layer brings the real shape of decoupling into focus. The AI stack now resolves into three distinguishable layers. The foundation layer, where a small number of frontier models sit, is consolidating. The orchestration layer, where frameworks like OpenClaw assemble models into systems of action, is becoming more open and portable. The application layer, where users meet these systems through products, interfaces, and workflows, is diffusing at speed.

This is no coincidence. The opening of the orchestration layer and the diffusion of the application layer do not disperse frontier capability; they deepen dependence on it. Frameworks like OpenClaw, by routing agents across multiple providers, reduce vendor lock-in at the orchestration level — but the quality of system cognition still depends on the quality of available models. Orchestration can amplify, structure, and stabilise intelligence. It cannot manufacture frontier-grade intelligence from non-frontier foundations. The easier it becomes to build agentic systems on top of the best frontier models, the stronger the incentive to use those models rather than to build alternatives. Application diversity has not distributed AI sovereignty. It has widened the demand funnel through which centralised intelligence is further reinforced.

The shift from conversation to action changes the stakes. Early AI systems answered questions and humans did the work; agentic systems increasingly do the work themselves — drafting and sending emails, scheduling and rescheduling meetings, organising files, moving funds, completing transactions. The user sees the result. The user does not see the chain of decisions and tool calls that produced it. Once AI not only recommends action but takes it, the distance between what the user sees and what the system does becomes a question this series will return to repeatedly in subsequent papers.

The application layer is expanding. The orchestration layer is opening. The foundation layer is consolidating. These three motions push the concentration v0.9 described in the same direction — deeper, not wider.

World Models and Physical AI

The same pattern of concentration is extending beyond language. Work on systems that model and predict the physical world is moving from research into early product discussion, on three parallel tracks. Major frontier labs including xAI, Meta, and Google DeepMind are extending their language-model capability into the physical domain, with humanoid programmes such as xAI’s Optimus push moving from prototype toward deployment. The Chinese ecosystem is doing parallel work, anchored by Unitree, the broader humanoid manufacturing base, and industrial-AI deployment programmes already operating at scale.

A third track has emerged through dedicated start-ups built specifically around world-model construction, establishing themselves as a distinct strategic line in the field. Two of the most credentialed researchers in artificial intelligence have stepped away from their institutions to lead this work. Fei-Fei Li — whose ImageNet project a decade earlier catalysed the deep-learning era — has taken extended leave from

TechCrunch, on the same date, alongside coverage in Chinese state and financial media including Caixin and China Daily and the NDRC’s own statement of the same date. On OpenClaw, see Chinese and international technology and financial media coverage of its rapid 2026 adoption — including reporting on queues for installations at major Chinese technology firms, local-government subsidies for deployment events, and investor enthusiasm — together with documented forking and repackaging of the codebase and its precursors (notably Clawdbot and Moltbot lineages) on GitHub. Adoption here refers to observable ecosystem replication — forks, derivative projects, developer commentary, media coverage — rather than to audited user metrics, which remain limited.

Stanford to found World Labs, a *spatial intelligence* company whose core mission is the construction of large-scale world models.

Yann LeCun, recipient of the A.M. Turing Award for his foundational work on deep learning, has departed Meta and co-founded AMI Labs (Advanced Machine Intelligence) to pursue the same class of systems. His public position is that large language models alone cannot reach artificial general intelligence. AGI is the threshold at which a system exhibits cross-domain autonomous reasoning and generalisation comparable to human cognition. World models, LeCun argues, are the architecture that can.

This view is no longer confined to research-side advocacy. In April 2026, the Goldman Sachs Global Institute framed world models as a decisive next step in AI and argued that current AI-infrastructure projections, built around language-model scaling, materially understate the capital, simulation, and physics-engine demands of the world-model track. In the same week, MIT Technology Review's new "10 Things That Matter in AI Right Now" list identified world models as one of the field's most consequential current developments. The thesis that the next layer of AI capability lies beyond language is now converging across research labs, finance, and technology press — within weeks of this paper's planned publication.

World models train on different data altogether. They learn from what cameras see, what microphones hear, what touch and motion sensors record across thousands of synchronised channels — visual, acoustic, kinetic, and physical streams drawn from real and simulated environments rather than from text. The sensing, modelling, validation, and deployment infrastructure these systems need exists today in fragments. It is not yet at the scale or maturity that broad deployment will require.⁹

Humanoid robotics makes the point concrete. Over the past eighteen months, Unitree, Figure, Tesla, and the broader humanoid ecosystem have moved these systems from laboratory curiosity to commercial roadmap. Chinese activity is especially visible — factory pilots, training centres, industrial-policy initiatives. Humanoid robots are not world models. They are among the most important physical platforms through which the multimodal embodied data world models need can be generated at scale, alongside simulation environments, industrial sensor networks, and autonomous vehicles.¹⁰

Language-model capacity already concentrates in the G2. World-model capacity is likely to concentrate further. The data, the robotics manufacturing base, and the sensor infrastructure all concentrate there more

⁹ On Fei-Fei Li's foundational work, see J. Deng et al., *ImageNet* (2009), and A. Krizhevsky, I. Sutskever, and G. Hinton, *AlexNet* (2012). On Li's 2024 leave from Stanford and the founding of World Labs, see World Labs' launch communications and September 2024 coverage of its USD 230 million funding round in *Forbes*, *Reuters*, and *TechCrunch*. On Yann LeCun, see ACM, 2018 A.M. Turing Award announcement (March 2019), jointly to Y. Bengio, G. Hinton, and Y. LeCun. On LeCun's late-2025 departure from Meta and the co-founding of AMI Labs, see *TechCrunch* (December 2025; March 2026), *MIT Technology Review* (January 2026), and AMI Labs' public materials at amilabs.xyz. On LeCun's published position on LLMs and AGI, see his *MIT Technology Review* interview (January 2026) and the AMI Labs mission statement. On the convergence of finance-side and technology-press analysis on the world-model thesis, see G. Lee and D. Keyserling, *When AI Learns How the World Works* (Goldman Sachs Global Institute, April 2026), and *MIT Technology Review*, "10 Things That Matter in AI Right Now" (April 2026).

¹⁰ On the frontier-lab physical-AI push, see Tesla and xAI's public communications on the Optimus humanoid robot and the Digital Optimus / Macrohard initiative, Meta's world-model research portfolio (including V-JEPA 2), and Google DeepMind's Gemini Robotics publications. On humanoid commercialisation 2024–2026, see Unitree's product and shipment communications, Tesla's Optimus programme updates, Figure's pilot-deployment announcements, and standard technology-press coverage. On Chinese humanoid deployment, training infrastructure, and industrial-policy support, see Chinese and Hong Kong technology-media coverage in *36Kr*, *TechNode*, and the *South China Morning Post*, alongside Chinese state-media coverage of national humanoid-industry initiatives.

strongly still. The dynamic that shaped the language-model era is extending into the era of physical intelligence. The gap v0.9 described is not stabilising at the language-model frontier. It is moving.

This is what makes the series' first paper urgent rather than thematic. The Language paper is not a neutral topic. Language is where the gap first opened. Language is where everything else rests. A country that cannot govern in its own language against AI systems that think in another will not govern the embodied, physical-world systems that come next. The same is true for every domain the series addresses. The work begins with language because language is where the gap first became visible and where it is still most correctable.

The framework has held. A reviewer could reasonably have argued a year ago that the G2 framing overstated a temporary condition. The evidence since does not support that reading. The gap is widening along every dimension the framework identified — capital, compute, data, talent, models, infrastructure — and now extending into a further dimension, physical-world data, that will widen it more.

The decoupling described in Section III is intensifying, not normalising.

Each subsequent paper will open with an update of this kind. The framework will be tested against the previous four months. The analysis will state openly where it has held and where it has had to evolve. AI does not wait. The project will not pretend it does.

The updates will track both G2 ecosystems on equal terms, the U.S. and China, without filtering one through the lens of the other. Most current AI analysis treats one ecosystem as the default frontier and the other as competitive context. Code After treats both as the frontier, because both are.

The project exists to connect what is happening at the frontier to the lives AI is rebuilding. It is happening to everyone at once — the lawyer, the accountant, the doctor, the teacher, the CEO, the politician, and the citizen scrolling the news on their phone. The reach is the goal. The series is built to deliver the analysis to readers who need it, in the languages they read in, at a pace that matches what is actually changing around them. How the series is built to do this is the subject of the sections that follow.

Code After is a framework, a series, and a living record of its own claims against what actually happens.

V. The Series

Expert opinion on AI ranges from the age of abundance to existential collapse. Both extremes are widely held by serious people. Code After takes neither position as its starting point. The project proceeds from a different premise: AI will change people's lives and the institutions they depend on in ways that produce both gains and costs, and the analytical task is to examine, domain by domain and in real time, what is actually changing, what dangers and opportunities are emerging, and what institutions can still do in the time the change leaves them.

AI's reach is too broad for any single project to cover completely. The current series selects six domains in which its immediate institutional effects are most visible and most consequential — language, education, work, evidence, measurement, and jurisdiction. Each will surface its own combination of opportunities and dangers. The Language paper diagnoses what the framework now identifies as a structural threat: there are roughly seven thousand living languages spoken today, and frontier AI's linguistic base is accelerating its convergence onto a small handful of dominant ones.¹¹

With few exceptions, the institutional standing of those seven thousand languages now faces structural exclusion from the AI layer. The disruption is not a local technical iteration. It is a structural reorganisation of the global linguistic ecosystem, and its urgency and breadth far exceed the present awareness of most citizens and policymakers. Other domains will need treatment in subsequent series as the transition deepens. Code After is the beginning of a longer project, not its complete map.

Each paper is archived on Zenodo in both English and Chinese originals with a persistent DOI, and disseminated through ResearchGate and SSRN. Local-language editions will be introduced progressively. Where possible, they appear in parallel with the dual originals rather than in sequence, to reduce the lag between analysis and reach. The project's primary home is codeafter.ai. The site is being built to host all of the project's outputs — papers, companion essays, distribution materials, local-language partner editions as they appear, Updates as each paper publishes — and to evolve as audiences and their needs evolve. It is built for the readers, not for the authors. Section VII develops the full publication architecture.

The six subjects are language, education, work and professional credentialing, evidence and truth, measurement and value, and jurisdiction and territory. Each foregrounds a core institutional function. Language is representation. Education is capability formation. Work is value allocation. Evidence is truth

¹¹ On the number and status of living languages, see D. Eberhard, G. Simons, and C. Fennig (eds.), *Ethnologue: Languages of the World*, 29th ed. (SIL International, 2026), reporting 7,170 living languages in use today; see also *Ethnologue*, "How many languages are there in the world?", noting that roughly 44% of all languages are endangered, often with fewer than 1,000 users remaining, and that the world's twenty largest languages are spoken natively by more than 3.7 billion people — 0.3% of the world's languages accounting for nearly half of its population. On global linguistic-diversity decline, see D. Harmon and J. Loh, "The Index of Linguistic Diversity: A New Quantitative Measure of Trends in the Status of the World's Languages," *Language Documentation & Conservation* 4 (2010), finding a 20% decline in global linguistic diversity between 1970 and 2005, with sharper declines among Indigenous languages of the Americas and the Pacific. On how frontier AI systems are deepening linguistic exclusion of non-English and low-resource language communities — adding a new vector of pressure on top of an already-fragile ecosystem — see Stanford HAI, "Mind the (Language) Gap: Mapping the Challenges of LLM Development in Low-Resource Language Contexts" (2025); Stanford Report, "How AI is leaving non-English speakers behind" (May 2025); and Ada Lovelace Institute, "Now you are speaking my language: why minoritised LLMs matter" (28 November 2024). On UNESCO's monitoring of language status, see UNESCO, *World Atlas of Languages*, current online edition; see also UNESCO, *Atlas of the World's Languages in Danger*.

verification. Measurement is economic description. Jurisdiction is authority. The ordering is not arbitrary. It tracks the order in which the institutional functions depend on one another.

Language comes first because linguistic sovereignty is prior to the rest. A country that cannot govern in its own language against AI systems that think in another loses ground in everything else. Education, work, evidence, measurement, and authority all operate through language. When the language fails, they fail with it.

Education follows because the gap between work produced and capability acquired is already visible in daily life. Parents, teachers, and students are living it now. The credentialing systems that convert education into economic opportunity are losing the signal they were built to carry.

Work and professional credentialing extends the education question into the labour markets where its consequences land. The signals by which skill is recognised, hired, and paid are losing their clarity. The professions that rest on those signals are adjusting in ways that have not yet been named.

Evidence and truth addresses the systems through which we tell authentic from fabricated content — in courts, newsrooms, archives, and the ordinary verifications of daily life. Those systems are straining against tools built to produce plausibility at scale. The strain is most dangerous where institutional legitimacy depends on the distinction holding.

Measurement and value extends v0.9's Measurement Gap beyond accounting. The economic infrastructure of productivity metrics, national statistics, and economic categorisation rests on industrial-era measurement that cannot price or count intelligence that updates itself. The frameworks built on those measurements are beginning to drift from the economies they are meant to describe.

Jurisdiction and territory closes the series by returning to the sovereignty question with which v0.9 began. The physical basis of sovereign authority is being renegotiated against computation that recognises no border and cognition that no longer stays where it was produced. The paper will rest on what the five earlier papers establish.

Each paper follows the same structure. It names the subject matter and establishes why AI is rebuilding its foundational logic. It diagnoses the gap that has opened between the inherited framework and what AI now does, naming it precisely. It traces the decoupling mechanics across the gap's operating levels. It applies the Incorporation Heuristic — $I = V \times W \times N$ — developed in v0.9, showing what Visibility, Workability, and Necessity mean in the domain. It proposes the structural response the domain requires.

The structure is not optional. Every paper must reuse and stress-test the same engine. A subject that cannot sustain the full five-move structure is not a Code After paper, and will not be written as one. The template is a methodological filter. The discipline is intentional.

VI. Layered Outputs

A single register cannot serve every reader the project is built to reach. The academic reader needs citations, methodological detail, and engagement with the relevant literatures. The policymaker needs concrete implications for the decisions they face. The professional reader needs diagnosis in the domain they work in. The general reader needs access to the argument without first becoming fluent in the disciplines it draws on. Writing for all four at once produces a register that serves none of them well.

The series addresses this through three layers: academic papers, companion essays, and a distribution layer of shorter public-facing pieces.

The academic paper is the analytical core. Full density, full citation, written for the scholarly readers whose engagement will test and refine the framework. Each of the six papers in the series will be led by an academic paper of this kind.

The companion essay is the reach. Shorter, less technical, written for the general reader and the policymaker. It carries the argument without the scholarly apparatus, shows what the argument means in its domain, and does most of the work of meeting readers who would not otherwise encounter the analysis. The companion essay deserves at least the editorial attention the academic paper receives, because most readers will arrive through it.

The distribution layer is the entry point. Op-eds, policy briefs, presentations, interviews, and explanatory material on codeafter.ai will extend specific arguments from the series into specific public conversations as occasions require.

The three layers work together as a pathway rather than as parallel products. The academic paper establishes the analysis. The companion essay translates it. The distribution layer delivers it into the conversations where the argument has to land. A reader entering through a distribution piece finds the companion essay standing behind it. A reader entering through the companion essay finds the academic paper behind that. A reader entering through the academic paper finds the framework that ties the whole series together.

The aim is fidelity across registers, not simplification. Each layer is written at the density its audience can work with. Each carries the same argument. The academic paper holds the full analysis. The companion essay carries the argument in plainer terms. The distribution layer meets readers where the conversations they are already in.

The discipline applies across languages as well as registers. Translation of frontier material whose vocabulary is not yet stable in the target language requires skilled adaptation, not mechanical conversion. Short declarative sentences travel across languages better than long subordinate constructions. The voice discipline is therefore part of the multilingual architecture from the start, not a stylistic preference layered on top.

The specific forms will develop as the series progresses and as readers emerge. The commitment to a layered architecture is fixed.

VII. Publication Architecture and Bilingual Commitment

The publication decisions for the series follow from the constraints identified earlier. AI-era institutional change moves faster than the monograph cycle, and the readers most affected by the change cannot wait for the cycle to complete.

Each paper is archived on Zenodo as the versioned record, with a persistent DOI, and disseminated through SSRN, ResearchGate, and codeafter.ai. The persistent DOI provides the scholarly record. Open access provides the reach. Version control allows the work to be updated as the subject evolves.

A v1.0 edition of the primary manuscript will be produced once the series has completed. It will consolidate the framework with the accumulated findings of the six papers and will be submitted to a major university press for formal academic publication. The sequence — open-access series first, consolidated monograph later — is not a retreat from academic ambition. It is the route to it that the subject matter permits.

The project is bilingual at the structural level, not at the translation level. English and Chinese are the two primary languages of frontier AI development. The training data, the research literature, the model documentation, and the technical discourse all run in these two languages. A project that analyses this world in only one of them captures only part of what it claims to describe. The system being analysed is already bilingual. The analysis must be as well.

codeafter.ai is bilingual English–Chinese from launch. Chinese editions of each paper are produced in parallel with the English originals through direct authorship, collaboration, or high-quality translation as circumstances permit, with cadence kept as close to the English release as production allows.

Additional language editions — Spanish, Arabic, French, Hindi, Japanese, Korean, Swahili, Indonesian, and others — will follow where the project's purposes require them and where genuine collaborators can be found. The initial selection itself is part of the project's analytical work. Rather than translate into a fixed set of high-prominence languages, the series uses each language edition as evidence about the conditions under which AI-relevant legal and institutional argument travels or fails to travel. The criteria are language family, resource tier, script system, and the presence of local collaborators positioned to extend the work. The Language paper, first of the six, will itself inform which additional languages deserve priority by naming the conditions under which the Linguistic Gap is most urgent.

Collaborators on local-language editions are acknowledged formally and credited as contributors to the series, not as translators of completed work. Local editors are responsible for reconstructing the argument in the register of their own language, embedding cases, references, and institutional context that anchor the work for local readers. They also identify, in that setting, how AI's effects are experienced, mediated, and contested — what the framework's predictions actually look like on the ground in their jurisdiction. Each local edition is, in this sense, both an output and a research site.

The publishing architecture and the layered outputs together form a system designed against latency. Analysis does not wait for translation; translation does not dilute analysis. Given that language is where institutional reach first breaks down under AI, parallel-language publishing is not a publishing strategy but a methodological premise. The Language paper will map that break in full, and will show why the erosion of linguistic sovereignty is now moving faster than most governance frameworks can perceive it.

VIII. The Partworks Model

The publication model the series is adopting has a precedent worth naming. Partworks publishing — the serial release of a larger body of work in instalments, with local editions produced by regional publishers under licence — was one of the most effective mass-market knowledge-distribution architectures of the late twentieth century. Italian publishers led the international expansion. RCS Libri (Rizzoli), De Agostini, and Fratelli Fabbri Editori exported partworks titles across markets and subject categories well beyond Italy.

I know this model from direct experience. Beginning in the early 1990s, I published the first Chinese editions of *PC Magazine*, *PC Week*, *PC Computing*, *Institutional Investor*, and *Popular Science*, among other titles, under licence from their originators. ZDNet was not a magazine but one of the earliest digital platform plays in China, which I built and operated under licence from the same period. I also published hundreds of global book titles in technology, finance, and the sciences, and produced partworks in China through a joint venture with RCS Libri. The computing titles in particular put me at the start of the computer and internet revolution in China. AI is the revolution now succeeding it.

What that period taught was practical. Serious content reaches general publics most reliably through fixed cadence, cumulative architecture, and local editorial partnership — not through centralised publication in a single jurisdiction or language. The knowledge travels because the structure lets it.

The Code After Series adapts that logic to the AI era. The cadence is four months per paper rather than fortnightly. The distribution is open access rather than subscription, because the mission is public good. The aim is not to reproduce the partworks model but to apply its underlying principles to a subject that no community will be able to retire from. AI will not go away for any community. The conversations and contestations it requires will happen in every community's own language — or they will not happen at all. A multilingual mechanism for sustained engagement is the only architecture that lets those conversations be honest in their own register and accountable to their own institutions.

A. Incremental Entry

The architecture accommodates partners joining at any point. A publisher may take on Paper 1 alone and decide whether to continue. A publisher may enter at Paper 2 or 3, with earlier papers available as back-list. A publisher may take on all six as a committed programme, or select a subset by domain relevance to the local market. Arriving later does not foreclose access to what came before. Partners can begin when they are ready, and they can always go back.

B. The Asymmetry to Acknowledge

The partworks analogy under-describes the operational challenge. Historical partworks succeeded because local publishers carried commercial risk and the originator supplied content. The Code After version inverts this. The originator carries the content risk. Local partners are asked to invest in adaptation for markets where commercial return is uncertain. The model can work — but only with partners who recognise what they themselves bring, which is editorial judgement, local institutional knowledge, distribution capacity, and the situated evidence that no central operation can produce. The partnership cannot be framed primarily in terms of what the partner receives. The partners most worth having understand this without needing it explained.

C. Licensing and Terms

The series is released under CC BY-NC-ND. The project operates a dual-licence model: the public licence governs unmodified, non-commercial sharing, while bespoke partnership licences cover adaptation and commercial use by official partners. Non-commercial academic and policy sharing of the unmodified original materials is permitted under the public licence with attribution and without separate agreement.

Local-language editions are adaptations and require a separate licence, whether or not they are produced commercially. Non-commercial institutional editions may be handled through simplified academic or institutional terms. Commercial local-language editions — print runs, subscription distribution, bundled-with-subscription delivery, co-branded professional publications — require commercial terms in addition, negotiated bilaterally with the partner.

The default structure has three elements: an exclusive local-language licence for a defined language, territory, format, and term; a revenue-share arrangement calibrated to the partner's market and production costs; and editorial cooperation that preserves the framework's analytical structure, terminology, and attribution across editions while leaving local editorial decisions to the local editor. The terms scale to the partner's situation. The editorial cooperation does not.

D. Contact

Partners interested in discussing a local-language edition should write to partners@codeafter.ai. The same address handles regional imprints, institutional collaborations, and bespoke publication arrangements. Inquiries are welcome from publishers, academic institutions, policy research organisations, professional services firms, and any other organisation positioned to produce serious work in a language the series does not yet serve. A short note describing the partner, the proposed language or market, and the interest in the series is sufficient to open the conversation.

IX. The Voice

A word on register, because the writing is itself part of the argument.

Frontier scholarship has accumulated habits that often make it unreadable to readers outside its own conversation. The habits are not unreasonable. They signal disciplinary membership. They protect against overclaim. They preserve the epistemic caution the disciplines have earned through long practice. But the habits have a cost. They wall off the frontier from the readers who most need access to it, and the wall is growing.

The Code After register is designed to remove the wall without removing the rigour. Claims are stated directly, in short declarative sentences where possible, with qualification placed where the argument requires, not as default posture. Jargon is reserved for terms that do real analytical work. Examples are chosen to land for a reader outside the discipline the example comes from. The writing is serious without being solemn, confident without being aggressive, composed without being cold.

The register is a discipline, not a finished achievement. It is enforced in editing as well as in initial drafting. Every paragraph is tested against the standard, and where the standard slips the paragraph is rewritten.

The register has to carry across translation without losing its character. Writing for translation from the start is writing better English. Short sentences survive. Concrete nouns survive. Active verbs survive. Excessive hedging and heavily nested subordinate clauses do not, and have been pared back for reasons beyond clarity in English alone.

This matters for the bilingual and multilingual architecture the series has committed to. The voice is a shared standard across every edition. Chinese editions, local-language editions, and future adaptations are not expected to reproduce English phrasing. They are expected to reproduce the same clarity, the same discipline, and the same directness in their own language. The objective is equivalence of function, not literal correspondence of form. Collaborators and publishers joining the series inherit the standard along with the framework.

The register is the first test of the project. If the writing cannot reach the reader, the mission fails regardless of the analytical quality behind it. The series is a research program with public-facing outputs, not a set of essays with a shared brand — and the voice is what makes that distinction operational. Every paper in the series will be judged on whether the voice holds. So will this document.

X. Why This, Why Now

The AI transition is not negotiable. It is not escapable. It is not reversible on any timescale that matters to the people alive now. It will affect the countries that built the AI foundations and the nations that did not, the industries that invested early and the companies that did not, the generations that grew up with it and the generations that will inherit what it becomes. It will not wait for the institutions responsible for governing it to catch up with what it is. It will not wait for the public responsible for those institutions to catch up with what is happening to them.

The Code After Series is an attempt to give institutions and the public a framework they can use in the time that remains before the new architecture hardens. The attempt may fail. The framework may be wrong. The writing may not reach. The pace of AI may outstrip the pace of the analysis. The readers who need the work most may never encounter it. These failure modes are understood, and accepted as the cost of attempting the work at all.

What is not acceptable is not attempting it. The subject matter is too large and the stakes too widely distributed for the absence of the work to be defensible. If the series fails, it will have failed at something worth failing at. If it succeeds, it will have helped give the majority of the world a framework through which to understand and engage with a transition that has so far been conducted without them.

The work now moves from argument to publication.

*The Code After Series is hosted at codeafter.ai. The primary manuscript, *Code After: Law, Accounting, and the Governance of Artificial Intelligence* (v0.9, April 2026), is archived on Zenodo (DOI: 10.5281/zenodo.19537473). The first paper of the series, *Code After Language: The Linguistic Gap and the Sovereign Language Stack*, will appear in September 2026. A v1.0 edition of the primary manuscript will be produced upon completion of the series and submitted for formal academic publication. Publishers and institutional partners interested in local-language editions should write to partners@codeafter.ai.*

Appendix A. Terms from v0.9 Used in This Paper

Categorical Gap (v0.9, Part II) The structural mismatch between legal categories designed for deterministic, territorial, human-driven activity and AI systems that are none of these things — producing a thinning of meaning in which legal classifications cannot attach cleanly to technical realities.

G2 (v0.9, Part I) An analytical category referring to state-level AI ecosystems that meet the threshold of vertically integrated capability across hardware, compute, model development, talent concentration, and capital scaling. The United States and China are the clearest G2 states. The term is descriptive rather than geopolitical.

G3 (v0.9, Part I) States whose regulatory frameworks propagate globally through market size, legal-harmonisation capacity, and institutional legitimacy, even without full integration across the AI stack. The EU is the clearest G3 actor alongside the G2 states.

Incorporation Heuristic ($I = V \times W \times N$) (v0.9, Part IV) An analytical model explaining why some sovereign rules propagate into global practice and others do not, as a function of three variables: Visibility (detectability and credible enforcement), Workability (operational implementability), and Necessity (market gravity making exit irrational). The multiplicative structure means failure at any variable is dispositive.

Measurement Gap (v0.9, Part II) The rupture that forms when industrial-era accounting frameworks confront probabilistic, self-updating capital — limiting the State's capacity to price AI's contribution even when it can regulate AI's conduct.

Rule-Execution Gap (v0.9, Part I) The structural separation between the declaration of a rule and its realisation inside the systems meant to execute it — the distance a rule must travel from legislative intent to technical implementation.

Visibility Gap (v0.9, Part I) The structural asymmetry in which regulators cannot observe the systems they claim to govern — composed of three dimensions: technical opacity, jurisdictional fragmentation, and temporal acceleration.

Appendix B. Terms Developed in This Paper

Bilingual Structural Commitment (Section VII). The project's commitment to operate structurally in English and Chinese, not merely in translation — the two primary languages of frontier AI development. *codeafter.ai* launches as bilingual from the outset, with additional local-language editions produced in parallel as both outputs and research sites.

Code After Register (Section IX). The voice discipline applied across all editions and languages — direct claims in short declarative sentences, qualification where the argument requires, and jargon reserved for terms doing real analytical work. The standard governs editing and drafting alike, and distinguishes the series as a research programme from a set of essays under a shared brand.

Companion Essay (Section VI). The second layer of Layered Outputs. It is shorter and less technical than the academic paper, accompanies each substrate paper, and is written for the general reader and the policymaker, carrying most of the project's audience-reaching work.

Decoupling (Section III). The structural separation between an institution's inherited operating assumptions and the actual behaviour of the systems it attempts to govern — the mechanism through which Pre-Code instruments fail to reach Post-Code actors. AI is a general-purpose decoupling force acting on institutional systems; the series is the demonstration of that claim across six areas of institutional life.

Decoupling Morphology (Section III). The recurring sequence that follows a coupling break: a stable coupling breaks; legacy frameworks continue to issue outputs that no longer reach; technical intermediaries arise to restore function; legitimacy lags operational reality; and a new operational constitution settles into place without explicit ratification. The content is specific to each area of institutional life; the structure is identical across cases.

Distribution Layer (Section VI). The third layer of Layered Outputs — op-eds, policy briefs, presentations, interviews, and explanatory material on *codeafter.ai* that extend specific arguments from the series into specific public and policy conversations.

Five-Move Template (Section V). The standard analytical structure each substrate paper in the series applies to its domain: name the subject, diagnose the gap, trace the decoupling mechanics, apply the Incorporation Heuristic, propose the structural response. A paper that cannot execute all five moves is not a Code After paper.

Layered Outputs (Section VI). The three-tier publication architecture for each substrate paper: an academic paper as analytical core, a Companion Essay for reach, and a Distribution Layer of shorter public-facing pieces. The three layers work as a pathway rather than as parallel products.

Linguistic Gap (Section V; subject of Code After Language). The structural mismatch between the seven thousand living languages spoken today and frontier AI's accelerating convergence onto a small handful of dominant ones. The first domain-specific Gap the framework has diagnosed beyond the four Gaps of v0.9, and the subject of the first paper of the series.

Linguistic Sovereignty (Sections V, VII). The capacity of communities to conduct education, administration, law, and public reasoning in their own languages — the analytical category through which the Linguistic Gap operates. Its erosion is now moving faster than most governance frameworks can perceive, which the Language paper will diagnose in full.

Partworks Adaptation (Section VIII). The series' publication model, drawing on the partworks tradition of serial release and local adaptation and applied to the AI era — open access rather than subscription, four-month rather than fortnightly cadence, and an inverted commercial-risk structure under which the originator carries content risk while local partners invest in adaptation.

Post-Code Actor (Sections I, III). Any AI system whose outputs materially participate in institutional processes. The term actor is functional rather than metaphysical — a system whose outputs participate materially in institutional processes regardless of whether the system itself possesses agency in the human sense.

Post-Code Condition (Sections I, III). Institutional systems operating in the presence of probabilistic, adaptive, partially opaque computational actors whose outputs cannot be fully reduced to deterministic instruction or audited through inherited governance frameworks. The Post-Code and Pre-Code conditions together name the structural break the Code After project is built to address.

Pre-Code Condition (Sections I, III). Institutional systems built on the assumption that code operates as deterministic instrumentation, subordinate to human intention and traceable through stable chains of execution and accountability. The distinction is structural rather than chronological — Pre-Code does not name a period before software but a regime in which code, however sophisticated, behaved as deterministic instrumentation subordinate to human intention.

Probabilistic Intelligence (Sections I, III). The AI itself — the underlying technology from which Post-Code actors are built. The first level of a three-level hierarchy used throughout the series: probabilistic intelligence (the technology), Post-Code actor (the technology in operation), quasi-agent (the technology that shapes outcomes).

Quasi-Agent (Sections I, III). A Post-Code actor capable of shaping outcomes without being a legal subject in the inherited sense. The third level of the actor hierarchy used throughout the series — the AI that does not merely inform institutional decisions but moves them, and that consequently sits at the centre of the decoupling problem.