

CODE AFTER

Law, Accounting, and the Governance of Artificial Intelligence

Richard Yan

Email: richard.yan@codeafter.ai

ORCID: <https://orcid.org/0009-0000-7611-6323>

v0.9 — Open-Access Pre-Publication Release

Released: April 12, 2026

Copyright © 2026 Richard Yan.

Licensed under Creative Commons Attribution – NonCommercial - NoDerivatives 4.0 International (CC BY-NC-ND 4.0). Attribution required. Commercial use and derivative works are prohibited.

This v0.9 edition is an open-access pre-publication release archived on Zenodo to provide a persistent DOI and establish a dated scholarly record. A final, expanded v1.0 edition will follow.

Suggested citation: Yan, R. (2026). *Code After: Law, Accounting, and the Governance of Artificial Intelligence* (v0.9). **DOI:** 10.5281/zenodo.19537473

Abstract

Artificial intelligence (AI) is often framed as a technological innovation subject to regulatory oversight. This manuscript advances a different thesis: AI represents a structural shift in the conditions under which governance operates — one that alters the basic prerequisites of public authority. Traditional sovereign models presuppose visibility, interpretability, and enforceable rule execution. AI systems unsettle these assumptions in ways that exceed conventional regulatory categories.

The manuscript identifies two core constraints: the Visibility Gap, describing the divergence between system behaviour and human oversight, and the Rule-Execution Gap, describing the separation between formal rules and their operational realisation for AI. Together, these gaps give rise to a Visibility Monopoly, where hyperscalers' privileged access to their own models combines with the Translation Layer — the global advisory ecosystem of law firms, audit networks, and consultancies that operationalise compliance — to determine how behaviour, value, and agency become legible. This position, historically occupied by the State, now shapes every downstream governance response.

The manuscript develops the Dual-Compiler Thesis, arguing that governance depends on two constitutive institutional languages: law, which defines permission for AI action, and accounting, which measures AI through classification and valuation. The Incorporation Heuristic models the capacity of nations to internalise AI activity based on visibility, workability, and necessity — reframing regulatory competition as differential alignment between national capability and computational reality. Gateway Rules and the Golden Image describe how regulatory requirements propagate across borders via system design and professional services infrastructure rather than formal sovereign adoption.

The manuscript then turns to the material foundations of AI governance. It identifies capacity primacy — the structural condition in which control of computational resources conditions the effectiveness of legal and regulatory authority. Five non-substitutable constraints (energy, semiconductors, capital, talent, and time) determine whether a state can govern AI on its own terms. Legitimacy Debt accumulates when formal responsibility and real decision authority drift apart, producing a widening gap between statute and practice.

Finally, the manuscript proposes a Third Architecture: a governance framework built on a Sovereign API that translates democratic authority into machine-executable constraints; a Protocol of Federation that embeds national rules within networked economic infrastructure; and Parametric Democracy, which preserves democratic legitimacy by shifting regulation from rule authorship to constraint calibration. This architecture is designed to preserve sovereign diversity while establishing the institutional conditions under which AI systems remain accountable to public authority.

This v0.9 open-access release presents the core conceptual architecture of *Code After* to support early scholarly, policy, and institutional engagement. A forthcoming v1.0 edition will integrate full citations, empirical grounding, and jurisdictional applications.

TABLE OF CONTENTS

PREFACE

- Why This Book Exists
- The Argument
- The Stakes
- Why a New Constitutional Framework Is Needed
- A Note on Authorship

PART I — THE STRUCTURAL DIAGNOSIS

Why Sovereignty No Longer Compiles: Governance after AI

Chapter 1 — The World States Cannot See

- I. The Illusion of Control
- II. The Three Dimensions of the Visibility Gap
 - A. Technical Opacity — The Black Box Problem
 - B. Jurisdictional Fragmentation — The G2 Bifurcation and G3 Ecosystems
 - C. Temporal Acceleration — The Speed Chasm
- III. Why the Gap Is Built In, Not a Passing Phase
 - A. The Divergent Clocks of AI and Authority
 - B. Why Expertise Concentrates at the Frontier
 - C. The Geometry of Authority versus the Geometry of AI
- IV. Consequences of Governing Blind
 - A. Regulatory Drift — Rules Without Reach
 - B. Perverse Incentives — The Performance of Compliance
 - C. Legitimacy Decay — Authority Without Capability
 - D. The Governance Vacuum — When Authority Has No Operator
- V. Who Does See — The Visibility Monopoly
- VI. The Road Ahead — What the Visibility Monopoly Makes Possible

Chapter 2 — The Rule-Execution Gap

Why Governance Fails at Implementation, Not Legislation

- I. The Void: Defining the Rule-Execution Gap
 - A. The Fairness Trap: When Math Defies Doctrine
 - B. The Transparency Paradox: The Fidelity Trade-Off
 - C. Why the Rule-Execution Gap Persists
- II. The Forces Widening the Gap
 - A. The Epistemic Asymmetry — The Tacit Knowledge Problem
 - B. Temporal Acceleration — The Obsolescence Trap
 - C. Jurisdictional Fragmentation — The Compliance Gridlock
- III. Why Traditional Solutions Cannot Close the Gap
 - A. More Resources — The Capacity Fallacy

- B. Faster Processes — The Speed Fallacy
 - C. International Harmonisation — The Consensus Fallacy
 - D. More Disclosure — The Transparency Fallacy
 - IV. Professional Services as Execution Infrastructure
 - A. Cross-Jurisdictional Observation — The Visibility Advantage
 - B. The Mechanics of Translation — Writing the Operational Code
 - C. Synthesis Authority — Reconciling Incompatible Sovereigns
 - D. The Inversion of Governance
 - V. Blended Templates as Operational Legislation
 - A. The Algorithm of Rule Selection
 - B. The Maximal Compliance Strategy — The Ratchet Effect
 - C. Conflict Resolution — Structural Workarounds
 - D. Template Propagation — The Governance Contagion
 - VI. The Translation Layer — A Governance Class in Formation
 - VII. Conclusion — Where Law Fails and Governance Begins
-

PART II — THE GRAMMAR OF GOVERNANCE

The Dual Languages of Governance

Section I — Law as the Language of Power

- The First Function: Naming What Exists
- The Industrial Logic versus the AI Universe
- The Categorical Gap in Practice
- The Tripartite Failure and the Broken Translator
 - A. How Law Decides What Is AI
 - B. The Five Functions of AI Legal Language
 - C. Limits of the Language of Law
 - D. The Translation Imperative
 - E. The Adaptive Constitution: When Private Law Outpaces Public Law
 - F. The Constitutional Consequence

Section II — Accounting as the Language of Measurement

- The Industrial Logic and the AI Mismatch
- The Rise of the Value Intermediaries
 - A. The Mechanics of Recognition
 - B. The Recognition Selection
 - C. Accounting Authors Economic Reality
 - D. The Measurement Gap
 - E. Value Creation through Translation
 - F. The Invisible Charter
 - G. The Delegated Interpreter

PART III — THE ARCHITECTURE OF THE GLOBAL AI STACK

- I. From Architecture to Implementation
 - II. The AI Governance Machine
 - III. The Arrival of Agents
 - IV. The Geopolitics of Vertical Integration
 - V. The Vertical Geometry of Global Influence
 - VI. The G3 Fracture
 - VII. The Offramp from Dependency
-

PART IV — THE GOVERNANCE ENGINE

Design, Propagation, and the Incorporation Heuristic

- I. Why the World Needs a Place Where Systems Can Talk
- II. The Architecture of the Translation Layer
- III. The Valuation Translator — How HKFRS Bridges GAAP and CAS
- IV. The Legal Translator — Where Rights and Obligations Compile Across the G2 Divide
- V. The Physical Translator — Hong Kong as Data Port in the AI Era
- VI. The Systemic Risk of Losing the Interface
- VII. The Sovereignty Paradox: Maps versus Networks
- VIII. The Incorporation Heuristic
 - A. Visibility: The Detection Test
 - B. Workability: The Implementation Test
 - C. Necessity: The Market Gravity Test
 - D. The Result: The Framework Sort
 - E. Raising Your Score — Strategies for Increasing Incorporation Weight
- IX. The Power Hierarchy — Root Access, Admin Access, and User Access
 - A. The Three Global Centres (Root Access)
 - B. The Power Users (Admin Access)
 - C. The Permission Matrix — A Hierarchy That Is Real, but Not Closed
- X. Gateway Rules — How Global Norms Move Across the World
 - A. Why Gateway Rules Spread — The Economics of Convergent Design
 - B. The Golden Image — How Templates Spread the Rule
 - C. Three Ways Rules Travel
 - D. From Description to Design
 - E. Strategies for Non-G3 Countries — Designing for Incorporation
- XI. Conclusion — The Invisible Constitution

PART V — THE MATERIAL CONSTITUTION OF THE AI ERA

Compute, Capacity, and the Constitutional Limits of Governance

- I. The New Assets: Compute, Data, Models
 - A. Compute — The Engine
 - B. Data — The Fuel
 - C. Models — The Territory
 - D. The Triad in Practice — A Single Scenario
 - E. Transition — From Assets to Constraints
 - II. The Legitimacy Crisis — The Widening Gap Between Statute and AI Governance
 - A. The Shift of Authority — From Institutional Control to AI-Driven Execution
 - B. The Five Non-Substitutable Constraints
 - C. The Gradients of Operational Agency
 - D. Legitimacy Debt — The Silent Failure Mode
 - E. What Legitimacy Collapse Looks Like
 - F. The New Test of Legitimacy — Consent versus Execution
 - G. The Five Design Conditions for AI Legitimacy
 - III. Code After: The Operating Manual for the AI Era
 - A. The Sovereign API
 - B. Parametric Democracy — Governing at Machine Speed
 - C. The Protocol of Federation — Interoperability Without Harmonisation
 - D. The Four-Tier Constitutional Design and the Role of Licensed Operators
 - E. Calibration Events — Democratic Revision at Machine Speed
 - F. The Hard Problems — Inherent Tensions That Remain
 - G. The Interdisciplinary Necessity
 - H. Closing the Loop — The Visibility Gap Revisited
-

PREFACE

Why This Book Exists

For roughly three centuries in the Western legal-bureaucratic tradition, sovereignty rested on a single architectural assumption: the State occupied the highest vantage point. From that altitude, it could observe territory, classify activity, and issue commands that flowed downward through the transparent medium of law. Visibility and authority were inseparable; to see was to govern.

That architecture no longer holds.

We have crossed the threshold from the Information Age — when computers processed data for humans to decide — into the AI era, where software acts on its own. AI systems now classify, recommend, negotiate, allocate, and enforce at scales and speeds that exceed the observational capacity of any sovereign. The State has not lost its legal authority; it has lost its physics. Its inherited tools now operate under physical and informational limits they were never designed to meet.

This book is a systematic diagnosis of that condition — and a structural proposal for what may succeed it.

It is also a versioned release. Because AI capabilities and regulatory frameworks evolve faster than traditional publishing cycles, this manuscript is issued as *Code After: v0.9 (Open Access Edition)* — an interim release intended to enter scholarly and policy debate while the full edition is prepared.

One scope clarification is necessary. This analysis focuses on the oversight of centralised, cross-border AI systems — the computational platforms operated by frontier AI labs, hyperscalers, major technology firms, and multinational organisations through which regulatory obligations are translated and enforced. Open-weight models introduce a parallel dynamic, extending AI activity beyond centralised compliance perimeters. Their regulatory treatment will be addressed in the full edition.

The Argument

The argument unfolds in five movements.

Part I establishes the structural problem. It introduces the Visibility Gap and the Rule-Execution Gap, showing why inherited legal and oversight frameworks cannot regulate AI systems whose operations are opaque, high-velocity, and increasingly agentic.

Part II examines law and accounting as the dual constitutive languages through which economic reality is defined, measured, and governed. AI systems — like corporations and financial instruments before them — become legible to systems of control only when translated into these institutional grammars. The analysis identifies the Categorical Gap and the Measurement Gap, develops the Invisible Charter as the fused evaluative grammar through which AI systems become economically real, and shows how interpretive authority has migrated to professional intermediaries who now function as the Translation Layer.

Part III traces the migration of governance from horizontal, border-based frameworks to vertical, stack-based architectures. It introduces the AI Governance Stack, the Agentic Shift,

and the emergence of vertically integrated AI ecosystems, situating the United States (U.S.), China, and the European Union (EU) as three increasingly distinct institutional regimes whose divergence makes cross-border translation structurally necessary.

Part IV introduces the Translation Layer as operational machinery — beginning with Hong Kong as a reference implementation of the institutional infrastructure through which incompatible governance stacks currently communicate. From that structural analysis, it develops the Incorporation Heuristic as the logic shaping which rules travel, stabilise, or fail within globally networked markets, and presents the Sovereignty Paradox, the Gateway Rules framework, and the AI Power Hierarchy as structural consequences of translation under conditions of opacity and national competition.

Part V begins with the physical foundations of AI — the material constraints imposed by compute concentration, energy demand, and non-territorial strategic assets that shape national agency. It diagnoses the Legitimacy Crisis: the democratic deficit in the governance arrangements now operating. It concludes with the Third Architecture — the Sovereign API, Parametric Democracy, and the Protocol of Federation — as institutional designs for governing what follows.

The manuscript advances three contributions. First, it reframes AI governance by shifting the analytical centre from model behaviour to the institutional systems that render AI visible, legible, and governable. Rather than treating AI as a technical object to be regulated, the analysis shows how law, accounting, and administrative practice jointly construct the economic reality in which AI operates. Second, it advances accounting theory by demonstrating that the core operations of recognition, measurement, classification, disclosure, and consolidation no longer align when applied to probabilistic, networked, and emergent systems, producing a structural Measurement Gap that limits states' capacity to perceive, value, and govern AI-driven economic activity. Third, it develops a constitutional analysis of how governance functions when intelligence becomes distributed. The analysis formalises which rules propagate across borders and why, maps the resulting hierarchy of operational influence, and proposes institutional mechanisms through which democratic accountability can be extended to systems that currently operate beyond its reach.

The Stakes

AI is not unmanaged. It is shaped by default institutions, through mechanisms no one designed directly and without meaningful democratic consent. Professional services firms — the global law firms and the Big Four accounting networks — have quietly become the operational stewards of AI governance, not because they sought this role but because the structure of the AI era placed it on their desks. They convert legal requirements into operational systems and produce the templates that function as AI's invisible legislation.

This arrangement sustains a functioning global AI economy across a fractured world, but it carries a democratic deficit. No one voted for the Big Four, and no constitution imagines global law firms writing the operational rules for machine intelligence. This book makes that machinery visible, clarifies the available design choices, and proposes the institutional architecture that may succeed it.

Why a New Constitutional Framework Is Needed

AI governance does not fail for lack of principles. Ethical frameworks proliferate, national strategies multiply, and regulatory proposals accumulate. The constraint is more fundamental: rules do not execute themselves.

Between a legislative instrument and AI behaviour lies a structural gap. Law is abstract, territorial, and deliberative; computation is specific, distributed, and instantaneous. Effective authority depends on the machinery beneath it — the operational procedures, compliance architectures, and technical controls through which rules become executable inside AI ecosystems. This requires design, not aspiration. It requires translators capable of conversion and a constitutional framework built for the operating realities of the AI era rather than the assumptions of earlier regimes.

A Note on Authorship

This book is written by a single author, but it did not emerge from a solitary path. It began as a shared inquiry with my daughters, Natarsha and Jacqueline, as they entered their studies in law and accounting at Berkeley Law and NYU Stern. Their questions, research contributions, and critical engagement — spanning constitutional structure, institutional design, accounting standards, and the mechanics of global practice — shaped the conceptual terrain of this project and helped reveal where doctrine no longer maps cleanly onto operational reality.

This book reflects that shared journey. It is not a co-authored work, but it is collaborative in origin, grounded in the belief that scholarship is a long arc built across generations. The conceptual frameworks developed here are intended as a foundation for independent scholarship — and my hope is that Natarsha and Jacqueline will extend, challenge, and build upon them in their own right.

PART I — THE STRUCTURAL DIAGNOSIS

Why Sovereignty No Longer Compiles: Governance after AI

For much of the modern era, states governed from a position of informational advantage. Authorities assumed they could see enough of the world to classify it, regulate it, and enforce their decisions through law. That assumption no longer holds.¹

The technologies that once waited for human direction now act at machine speed. AI systems classify, allocate, negotiate, and enforce at scales that exceed the observational reach of public institutions. The State's authority remains intact on paper, but its ability to perceive and operationalise that control has eroded.

The crisis of AI governance is not a failure of political will but of institutional capacity: the State is attempting to oversee an ecosystem that moves faster, operates deeper, and obscures more than the machinery built to regulate an earlier technological order. Two fractures define this condition.

The first is the Visibility Gap — a structural asymmetry in which regulators cannot observe the systems they claim to govern. The operative logic of the AI economy is written in dynamic code, legible primarily to the private actors who build and maintain it. Opacity is not incidental; it is a predictable consequence of how these systems are trained, engineered, and deployed.

The second is the Rule-Execution Gap — the failure that emerges when legal mandates cannot be rendered into operational control. Principles such as fairness, safety, or non-discrimination require translation into technical specifications, and that translation occurs outside the State. Authority falters not at the level of intent but at the level of implementation.²

These are not temporary inefficiencies. They are symptoms of an operating system that no longer matches the world it was built to govern. Before the Code After can be built, the reasons the code before can no longer run must be understood.

¹ Scott, *Seeing Like a State* (1998); Krasner, *Sovereignty: Organized Hypocrisy* (1999).

² The conceptual framework developed in this manuscript — including the diagnostic constructs (Visibility Gap, Rule-Execution Gap, Categorical Gap, Measurement Gap, Visibility Monopoly), the analytical architecture (Dual-Compiler Thesis, Incorporation Heuristic, Topology of Power, Translation Layer, Invisible Charter, Invisible Constitution, Gateway Rules, Golden Image), and the prescriptive design (Third Architecture, Sovereign API, Audit Tier, Protocol of Federation, Parametric Democracy, Licensed Operator) — represents original contributions introduced here. Terms are defined at first use throughout the text, and a Glossary is provided at the end of the manuscript.

Chapter 1: The World States Cannot See

On a recent Tuesday morning, a credit-scoring AI operating across more than fifty countries made millions of lending decisions. It approved mortgages in Manchester, denied car loans in Mumbai, and recalibrated credit limits in Mexico City. The example is illustrative, but it reflects a real pattern: AI-driven decisions now occur continuously, across borders, and at scales that no single regulator can directly observe³.

Although these decisions were globally distributed, the model's training data, cloud infrastructure, corporate parent, and AI supply chains all intersected with three of the most consequential AI ecosystems. The U.S., China, and the (EU) shape the technical, economic, and regulatory conditions under which such systems operate. They receive particular attention in this book for that reason.

The pattern extends beyond financial services. Global supply chains are increasingly orchestrated by optimisation models that draw on port data, customs risk signals, congestion forecasts, and demand projections from multiple countries. No single state typically has access to the full dataset, the model design, or the decision logic that shapes routing, prioritisation, and delay. Each country therefore sees only a fragment of a system whose behaviour emerges from interactions that span borders and elude any single regulator's full view.

In both cases, the same structural condition holds. No public authority sees the system as a unified whole. American regulators observe only the domestic segments of data flows. European supervisors engage with documentation prepared under the General Data Protection Regulation (GDPR), the EU's comprehensive privacy law.⁴ Chinese officials monitor activity on servers within their borders. Each interacts with a fragment of the system; none has visibility into its full operational footprint.

This matters because sight is the precondition of authority — and control is weakened without it.

The global AI economy now operates through infrastructures and decision-making pipelines that frequently outpace the observational reach of any single state. The challenge is not merely regulatory lag — the familiar refrain that “law is slow and technology is fast” — but a deeper structural mismatch: regulators are attempting to oversee an ecosystem that is often opaque by design.

States retain hard sovereignty — the coercive authority to shut down data centers, restrict compute, or block market access. What they have lost is execution optionality: the ability to shape model behaviour directly through institutional command. Recognising this asymmetry clarifies why so much contemporary AI oversight feels misaligned with the world it seeks to regulate. Legislation arrives beside the point, harmonisation efforts struggle against operational reality, and the centre of gravity has shifted from rulemaking to translation.

³ Financial Stability Board, *Artificial Intelligence and Machine Learning in Financial Services* (2017); Bank for International Settlements, *Machine Learning in Credit Risk* (2023).

⁴ Regulation (EU) 2016/679 (General Data Protection Regulation).

Translation here carries a specific meaning: not a linguistic activity, but a technical one performed by actors positioned to observe what states cannot. Professional services firms perform this function using two of the principal operational languages that travel across time and borders: law, the language of permission, and accounting, the language of measurement. Together, these languages allow global AI to run, regardless of the political and institutional values that surround it.

I. The Illusion of Control

As AI spreads across the globe, the central question is no longer what rules legislators write, but how oversight actually works — and where it truly resides.

In the U.S., authority fragments across multiple agencies — the Federal Trade Commission (FTC), the Securities and Exchange Commission (SEC), and the National Institute of Standards and Technology (NIST) — each gripping a different piece of the elephant.

In Brussels, the EU Artificial Intelligence Act (EU AI Act) constructs a cathedral of risk frameworks: an elaborate edifice of classifications, obligations, and oversight mechanisms.⁵

In Beijing, the Cyberspace Administration of China (CAC) enforces a perimeter model of AI oversight, closely monitoring data flows, content, and the technical backbones that support computation within China.

Across the world, the machinery of regulation appears formidable. Tens of billions of dollars fund regulatory capacity across finance, health, environment, and technology. Hundreds of thousands of officials — spanning national regulators, subnational agencies, supranational bodies, and international organisations — promulgate rules, conduct inspections, and enforce compliance. Entire industries support this apparatus: regulatory affairs outsourcing, compliance services, and RegTech platforms that automate monitoring and reporting. Academic literature proliferates. Standards bodies issue frameworks. By every visible metric, the regulatory landscape appears fully engaged.

Yet the entire apparatus rests on an assumption that no longer holds. In the pre-AI world, economic activity was comparatively legible, slower, and territorially bounded. Regulators could see factories, inspect ledgers, audit processes, and compel disclosures. Observation was imperfect, but it was possible. The State's sensory apparatus — reports, filings, inspections, subpoenas — was built for a world in which the relevant activity remained within its field of view.

In the AI era, that visibility is under strain. What regulators once monitored directly now unfolds across distributed, dynamic, and partially autonomous environments. The operative logic is encoded in proprietary models, executed across borders, and updated continuously. The relevant activity no longer sits in places the State can easily observe; it lives in technical layers that are often inaccessible to public authority.

⁵ Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act).

This reflects a deeper systems-level mismatch: effective oversight of AI depends on forms of institutional sight that no single state fully possesses. It spans technical observability — access to model design, training data, emergent behaviour, failure modes, and security vulnerabilities — and extends to economic sight into valuation, capital flows, incentives, and market concentration. It culminates in classificatory legibility, the capacity to render system behaviour into categories that law and accounting can recognise, measure, and act upon.

No country holds these capacities in complete form. The reasons are structural. They arise from how sovereignty, technology, and economic organisation now interact in an integrated AI ecosystem. States are attempting to steer activity whose technical behaviour is controlled by private actors, whose economic dimensions are distributed across cross-border networks, and whose classificatory forms do not align with the representational grammar of territorial authority.

What emerges is the appearance of regulatory control: the belief that governments can direct what they cannot fully see.

II. The Three Dimensions of the Visibility Gap

The Visibility Gap has three interlocking dimensions, each representing a distinct form of informational deficit. Together, they generate systematic blindness precisely at the level where oversight decisions must be made.

A. Technical Opacity — The Black Box Problem

AI models are often opaque both by nature and by design. Frontier-scale deep learning systems contain billions or trillions of parameters whose interactions generate behaviours that cannot be fully predicted, reconstructed, or explained — not even by the teams that build them. The opacity is intrinsic to high-dimensional statistical algorithms and cannot yet be fully engineered away.

Consider the credit-scoring model introduced at the start of this chapter. Its training data, model design, and fine-tuning process are typically proprietary, reflecting the trade-secret norms of commercial credit modelling. Its decision logic emerges from statistical correlations distributed across millions of parameters — patterns that do not naturally reduce to a stable, global set of human-interpretable rules.

Even if regulators obtained full access to the underlying information — which they typically cannot — they would still confront the interpretability problem: knowing what a model outputs does not, by itself, explain why it outputs it. Sight is not the same as comprehension.

Technical opacity is compounded by a form of systemic adverse selection. The observational capacity required to understand state-of-the-art AI models has been effectively privatised by the market price of talent. The expertise necessary to interrogate model behaviour — machine-learning research, systems engineering, security analysis — has migrated overwhelmingly into private firms.

This pattern reflects a broader labour-market dynamic: advanced AI capabilities concentrate around a small number of firms and countries. The migration is not accidental. It reflects the

logic of modern labour markets: technical competence flows toward the actors best positioned to monetise it. Governments are left competing in a labour market they cannot afford to win.

At this point, a familiar objection arises: opacity might seem like a resource problem. If governments simply hired more engineers, built more supercomputers, and expanded regulatory budgets, they might believe they could eventually peer inside the box.

This assumption underestimates the structural character of the problem. Even the U.S. and China — home to the world’s deepest reservoirs of technical talent, compute capacity, and frontier model research — face severe limits in seeing, explaining, and auditing the platforms they themselves produce, let alone those produced by each other. Their laboratories operate near the limits of scale. Their firms train the largest models. Their engineers build the most complex pipelines.

Yet the opacity of AI models exceeds even their internal capacity for explanation. The creators cannot fully interpret their own creations. The black box is not hiding from regulators. It is hiding from everyone.

This is why model opacity cannot be solved through additional disclosure, more audits, or stronger compliance regimes. These tools assume that observation is primarily a function of effort. It is not. Observation is a function of technical design.

B. Jurisdictional Fragmentation — The G2 Bifurcation and G3 Ecosystems

Even if technical opacity could be overcome — a counterfactual — the second dimension of the Visibility Gap would remain: no state can see all components of the AI infrastructure at once.

Global AI is not a single universe. It is a multipolar but highly uneven landscape built around a small number of countries capable of operating across the full span of the AI pipeline. A fully integrated state controls the critical components: advanced chip design and fabrication, hyperscale cloud infrastructure, large-model training pipelines, data reservoirs, global deployment channels, deep talent pools, and domestic capital markets able to finance and scale AI.

Full capability does not mean self-sufficiency in every component; it means strategic control over the critical resources required to build, train, and deploy large-scale AI systems without dependence on external political decisions.

At present, the clearest states meeting this threshold are the U.S. and China. Others — most notably the EU — wield significant regulatory influence but do not yet control the underlying technical base. Europe remains dependent on foreign cloud providers and semiconductor supply chains, even as policy debates now centre on building a “EuroStack” or “public AI stack” to reduce this structural reliance. These initiatives signal a strategic ambition to expand Europe’s technical footing, though they remain early-stage rather than fully realised.

Most other significant AI economies — including the United Kingdom (UK), Japan, India, Korea, and Brazil — operate within what this manuscript terms stack dependence: they

regulate models built on chips they did not fabricate, trained on compute clusters they do not own, and deployed on platforms whose core designs are external. Their authority rests primarily at the regulatory, compliance, and market interface rather than in shaping the technical base itself.

This distinction defines the G2: an analytical category referring to state-level AI ecosystems that meet the threshold of vertically integrated capability across hardware, compute, model development, talent concentration, and capital scaling. In the current global landscape, the U.S. and China are the clearest G2 states. The term is descriptive rather than geopolitical, identifying countries capable of sustaining the full technological and financial span required to develop and deploy advanced AI at global scale.

A complementary category, the G3, refers to states whose regulatory frameworks propagate globally through market size, legal-harmonisation capacity, and institutional legitimacy, even without full integration across the AI stack. The EU is the clearest G3 actor, exercising regulatory gravity that shapes global norms despite the absence of full technology control. These categories describe current structural positions rather than permanent assignments; states can move through the hierarchy as their technical and regulatory capabilities evolve.⁶

This bifurcation produces a distinct form of informational failure: fragmented oversight. While the G2 possess vertical visibility over their own ecosystems, they do not possess full insight into each other's — and when the frame widens to include every other major AI power, no state in the system commands horizontal visibility across the horizon where all stacks inevitably interact.

Even with the world's most advanced AI ecosystem, the U.S. regulator can scrutinise segments of the codebase and portions of the training pipeline, but it still sees only narrow, sector-specific slices of how these systems circulate globally. Its regulatory sightline is uneven across the stack: strong over model development and platform-level diffusion, weaker over downstream uses that fall outside its legal reach.

The Chinese regulator exercises significant oversight over infrastructure — chips, supply chains, and datacenter physics — but has more limited sightlines into downstream monetisation, offshore diffusion, and cross-border model behaviour. China enjoys full vertical integration, yet its regulatory vantage point is anchored in hardware, data, and domestic compute rather than in the global pathways through which its models circulate.

The EU reviews compliance filings and assurance materials, but typically lacks systematic sightlines into the hardware and compute physics that sit outside its legal reach. Its regulatory grammar travels globally, yet its oversight remains concentrated in documentation, assurance processes, and market-access controls rather than the technical base.

Stack-dependent regulators — in the UK, Japan, India, Brazil, and elsewhere — primarily encounter the outputs of deployed models, with limited sightlines into upstream training data, model-development pipelines, or underlying model behaviour. Their vantage point sits downstream of the technical base.

⁶ The G2 and G3 are analytical categories introduced in this manuscript to describe structural positions within the global AI ecosystem. They are descriptive rather than geopolitical and should not be confused with prior uses of "G2" in international relations, military, or other discourse.

Each state sees a different slice of the landscape. None sees the system as a whole.

More importantly, no state can track how changes in one domain ripple through the others. A shift in Chinese data-localisation rules can alter American model-retraining cycles; a new European privacy standard can redirect global capital flows; a diffusion restriction in India can cascade into model behaviour. Each example traces a different propagation path — technical, financial, behavioural — underscoring why fragmented sightlines are not merely inconvenient but systemically disabling.

This is not a coordination problem. It is not a failure of diplomacy, harmonisation, or intergovernmental process. It is a fundamental mismatch between territorial authority and transnational AI — a world in which the borders that define political power no longer match the operational reality of the AI economy.

C. Temporal Acceleration — The Speed Chasm

Even if technical opacity were solved and territorial fragmentation aligned, a third dimension would persist: the frontier moves faster than oversight can track, interpret, or respond.

AI systems operate on timescales that exceed the State's capacity to observe or intervene. Decisions that once unfolded over hours or days now occur in milliseconds. A single credit-risk model can reassess large numbers of lending positions in near real time — a tempo no public institution can match. The asymmetry is inherent: the systems being governed move faster than the machinery built to govern them.

Oversight in the AI era is defined by asymmetric iteration rates. AI development and deployment accelerate through scaling, automation, and rapid feedback loops, while institutional response moves through consultation, deliberation, and negotiation. The gap between these tempos widens with every cycle, creating a governance environment in which acceleration itself becomes a source of institutional failure.

Consider the timing mismatch in the EU. The EU AI Act — the world's most comprehensive attempt to regulate AI — was proposed in April 2021. Midway through its negotiation, the release of ChatGPT (built on GPT-3.5) in late 2022 forced a rapid, midstream revision. The resulting amendment attempted to retrofit the Act to capabilities the original draft could not have anticipated. Yet even this intervention could not accelerate the legislative machinery: the Act entered into force only in mid-2024, with full implementation stretching into 2026. In AI time, a two-year lag is not a delay; it is a systemic failure.⁷

GPT-3.5 was not simply a technical milestone; it was the public inflection point — the first AI model to reach global consumers at scale, triggering what became a planetary reckoning with the arrival of machine intelligence. The dramatic demonstration of capability — amplified by warnings from senior figures in the field, including Geoffrey Hinton —

⁷ Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act). The Commission proposal was published on 21 April 2021 as COM(2021) 206 final. The Act entered into force on 1 August 2024.

transformed an obscure technical development into a mainstream technological rupture. It happened overnight.⁸

While the EU deliberated, AI had already cycled through successive capability regimes: the transition of GPT-3 from experimental tool to production instrument, the emergence of GPT-4 as a reasoning engine, and the rise of agents able to execute multi-step workflows. By the time the Act reached political consensus, the frontier had already shifted again.

Democratic deliberation cannot match the speed of AI because legitimacy takes time. Consultation, debate, and review are essential safeguards, not inefficiencies. These processes cannot be compressed without eroding the very authority they are meant to confer. The consequence is a systemic mismatch: public decision-making moves in procedural time, while AI advances in exponential time.

Authoritarian regimes do not escape the speed chasm either. Centralised authority can compress deliberation, but it cannot compress comprehension. The decision latency of any bureaucracy — democratic or otherwise — is measured in weeks or months. The AI cycle is measured in days; the inference cycle of an agent is measured in milliseconds.

Even the G2 cannot regulate at the speed they innovate. Their laboratories iterate on weekly cycles; their legislatures and ministries move on annual ones. The widening tempo gap is structural. The State is always responding to the last generation of models, while the systems operating within its borders run on a clock the law has not yet reached.

III. Why the Gap Is Built In, Not a Passing Phase

It is tempting to believe the Visibility Gap is temporary — a friction that will fade once regulators hire more experts, coordinate internationally, or adopt more agile processes. This view misdiagnoses the problem. The gap is not a transitional failure of capacity. It is a system-level consequence of how AI advances, how markets allocate expertise, and how state authority is organised. These forces do not weaken over time; AI compounds them.

Three forces make the Visibility Gap durably persistent.

A. The Divergent Clocks of AI and Authority

The core asymmetry is straightforward: AI compounds; public oversight resets. AI and state response operate on fundamentally different clocks. AI advances through recursive optimisation loops: scaling laws, fine-tuning cycles, reinforcement from human and synthetic feedback, and continuous redeployment. Each iteration builds directly on the last, compressing learning into ever shorter cycles. AI improves by running.

Public authority iterates differently. It requires processes that are intentionally slow because legitimacy depends on them. These cycles do not accelerate with experience; they reopen with every new issue, every new stakeholder, every new political change.

⁸ OpenAI released ChatGPT on 30 November 2022. For Geoffrey Hinton’s public statements on AI risk, see interviews with the New York Times (1 May 2023) and CBS (25 March 2023).

The consequence is not simply that AI moves faster than the State. AI compounds while public decision-making renews process. One system accumulates capability; the other re-initiates procedure. These clocks diverge by design, not by accident. No institutional reform bridges the distance between compounding iteration and procedural renewal.

B. Why Expertise Concentrates at the Frontier

In late 2023, the U.S. federal government opened senior machine-learning research roles for its newly created AI Safety Institute. The salary band — roughly \$180,000 — was competitive by civil-service standards. In the same period, Anthropic hired a comparable researcher at a total compensation package exceeding \$1.2 million, with access to frontier-scale compute clusters and a team of more than one hundred PhDs working on large-scale optimisation problems. The federal posting remained open for months. Anthropic filled its role in days.⁹

This divergence is not a budgetary anomaly. It is the predictable outcome of market incentives. Frontier AI requires rare combinations of expertise: machine-learning research, distributed systems engineering, security analysis, and large-scale optimisation. As models scale, the marginal value of technical insight increases, and the market bids that insight upward. The private sector can offer compensation, equity, compute access, and research environments that no public institution can match.

The consequence is systemic adverse selection. The people most capable of understanding frontier AI are inevitably drawn to the labs that create it. Governments cannot reverse this dynamic. They cannot replicate frontier-scale compute environments, they cannot match equity-based compensation, and they cannot offer the density of technical peers that AI labs attract. Even the G2 — with the deepest talent pools on Earth — face the same challenge. Their national laboratories and ministries do not house the teams that train trillion-parameter models; their regulators do not employ the engineers who design the systems they are tasked with overseeing.

C. The Geometry of Authority versus the Geometry of AI

No country currently sees full AI stacks. The remaining question is whether this discontinuity is contingent — solvable through coordination — or inherent. It is the latter, for a reason not yet identified in the preceding sections: the geometric mismatch between territorial authority and global compute placement — a mismatch that persists regardless of how political boundaries are redrawn.

States govern through territorial reach: borders define what can be observed, measured, and constrained. AI does not operate within these boundaries. It runs across cloud regions, data centres, and supply chains organised around efficiency, latency, and capital flows — a

⁹ Based on publicly reported federal salary bands and industry compensation data. Figures are approximate and illustrative.

topology that cannot be reshaped to match territorial lines without dismantling the global AI economy in the process.

This is what makes the mismatch geometric rather than diplomatic — it remains invariant under deformation. Borders can shift, territories can merge, treaties can proliferate — yet the compute remains borderless. No reconfiguration of political boundaries produces a world in which frontier AI falls within a single country’s line of sight.

Even the G3 cannot escape these limits: the U.S. shapes chip design and controls access to leading-edge fabrication but does not operate the fabs themselves; China controls domestic deployment and data flows but not the global capital that finances AI; the EU defines regulatory power but not the hardware or models on which that power must act. Each commands a slice of the system, but none commands the whole. Territorial authority is bounded; AI is borderless. Their geometries do not meet, and no degree of coordination can force them into alignment.

IV. Consequences of Governing Blind

When state authority loses traction over the AI it is meant to oversee, the failure is not procedural but constitutional. The machinery of public administration continues to operate — legislatures legislate, agencies issue rules, courts adjudicate — but the link between institutional action and AI behaviour weakens. Over time, this loss of coupling produces three internal distortions that reshape the State from within.

These are not enforcement failures or coordination gaps. They arise when institutions attempt to steer AI whose internal dynamics they cannot interrogate, model, or predict. A political centre that cannot establish causal influence over the platforms under its charge does not merely exercise authority poorly; it rules symbolically.

A. Regulatory Drift — Rules Without Reach

When a public authority cannot observe the AI models it seeks to oversee, regulation becomes an exercise in projection. Institutions legislate toward an imagined AI rather than the one operating in practice. The gap widens between the formal requirements written into statutes and the real behaviour of deployed systems.

This gap produces a distinctive failure mode: rules that exist on paper but do not bind in practice. Agencies issue frameworks, guidance, and compliance expectations, yet these instruments rarely penetrate the decision-making surface where outcomes are actually produced. Firms respond by constructing compliance covers — policies, checklists, model cards — that satisfy external expectations without altering internal logic. Oversight becomes more visible even as control becomes less real.

Over time, this dynamic becomes self-reinforcing. As visibility declines, regulators compensate with additional rules; as rules proliferate, firms invest more in producing documentation that meets procedural expectations than in aligning system behaviour with regulatory intent. The legal process becomes increasingly expressive — signalling priorities, values, and aspirations — while losing traction over the systems that shape outcomes.

This is regulatory drift: a world in which legal authority persists, but its operational reach diminishes. The public sector continues to issue commands, but the systems under its charge evolve along trajectories the law can no longer detect, interpret, or redirect.

B. Perverse Incentives — The Performance of Compliance

When the field of view narrows, incentives reconfigure. Companies shift from optimising for system performance to optimising for the narrow band of signals regulators can actually perceive. In that environment, documentation becomes the dominant channel of interaction between firms and the State.

The consequence is a distinctive institutional failure: the substitution of representations for assurance. In the absence of real observability, firms generate documents that are easy to audit but only loosely connected to the systems they describe. Model cards summarise design intent rather than trace operational behaviour. Risk matrices classify hypothetical hazards rather than failures observed in production. Compliance templates proliferate that satisfy audit logic without touching the decision-making mechanisms that drive outcomes.

These artefacts flourish precisely because they are legible. They translate dynamic, opaque systems into static representations that fit the bureaucratic machinery of review. But legibility is not control. As the gap between documentation and system behaviour widens, compliance becomes a performance staged for oversight bodies that lack the means to interrogate the underlying reality.

Over time, this dynamic produces a parallel oversight process — a meticulously maintained ecosystem of policies, templates, and attestations that evolves independently of the models it purports to regulate. Internal incentives shift accordingly: teams invest more in producing documentation that satisfies external expectations than in building mechanisms that would meaningfully constrain model behaviour. The compliance apparatus becomes increasingly elaborate, increasingly sophisticated, and increasingly detached.

C. Legitimacy Decay — Authority Without Capability

When public bodies lose traction over the systems they are meant to oversee, the resulting damage is not administrative. It is constitutional. The modern regulatory state rests on a premise that it can understand the domains under its charge well enough to act meaningfully within them. When that premise falters, legitimacy begins to erode long before authority formally collapses.

The Visibility Gap accelerates this erosion. As AI ecosystems evolve beyond the State's capacity to observe or interpret them, confidence in governmental competence begins to fracture. Citizens encounter rules that cannot be enforced. Firms receive directives that cannot be operationalised. Regulators themselves sense the widening gulf between the authority they hold in statute and the capability they possess in practice. The belief that the State can meaningfully steer these systems weakens across all constituencies simultaneously.

This produces a quiet legitimacy crisis — one that accumulates through small, unremarkable failures rather than dramatic breakdowns. Investigations stall for lack of technical insight. Enforcement actions target the visible rather than the consequential. Regulatory processes

generate text but not traction. Each instance is minor in isolation; together, they signal a public authority increasingly unable to shape the systems that shape society.

Legitimacy decays because it is recursive. State actors derive authority not only from formal mandates but from the shared expectation that they can act effectively. When that expectation erodes, compliance becomes discretionary, deference becomes conditional, and the State's coordinating power weakens. Authority persists in form, but its functional core begins to hollow out.

The legitimacy decay that follows — a government retaining the symbols of authority while losing the capability that once justified them — is incremental, cumulative, and largely invisible until it becomes structurally entrenched.

D. The Governance Vacuum — When Authority Has No Operator

The three failures described above — regulatory drift, the performance of compliance, and legitimacy decay — do not remain discrete. They interlock and reinforce one another. Over time, they collapse into a single condition: a regulatory vacuum, a structural space in which the State retains formal authority but loses the practical capability to exercise it.

A void of this kind is not an absence of order. It is an absence of public order. The global economy cannot function without risk assessment, standard-setting, assurance, and coordination. These are not optional features of complex systems; they are the infrastructure that makes complexity governable. When public authorities can no longer supply these public goods, the system does not stall. It reconfigures.

And in that reconfiguration, authority flows to the actors who possess what the State has lost: operational visibility into the behaviour of frontier AI systems. The void is not filled by insurgent bodies or novel political formations. It is filled by familiar intermediaries — consultancies, auditors, major law firms, and professional services networks — whose proximity to AI systems gives them the vantage point required to interpret, evaluate, and constrain behaviour in practice. (If platforms perform governance functions of their own, such as content moderation or safety gating, those functions operate in parallel rather than as substitutes for the State's regulatory role.)

These actors do not overthrow the State; they inherit the functions the State can no longer fulfil. They become the de facto interpreters of risk, the authors of compliance norms, the coordinators of cross-organisational behaviour. They supply the translation capacity that public authority now requires but can no longer produce.

Here the inquiry turns. The question is no longer whether governments can steer AI — their capability is structurally constrained under the conditions established in this chapter. The question is who fills the gap, and what it means for public authority when the practical work of oversight migrates to private hands.

Where the State cannot see, others already do.

V. Who Does See — The Visibility Monopoly

The erosion of state sightlines does not create an empty space. It shifts observational power toward organisations capable of operating simultaneously across the technical, legal and financial domains of the global AI economy. Not shadow governments. Not autonomous AI agents. Professional services firms.

Global law firms, the Big Four accounting networks, and elite strategy consultancies now hold a level of cross-jurisdictional, cross-stack visibility that public authorities increasingly struggle to match. They operate where governments cannot: across borders, across platforms, across industries, and across time. Their vantage is not positional. It is embedded in the role they occupy within the global economy.

They observe dimensions of the AI economy that rarely converge inside any public institution. They see how code is architected, fine-tuned, deployed, and modified in production, and how data and model outputs move across legal regimes with incompatible assumptions. They see how firms reconcile simultaneous obligations under divergent regulatory systems, and how capital markets interpret and price risks that regulators are still learning to define.

This visibility is not a byproduct of their work. It is the work. These firms do not simply interpret rules. They convert rules into operational reality. They do not only advise on compliance. They construct the controls that organisations rely on. They function as a Translation Layer between formal legal regimes and the technical and organisational systems through which AI is deployed.

Their influence is not supervisory but infrastructural. These firms shape outcomes through the artefacts they produce: the templates that become default contractual forms, the cross-jurisdictional mappings that harmonise divergent regulatory requirements, and the implementation playbooks that define what “compliance” means in practice. By standardising how obligations are interpreted and operationalised, they produce the operational definitions that most organisations follow — definitions that often travel farther and faster than the statutory requirements they were designed to implement.

Their position emerges from the system-level dynamics created by the Visibility Gap. When no single public body can observe the full system, the actors capable of seeing across its operational layers become the natural sites of coordination, interpretation, and implementation. They occupy the only vantage point from which fragmented legal requirements can be translated into coherent, executable systems — and in doing so, they become the practical governors of AI deployment.

Within these firms, practical knowledge about how AI oversight actually functions is accumulated, refined, and transmitted through the day-to-day work of audits, investigations, model evaluations, contract negotiations, and cross-border compliance builds. The Translation Layer is not merely where regulatory implementation is carried out; it is increasingly where the capability to govern AI is produced.

The Visibility Gap does not directly produce monopoly; it produces dependence. That dependence becomes a monopoly when it concentrates around actors whose position cannot be replicated or bypassed. A visibility monopoly, in the sense used throughout this manuscript, requires three structural conditions: privileged access to system-level

observability that no public authority can replicate; non-substitutability, meaning that the operational knowledge accumulated through cross-stack, cross-jurisdictional practice cannot be reconstructed by alternative observers without equivalent institutional reach; and functional dependence, in which regulators, markets, and firms rely on these actors not merely for advice but for the interpretive and implementation capacity on which compliance itself rests.

Hyperscalers satisfy the first condition through exclusive visibility into their own models — internal states, inference traces, system-level behaviour, training data provenance, and failure modes that no external actor can independently access. The Translation Layer satisfies the second and third: its cross-jurisdictional vantage cannot be substituted by any single public institution, and its interpretive authority makes regulators dependent on the operational knowledge it produces. Together, hyperscaler visibility and Translation Layer visibility constitute a structural position historically occupied by the State — one that now determines the conditions under which behaviour, value, and agency become legible. This is the Visibility Monopoly.

A practical illustration emerged during the early implementation of the EU AI Act. When the regulation entered into force in 2024, it did not self-execute. Within weeks, major accounting and advisory networks developed standardised risk-classification and conformity-assessment frameworks for multinational clients. By the time detailed regulatory guidance emerged months later, many firms were already operating under internally engineered compliance tools derived from these templates. The law was written in Brussels; the operational meaning was written in audit workpapers.¹⁰

Because visibility determines oversight, these firms increasingly shape how regulatory categories are interpreted, how obligations are implemented inside real workflows, how risk is recorded, and how responsibility is assigned. They serve as the interpreters of the AI era — turning the behaviour of frontier models into the concrete classifications that law, finance, and organisational decision-making require.

This arrangement is not temporary. It is a defining feature of the AI era.

VI. The Road Ahead — What the Visibility Monopoly Makes Possible

The Visibility Gap explains why traditional approaches to AI regulation struggle to keep pace with innovation. But diagnosis alone is insufficient. The more consequential question is what this new institutional condition makes possible — and for whom.

The actors capable of observing how AI actually functions across the global economy will consistently acquire influence over how it is managed in practice. For much of the world —

¹⁰ The timeline reflects publicly reported activity by major professional services networks in the period following the Act's entry into force. The specific pace and scope of template development varied across firms and should be treated as a general characterisation of industry response.

outside the G3 — many countries remain stack-dependent: financial centres, regulatory hubs, and the professional networks that intermediate global commerce.

As the U.S. and China consolidate increasingly distinct AI ecosystems, the rest of the world requires connectors and translators. The global economy cannot function as two sealed spheres. This shift elevates the importance of the actors who operate between rule and execution. Part IV scrutinises this space directly — the gap between legal pronouncements and operational reality.

If sight is the precondition of authority, execution is its instrument. And those instruments have moved. The road ahead is defined by this migration. In the AI era, authority no longer resides solely where rules are drafted; increasingly, it resides where those rules are made workable. Understanding this shift — and the forms of power it creates — is essential to understanding how the AI economy will move forward.

Chapter 2: The Rule-Execution Gap

Why Governance Fails at Implementation, Not Legislation

A government can govern only what it can see — but sight alone does not restore control. Chapter 1 identified the Visibility Gap: the structural blindness produced by opacity, fragmentation, and acceleration. Yet even perfect vision does not guarantee control. The deeper failure is execution.

Consider the cockpit of an aircraft. The radar is clear; the storm ahead is unmistakable. The pilot moves the controls. The dashboard responds. Every instrument signals compliance. But the flaps do not move, the plane does not turn; it continues into the storm.

This increasingly resembles the condition of the modern public authority in the AI era. It is operating a dashboard that lies: legislators mandate fairness, regulators publish frameworks, platforms certify adherence — yet the underlying model, the code, the optimisation loop continues unchanged.

The connection between the pilot's hand (the law) and the plane's mechanics (the code) has been severed. The dashboard reflects the appearance of compliance, not the underlying reality. Scholars of administrative law call this the gap between “law on the books” and “law in action.” In AI ecosystems, the gap is wider, faster, and structurally embedded.

The Rule-Execution Gap is the structural separation between the declaration of a rule and its realisation inside the systems meant to execute it. It marks the space between legislative intent and technical implementation — the distance a rule must travel as it moves from legislative chambers and regulatory agencies into the architectures, pipelines, and operational protocols that now run the AI economy. It is where oversight falters, not because public authorities lack formal power, but because they lack the institutional and technical machinery required to translate legal text into executable form.

The Visibility Gap blinds the State. The Rule-Execution Gap disarms it. Even when the State can see, it cannot act. The cockpit lights flicker with the illusion of control, but the control surfaces remain still.

The Rule-Execution Gap is clearest in places where rules collide with machinery. Inside firms forced to convert legal principles into operational protocols, the gap ceases to be abstract and becomes a daily managerial challenge. A familiar scenario in global finance makes the struggle unmistakable.

On a Wednesday morning in Singapore, a global bank convened an emergency meeting in a glass-walled conference room overlooking Marina Bay. Overnight, its credit-risk model — deployed across twenty-seven countries — triggered a spike in adverse-action notices for young borrowers in Southeast Asia. The model had not malfunctioned. It had simply learned.

Around the table sat three people who rarely spoke the same language, even when they used the same words.

Amira, the bank's regional general counsel, carried a binder of precedents under the U.S. Equal Credit Opportunity Act. Under its doctrine, she explained, the model must manage disparate-impact risk — statistical patterns that disproportionately harm a protected class — unless the bank can demonstrate a defensible business-necessity rationale.¹¹

Lukas, the EU compliance lead, countered with a different mandate. Under the transparency and documentation provisions of the EU AI Act, the model qualified as a high-risk system and therefore had to meet strict explainability and traceability requirements. Disparate impact alone was insufficient; the system's internal logic had to be intelligible to regulators and human reviewers.¹²

Chen, the machine-learning engineer flown in from Hong Kong, listened carefully. He knew the model's design: an ensemble combining gradient-boosted decision trees with a neural network layer — a widely adopted pattern in modern credit-risk modelling. He also knew the limitation that would be politically explosive: the model's internal decision pathways could not be fully reconstructed. Post-hoc interpretability tools such as Shapley Additive Explanations (SHAP) or Local Interpretable Model-agnostic Explanations (LIME) could approximate feature contributions, but they could not satisfy the EU's demand for structural explainability, nor the U.S. requirement for a robust business-necessity justification.

The meeting quickly devolved into a jurisdictional standoff. Satisfy U.S. disparate-impact doctrine, and the bank risked violating EU explainability rules. Satisfy EU transparency mandates, and it risked degrading performance in Asia. Satisfy China's algorithm-filing regime, and it risked exposing proprietary logic that breached confidentiality expectations everywhere else. The model had not broken the rules. The rules had broken each other.

By lunch, the team reached the only conclusion possible: they would create three different versions of the model — one for the U.S., one for the EU, and one for Asia — each with different fairness constraints, documentation requirements, and risk tolerances.

This was not an exotic solution. Large multinationals routinely maintain region-specific variants of the same model, each tuned to satisfy incompatible regulatory demands. In

¹¹ Equal Credit Opportunity Act, 15 U.S.C. § 1691 et seq.

¹² Regulation (EU) 2024/1689, arts. 9 (risk management), 11 (technical documentation), 13 (transparency), 15 (accuracy, robustness, cybersecurity).

industry practice, this is called geofencing the model — creating compliance variants that satisfy regulators but fragment the system.

The bank had not solved the fairness problem. It had partitioned it.

This is the Fairness Trap: when firms attempt to operationalise fairness across incompatible legal regimes, the outcome is not harmonisation but balkanisation. The Rule-Execution Gap is not a failure of compliance. It is the practical impossibility of converting legal norms into deployable, coherent technical specifications across jurisdictions that disagree on what fairness means.

The Rule-Execution Gap is not a passing byproduct of regulatory lag, nor a problem solvable by hiring technical talent or accelerating rulemaking cycles. It is fundamental: the workflows that turn rules into operational reality do not exist inside the State — and cannot be built there at scale. In the AI era, the decisive work of oversight occurs in the conversion of principles into procedures that systems can actually follow. That conversion does not happen inside the State. It happens inside the private firms that sit closest to the code and closest to the law at the same time.

In practice, these firms determine how obligations are operationalised and how compliance is demonstrated; the government may articulate the goals, but the machinery that makes those goals real now operates outside it.

I. The Void: Defining the Rule-Execution Gap

What does it mean for a rule to work?

In classical legal theory, the answer is comfortingly simple: a rule “works” when a legitimate authority writes it down, publishes it, and backs it with enforcement. Legislators articulate principles, regulators translate them into regulations, and courts tidy up whatever escapes the first two steps. The chain of custody is linear, civilised, and — on paper — reassuringly manageable.

AI disrupts this model. The distance between a legal principle and a line of code is not a gap. It is a design-level mismatch — a physics-grade disconnect between how law expresses intent and how systems execute behaviour. And that mismatch is not created by complexity alone. It is created by the collision between abstract statutes and probabilistic machinery that do not obey rules so much as navigate constraints.

A. The Fairness Trap: When Math Defies Doctrine

Consider the universal mandate that AI must be “non-discriminatory” — a requirement now embedded in every major regulatory instrument, from the EU AI Act to U.S. Executive Order 14110 to China’s algorithmic governance guidance. It is rooted in decades of anti-discrimination doctrine.¹³

¹³ Executive Order 14110, Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence, October 30, 2023.

But principles do not execute themselves.

To an engineer, “fairness” is not a moral aspiration; it is a statistical constraint. Computer science offers a menu of mathematically rigorous fairness definitions — at least twenty, depending on the granularity applied. Some focus on calibration, asking whether a risk score of “7” should correspond to the same probability of default for every group. Others emphasise error balance, such as equalised odds, which requires false positives and false negatives to occur at the same rate across groups. Still others prioritise demographic parity, asking whether approval rates should be equal across groups.

These are not philosophical positions. They are equations. And here is the crisis: it is mathematically impossible to satisfy all of these fairness conditions at once when base rates differ across groups. As shown by Kleinberg, Mullainathan, and Raghavan (2016) and Chouldechova (2017), calibration within groups cannot be simultaneously achieved with the relevant error-rate balance conditions except in highly constrained special cases. The maths refuses to cooperate.¹⁴

The statute demands non-discrimination. The code demands a trade-off. Optimise for calibration and equalised odds breaks; enforce equalised odds and calibration collapses; impose demographic parity and both accuracy and calibration degrade. The trade-off cannot be avoided. It must be chosen. The State cannot legislate the maths. It cannot outlaw an impossibility theorem. So it retreats to process: bias audits, risk assessments, documentation.

But the choice does not vanish. It simply moves — into the technical workflow where legal abstractions are compressed into a single operational decision.

B. The Transparency Paradox: The Fidelity Trade-Off

The same mismatch appears in mandates for “transparency.” GDPR, the EU AI Act, and the emerging wave of sectoral AI rules all require high-risk systems to be explainable. But technical reality forces a hard choice between fidelity (truth) and intelligibility (human usability).

Engineers can offer two forms of transparency. The first is the technical dump: model architecture, weights, and training data — perfectly faithful, yet functionally unreadable to any human decision-maker. The second is the simplified explanation: a system card or narrative summary that is understandable but necessarily inexact.

The law frames transparency as a right. The system treats it as a tunable constraint. And the transparency mandate dissolves into the same operational void as fairness: a principle that cannot be rendered into a single, coherent technical representation. Someone must decide where to place the slider — how much fidelity a regulator can absorb and how much simplification a system can tolerate. That decision is made upstream, inside the technical

¹⁴ J. Kleinberg, S. Mullainathan, and M. Raghavan, ‘*Inherent Trade-Offs in the Fair Determination of Risk Scores*’ (2016), Proceedings of ITCS 2017; A. Chouldechova, ‘*Fair Prediction with Disparate Impact*’ (2017) 5 Big Data 153-163.

workflow, long before the law arrives.

C. Why the Rule-Execution Gap Persists

The Rule-Execution Gap is the structural separation between the Rules Layer and the Execution Layer — two domains organised by different design logics. The Rules Layer must remain broad, principle-driven, and technology-agnostic; it speaks in norms, rights, and general duties. The Execution Layer must be precise, parameterised, and operational; it speaks in thresholds, metrics, and code paths. These layers rely on different institutional grammars. The gap widens as adaptive, machine-learning systems replace deterministic ones, because rules must increasingly shape systems whose operational logic evolves over time rather than remaining fixed at the point of implementation.

This gap is not a drafting flaw or a temporary mismatch awaiting better legislation. It is enduring. As long as law must remain general and code must remain exact, the firms that translate between them — the ones that choose the metric, set the threshold, and define the evidence — exercise the real control, whether the formal authority notices or not.

II. The Forces Widening the Gap

The Rule-Execution Gap may be inherent, but its size is not constant. In the AI era, three forces are actively widening it to unprecedented scale. These forces mirror those that produced the Visibility Gap in Chapter 1, but here they operate in a different register. There, they were epistemic barriers — obstacles to seeing. Here, they become operational frictions — obstacles to doing.

They do not stop formal authorities from observing the system. They stop them from acting on what they observe.

A. The Epistemic Asymmetry — The Tacit Knowledge Problem

The first force is not that regulators lack information. It is that they lack tacit knowledge — the kind of practical understanding that cannot be written down, cannot be legislated, and cannot be absorbed from a briefing memo.

In classical industries — aviation, banking, pharmaceuticals — regulators often know as much as the regulated. Inspectors at the Federal Aviation Administration (FAA) understand how lift works. Bank examiners at the Federal Reserve can read a balance sheet without hesitation. Reviewers at the Food and Drug Administration (FDA) can parse a clinical trial and understand what happened.

In the AI era, this parity breaks down. Execution requires knowledge that cannot be captured in a statute or a guidance note. It requires the intuition to know which optimisation routine will stabilise a loss function rather than destabilise it; the contextual judgement to understand why a dataset requires a specific cleaning step despite what the documentation claims; and the experiential muscle memory of deployment, earned only by running the system, breaking it, and repairing it under production conditions.

This is the tacit-knowledge gap. Law deals in explicit knowledge — text, definitions, categories. Execution depends on tacit knowledge — heuristics, trade-offs, engineering judgement. A regulator can mandate “fairness,” but cannot write the rule that says “do not set parameter X above 0.5,” because the regulator does not know parameter X exists. The issue is not headcount but practice: tacit knowledge resides where models are built, broken, and repaired — far from the chambers where statutes are drafted.

This produces a quiet inversion of authority. Regulators hold the mandate but lack the practical footing required to carry it out. Firms possess that footing — and the strategic incentives to shape how systems actually run. And so the gap widens, not because firms defy the law, but because they are the only ones whose fingers press the keys.

B. Temporal Acceleration — The Obsolescence Trap

The second force is one of timing. The speed of AI innovation does not merely outpace law; it overtakes law before the law can be implemented. This is not the familiar lament that “law is slow.” It is a deeper mismatch between two incompatible clocks: the cadence of rulemaking — deliberative, procedural, legitimacy-driven — and the cadence of technical change — iterative, experimental, competitive. One is designed to prevent rash action. The other is designed to reward it.

The history of online tracking makes the pattern clear. By the time consent pop-ups were standardised under the GDPR and reinforced through European Data Protection Board guidance in 2019–2020, the tracking landscape had already moved on. Client-side cookies — the very mechanism the pop-ups were engineered to police — were being displaced by server-side tagging, first-party data pipelines, and probabilistic fingerprinting techniques that had been accelerating since 2016. Browser vendors had begun deploying tracking protections that bypassed cookies entirely, and by 2025 industry analyses estimated that a majority of major websites relied on fingerprinting signals invisible to consent banners. The GDPR’s carefully engineered consent framework targeted a mechanism the industry had already begun abandoning.

In AI, this dynamic is amplified. A rule designed for one model generation often arrives just in time to regulate the one that no longer exists. A framework drafted for interpretable, rule-based decision trees becomes misaligned the moment it is applied to neural networks, whose internal logic is less “if-then” and more “trust me, it works”.

A rule written for static, train-once-deploy-once models becomes inadequate when applied to continuous-learning systems that update themselves post-deployment through fine-tuning, reinforcement learning, or federated learning. A policy crafted for chatbots — prompt in, text out — becomes irrelevant when applied to autonomous agents that chain actions, call tools, navigate environments, and make long-horizon decisions regulators never contemplated.

Regulatory frameworks emerging in 2024–2026 — the EU AI Act, U.S. executive orders, OECD guidelines — were drafted for the AI assumptions of 2020–2023: tabular machine learning, early large language models (LLMs), and static risk-assessment pipelines. Meanwhile, the industry has already shifted to multimodal agents, world-model-driven planning, inference-time scaling, and model families that behave less like software and more like autonomous problem-solvers.

The baseline process duties in these rules still matter, but the fit between rule and reality erodes with every shift in underlying practice. This is the Obsolescence Trap: by the time a rule is drafted, the practice has evolved; by the time it is enforced, the system design has changed; by the time it is revised, the field has moved again.

C. Jurisdictional Fragmentation — The Compliance Gridlock

The third force is spatial. AI operates across borders; law remains stubbornly territorial. A multinational deploying a single AI system does not face one set of rules. It faces hundreds — overlapping, contradictory, and often mutually exclusive — each backed by a government convinced it is the centre of the regulatory universe. And each is correct — within its borders.

The U.S. regulates privacy through sector-specific statutes, state-level variation, and ex-post enforcement. Health data falls under the Health Insurance Portability and Accountability Act (HIPAA). Financial data under the Gramm–Leach–Bliley Act (GLBA). Children’s data under the Children’s Online Privacy Protection Act (COPPA). Everything else is governed by the FTC’s authority over “unfair or deceptive” practices.¹⁵

By 2026, approximately twenty states will have enacted comprehensive privacy laws — each with distinct thresholds, exemptions, cure periods, and enforcement mechanisms. California’s regime is expansive enough to function as a quasi-independent privacy jurisdiction. A federal baseline remains absent. The compliance environment is defined by fragmentation and after-the-fact enforcement rather than unified statutory command.

The EU approaches regulation through comprehensive, procedure-driven frameworks. The GDPR enforces purpose limitation, data minimisation, and accountability through detailed statutory obligations. The EU AI Act adds a risk-based classification system that elevates documentation to a central governance tool: high-risk systems must produce records extensive enough to support full lifecycle auditability. Enforcement is coordinated through national Data Protection Authorities under the guidance of the European Data Protection Board. The EU model prioritises procedural completeness and formal accountability over speed, producing a highly structured but slow-moving regulatory order.

China regulates data as a strategic resource — because, through its legal lens, it is. The Cybersecurity Law, Data Security Law, and Personal Information Protection Law (PIPL) establish data-localisation requirements, cross-border transfer reviews, and security assessments for activities that may implicate national security — a category defined with broad scope. AI models must be filed, labelled, and auditable under the Algorithmic Recommendation Provisions and the Generative AI Measures. China’s framework is directive and state-centred, emphasising government visibility and control over data flows and system behaviour.¹⁶

¹⁵ Health Insurance Portability and Accountability Act of 1996 (HIPAA); Gramm–Leach–Bliley Act (GLBA); Children’s Online Privacy Protection Act, 15 U.S.C. §§ 6501–6506 (COPPA).

¹⁶ Cybersecurity Law of the People’s Republic of China (2017); Data Security Law (2021); Personal Information Protection Law (2021); Provisions on the Management of Algorithmic Recommendations (2022); Interim Measures for Generative AI Services (2023).

The gridlock emerges from the collision. The U.S. emphasises disclosure and litigation readiness. The EU emphasises purpose limitation and documentation. China emphasises state primacy, localisation, and oversight of data flows. No firm can build a separate compliance system for every jurisdiction. Operational reality demands a unified approach. But no country can write the rules for a unified global approach.

Europe cannot tell a U.S. firm how to satisfy China’s localisation mandates; China cannot tell a European firm how to comply with U.S. securities disclosure; the United States cannot tell anyone how to reconcile GDPR’s purpose limitation with China’s data-control regime. Each insists on its own logic, but none has the authority or line of sight to reconcile the others.

This is the Compliance Gridlock: incompatible commands from multiple sovereigns, all of which must be obeyed, and none of which can be reconciled by the states that issued them. International coordination sounds elegant — treaties, standards, harmonisation. But these require shared incentives, definitions, and enforcement. The U.S., the EU, and China share none of these — not even a common understanding of what “data” is.

And so the synthesis rule — the operating grammar that lets rival national commands coexist — is authored outside the State. The mechanics of this cross-border reconciliation are examined in Part IV, where the Translation Layer’s role in making incompatible commands run in the same system becomes explicit.

III. Why Traditional Solutions Cannot Close the Gap

Recognising the Rule-Execution Gap as inherent to AI governance does not mean regulators have ignored it. Across the G20 and beyond, governments have deployed every familiar instrument in the public-policy toolkit: more resources, faster processes, international coordination, mandated disclosure, and the occasional moonshot task force.

These efforts are earnest, sophisticated, and often heroic. They are also incapable of closing the gap — not because regulators are incompetent, but because they are attempting to solve a machine-intelligence challenge with industrial-era tools.

A. More Resources — The Capacity Fallacy

The instinctive response to the Rule-Execution Gap is to add capacity. Hire engineers, expand budgets, create internal labs. These moves help. They improve oversight. They sharpen questions. They reduce the number of hearings where everyone wishes the microphones had failed. But they do not close the gap.

A regulator can hire brilliant people and still lack what the task requires: not abstract expertise, but the lived fluency that comes from operating models under real-world load. That fluency does not arrive with credentials. It forms only inside labs that design, build, ship, scale, and repair models in motion. The State can recruit talent; it cannot replicate the conditions that produce operational judgement.

The economics of AI make the gap more entrenched. Top talent gravitates to where frontier models live: frontier labs, hyperscalers, start-ups and high-growth firms. The public sector cannot match compensation, stock options, compute, or the freedom to experiment. And even

when agencies recruit strong talent, they cannot recreate the ecosystem in which tacit knowledge forms. What remains is a persistent asymmetry: the State can study these models, but only firms learn by building and running them.

B. Faster Processes — The Speed Fallacy

A second plan is to accelerate rulemaking. Governments experiment with regulatory sandboxes, expedited consultations, agile rulemaking, and rolling guidance. These tools help. They move regulators closer to the frontier. They occasionally narrow the distance between discovery and response — but never enough to change the underlying dynamic.

Acceleration cannot overcome the widening gap between the pace of democratic decision-making and the speed of AI innovation. Democratic authority moves deliberately. Notice, comment, deliberation, revision, judicial review — these are not administrative delays. They are the processes through which law earns its authority. Compress them too far, and regulation does not become more effective. It becomes less legitimate.

Meanwhile, AI moves on a far faster clock: models shift weekly, platforms evolve monthly, and practices change long before regulations clear their procedural steps. The government operates on a rhythm that cannot make rules at the speed of silicon without losing the legitimacy that makes legislation worth following.

C. International Harmonisation — The Consensus Fallacy

If fragmentation creates execution gridlock, harmonisation seems like the obvious cure. And on paper, the world has never been more aligned. More than eighty countries and international bodies — from the OECD to UNESCO to the G7 — have endorsed high-level AI principles: fairness, transparency, accountability, human rights, safety. The consensus is impressive. It is also almost entirely non-operational at the level where execution occurs.¹⁷

These frameworks harmonise aspirations, not implementation. They declare that AI should be “trustworthy,” but they do not specify which fairness metric to apply, what thresholds constitute compliance, how to audit model internals, how to reconcile data sovereignty with global training pipelines, or who enforces what — and with what authority. The world agrees on principles. It disagrees on everything that makes principles real.

Fairness illustrates the pattern. The EU mandates mitigation but declines to choose a metric. The U.S., through NIST, recommends evaluating multiple definitions. China emphasises outcome equity without prescribing statistical form. Three regimes, three incompatible implementations of the same idea — and fairness is the easy case. The same divergence appears across transparency, data governance, harm definitions, and enforcement.

The major powers have no incentive to converge. The U.S., China, and the EU operate with incompatible logics. The U.S. is innovation-driven, litigation-heavy, and market-first. China is sovereignty-driven, security-first, and state-centred. The EU is rights-driven, process-first,

¹⁷ OECD, OECD/LEGAL/0449 (2019, amended 2024); UNESCO (2021); G7 Hiroshima Process Code of Conduct (2023).

and documentation-intensive. These are not stylistic differences. They are different theories of the State.

Even when harmonisation succeeds on paper, it often fails in practice. Treaties and multilateral initiatives produce elegant principles — and then exclude the major players or lack enforcement mechanisms. This is not convergence but fragmentation with better branding. Harmonisation without shared machinery for doing the work produces convergent vocabulary and divergent practice.

D. More Disclosure — The Transparency Fallacy

A final strategy is to mandate more disclosure: algorithmic impact assessments, model cards, data sheets, audit logs, transparency reports. These tools help. They improve situational awareness. They create paper trails. They help regulators ask better questions. But they do not close the gap.

Disclosure shows what a system is supposed to do. Execution reveals what it actually does.

Documentation is static, declarative, and often curated. System behaviour is dynamic, contextual, and sometimes opaque even to the people who built it. A model card can describe the training data. It cannot explain why the system behaved one way on Monday and another on Tuesday.

Transparency strengthens accountability at the Rules Layer. But the gap lives in the Execution Layer. And no amount of disclosure gives regulators the ability to specify how a model should be trained, how data should be structured, how risk should be mitigated in real time, how an agent should behave when interacting with tools, or how emergent behaviour should be constrained.

Each remedy — capacity, speed, consensus, disclosure — treats symptoms, not the underlying condition. The Rule-Execution Gap persists, but it does not stay empty. The space law cannot reach is filled — quietly, predictably — by the actors that can execute.

IV. Professional Services as Execution Infrastructure

If states cannot execute their own rules, someone must. And someone already does. The actors filling that void are not startups or platforms, but the institutions that have been present all along: professional services firms. Global law firms and the Big Four accounting networks — Deloitte, PwC, EY, and KPMG — have become the operational machinery of the AI economy.

They no longer confine themselves to advisory memos or statutory interpretation. Their work is execution: turning abstract requirements into controls, workflows, documentation regimes, and the templates global clients actually use. Public authorities issue requirements; these firms convert them into operational form. That conversion is their defining role.

These firms bridge the Rule-Execution Gap through three capabilities governments cannot match at scale. They possess template memory: accumulated exposure to hundreds of legal regimes at once, enabling rapid pattern recognition and conflict resolution. They command

implementation capacity: the ability to embed requirements through controls, workflows, and system-specific platforms. And they exercise synthesis authority: the mandate to reconcile conflicting demands into a single functioning template.

Once templates are set, they become the defaults that shape how systems actually run, often exerting more practical influence than the legal text that initiated them. These firms do not wield enforcement power. But they define the boundaries within which AI systems can be designed, built, deployed, evaluated, priced, and permitted to exist.

A. Cross-Jurisdictional Observation — The Visibility Advantage

A single Big Four network works with thousands of multinational clients across more than one hundred countries. Through this client work, it gains a form of cross-border insight no public authority can replicate: a practical view of how companies design, deploy, and adapt AI systems under multiple, sometimes conflicting, legal regimes.

Where public bodies see only the materials submitted to them, these firms see how requirements interact across borders — where they align, where they clash, and how companies reconcile those tensions in daily operations. They observe the engineering choices, the workarounds, the points where processes strain under competing obligations, and the adjustments companies make long before any regulator initiates a review.

They work simultaneously on legal requirements, financial reporting, and operational workflows — a combination that gives them a composite view few other institutions possess. In a world where AI moves globally while rules remain territorial, the firms with the broadest observational reach hold the decisive informational advantage.

B. The Mechanics of Translation — Writing the Operational Code

How does a professional services firm turn an AI-specific legal mandate into something a company can actually operate? Through a structured workflow that mirrors established compliance and assurance processes — systematic enough to repeat, flexible enough to adapt across industries.

Across thousands of engagements, a consistent pattern emerges: a five-stage process that generates the operational code AI platforms rely on. The process begins with rule identification — a comprehensive mapping of applicable obligations including statutory requirements, sector-specific rules, state-level privacy laws, industry mandates, contractual terms, and soft-law frameworks from bodies such as ISO, OECD, and NIST. Public authorities see their own regulations; these firms see the composite set of obligations client companies must satisfy simultaneously.

Next comes the feasibility assessment: determining which requirements are mathematically unrealistic, computationally intensive, or operationally incompatible. Broad principles such as “fairness,” “explainability,” or “continuous oversight” are converted into measurable thresholds, documentation practices, interpretability tools, monitoring cadences, and alert triggers. This is where legal ambition meets technical constraint — the point at which the rule begins to take operational form.

The third stage is conflict resolution — the engineering diplomacy. Where requirements conflict — data minimisation versus disclosure, localisation versus global pipelines, purpose limitation versus retraining cycles — the firm designs the accommodations. It builds data-segmentation solutions, jurisdiction-specific model variants, federated training pipelines, dual-layer audit trails, synthetic-data buffers, and contractual carve-outs. This work is rarely visible to regulators, but it is where incompatible demands are converted into a single workable configuration.

Once conflicts are resolved, the firm synthesises the surviving requirements into a unified operational template — the risk protocol, the fairness metric, the documentation workflow, the audit structure. This template becomes the rulebook the client implements — a rulebook no regulator drafted, but every regulator will eventually encounter.

Finally, the firm assembles the defensibility instruments — audit logs, model documentation, impact assessments, governance memos, board attestations, and evidence packets put together to demonstrate reasonable effort in future reviews. This becomes the evidentiary shield in disputes — and often the primary record regulators rely on to understand how the system operated.

The distinction is clear: regulators articulate principles; professional services firms design the protocols that make those principles actionable. And in the AI economy, the protocols determine how the principles work in action.

C. Synthesis Authority — Reconciling Incompatible Sovereigns

National directives collide, and no harmonisation mechanism resolves those collisions. The most consequential capability of professional services firms is their ability to navigate cross-border legal contradictions — commands issued by different governments that cannot be satisfied at the same time.

A global AI deployment is shaped not by a single legal regime but by a set of incompatible ones, each backed by national enforcement tools. One country demands expanded insight into models and data. Another restricts the disclosures the first requires. A third prohibits the same data from crossing its borders. Each authority issues its mandate with confidence. None provides a method for complying with the others. No regulator resolves these collisions. The translators do.

These professionals determine — quietly and with wide-ranging consequence — which requirement becomes the binding constraint, which risks can be absorbed, which must be engineered out, and which designs can survive contradictory directives. This work is not advisory. It is determinative.

Global AI deployments do not follow the text of national laws. They operate on a single, coherent blueprint engineered to withstand their contradictions. That blueprint is not issued by any legislature or regulator; it is assembled through the work of private firms converting incompatible demands into a workable design.

D. The Inversion of Governance

The Rule-Execution Gap does more than reassign tasks from public authorities to private actors; it alters the sequence through which rules take operational form. The firms that design the workflows, controls, and assurance mechanisms that AI systems follow now shape the terrain long before formal instructions arrive.

The shift is clearest when viewed as a sequence. The traditional chain runs: Legislature → Regulation → Implementation. The AI-era order increasingly reverses: Implementation → Standard → Regulation → Legislation.

Professional services firms develop the execution templates first — the blended standards, risk taxonomies, documentation practices, and audit structures that multinationals rely on to deploy AI globally. These templates diffuse across industries and national borders as working norms. Regulators encounter them next, often treating them as evidence of “reasonable” practice. Legislatures arrive later, formalising elements that have already taken root in practice. In AI ecosystems, implementation often leads and law follows.

V. Blended Templates as Operational Legislation

If professional services firms form the execution apparatus of the AI economy, then blended templates are the instruments through which execution becomes systematised and repeatable.

They are not merely documents. They are the operating schematics that turn abstract principles into concrete procedures, controls, and workflows. A template is the point where legal intent is refactored into system design — the logic that companies follow.

Templates exist because global firms cannot engineer bespoke compliance for every jurisdiction. The complexity is prohibitive. The cost is unsustainable. The operational risk is untenable. A multinational cannot maintain separate AI rulebooks for the EU, the U.S., China, India, and Brazil — and still more variants for sectoral and state-level directives. The result is a combinatorial explosion. When dozens of regulatory regimes speak at once, someone must choose what becomes the baseline and what becomes the deviation. Blended templates make that choice — not as replacements for statutes, but as the mechanism through which statutes acquire operational form.

A. The Algorithm of Rule Selection

If blended templates function as rule-forming instruments for AI, then one question defines their logic: which rules get in? Templates cannot absorb every directive. The global rulebook is too vast, too inconsistent, and too unevenly enforced. In practice, compliance architects triage — following a disciplined, repeatable process that determines what enters the baseline and what stays behind.

The first question is survival: does this rule control market access? If a requirement blocks deployment, it becomes the global baseline; everything else bends around it. The second question is pain: does this rule carry credible enforcement? A regime with meaningful sanctions and demonstrated willingness to use them does not request priority; it commands

it. The third question is defence: does this rule shape the evidentiary trail? Documentation requirements, attestations, and audit logs do not define the system's design; they determine how well it holds up in court. The fourth question is exposure: does this rule carry reputational consequences? Soft-law commitments and public expectations are adopted if they do not conflict with the answers above.

This triage determines which rules become the common baseline, which remain country-specific deviations, and which never advance beyond aspiration. Part IV formalises the logic beneath it.

B. The Maximal Compliance Strategy — The Ratchet Effect

When rules diverge but do not directly conflict, blended templates tend to settle on the most demanding interpretation. The outcome is dictated by economics, not moral aspiration. A single global rulebook is cheaper, safer, and easier to update than a patchwork of national and regional variants. Every deviation adds cost, complexity, and new failure modes. Uniformity wins.

The strictest rule often gets applied everywhere. When the GDPR requires explicit consent for certain data uses, the template shifts to explicit consent worldwide, even where implied consent would suffice. The same dynamic appears in retention limits, documentation burdens, and age-verification standards. A quiet ratchet operates: once the most demanding rule enters the template, it becomes the baseline across the world. No country negotiates global influence. The template confers it.

C. Conflict Resolution — Structural Workarounds

When national mandates directly collide, blended templates rely on engineered workarounds — technical solutions that make otherwise clashing obligations simultaneously survivable.

Common methods follow predictable patterns. Firms use data sharding to separate information so that no regulator, subpoena, or platform ever sees the whole dataset; each shard satisfies its local rule, and the global system avoids conflict by never assembling the full picture.

They deploy federated architectures, training models locally and sharing only weight updates so the model learns at scale while the underlying records never leave their country of origin. They construct contractual firewalls — corporate and legal arrangements that prevent a parent company from accessing a subsidiary's data, even if it wants to — turning the firewall into an enforceable control rather than a metaphor.

These workarounds do not appear in any statute. They appear in templates. And once they do, they become the applied meaning of the law — the mechanism through which conflicting national directives are made practical. The template does not settle the conflict on paper. It settles it in building the AI model.

D. Template Propagation — The Governance Contagion

Templates do not remain confined to a single client. They spread through the global AI economy with the steady rhythm of a governance contagion. The engine is not mysterious. It is a repeatable diffusion dynamic — the same pattern that has operated for decades in accounting, engineering, and finance. What is new is the speed and the stakes.

The sequence typically unfolds as follows. A Big Four firm develops a high-end blended template for a flagship client launching its first worldwide AI deployment into a maze of regulatory regimes. Successful deployments harden the template — each pass smoothing an edge, clarifying the logic, and strengthening the defensibility shield. The same form is then applied to hundreds of mid-market clients who cannot afford bespoke solutions but can afford the template. Competitors replicate it because clients expect an “industry standard” regardless of provider. The template becomes the de facto standard, not by decree but by diffusion.

Regulators eventually encounter the template during audits and, seeing it often enough, begin treating it as the standard expression of “reasonable compliance.” By the time a regulator writes a formal rule, the template has already become the industry or sector’s working norm — created with no vote, no regulator’s pen, and no public voice.

The template becomes applied law through use, not decree. Practice outruns doctrine. The models spread first. The law arrives later — if it ever does.

VI. The Translation Layer

A. The Formation — How a Class Emerged Without a Charter

Blended templates are instruments. The people who construct them constitute something larger: a governance class in formation. Professional services firms — and the hybrid professionals within them — have assumed a structural role that no prior institutional design anticipated.

Templates are not new, and neither are the translators who build them. For decades, professional services firms have converted legal requirements into operational controls — building the compliance frameworks through which multinationals navigate cross-border regulation. The work was complex, but it was sequential: rules stabilised first, institutional knowledge accumulated, and templates followed. The translator arrived after the terrain was mapped.

In the AI era, that sequence has collapsed. The terrain is shifting while the map is being drawn. These professionals are no longer interpreting settled law for familiar industries. They are constructing new operational rules for new systems that evolve faster than the statutes that govern them. They must reconcile regulatory philosophies that disagree not merely on the rules but on what AI is. The U.S., China, and the EU are not enacting different versions of the same law; they are producing different knowledge ecosystems — different assumptions about what machine intelligence constitutes, how it should be measured, and who bears responsibility for its actions.

The Translation Layer did not emerge through legislation, revolution, or institutional design. It formed the way most coordinating classes do: through function, not proclamation. Two forces produced it.

The first is the convergence of expertise. AI governance draws from engineering, law, accounting, audit, risk, operations, finance, and data management — disciplines that rarely share a corridor, let alone a methodology. No single field can translate legal rules into machine behaviour. A lawyer cannot build a scalable model pipeline. An engineer cannot interpret a cross-border discovery order. An auditor cannot reconcile conflicting privacy regimes. But the Translation Layer — collectively — can. This convergence produces a new kind of professional: the hybrid translator, fluent in legal principle, measurement discipline, and technical constraint. It is a role no university teaches, but every global system now requires.

The second is the emergence of a shared methodology — an invisible curriculum. Members of the translator class are trained not by universities or governments but through Big Law and Big Four onboarding, client-playbook reuse, clause banks, audit frameworks, and cross-engagement knowledge transfer that carries methods from one industry to the next and one country to the next. Over time, this produces a shared worldview about what counts as reasonable, defensible, compliant, or safe enough. These norms are not written in law or standards. They are written in practice. And practice, repeated at scale, becomes the operating default.

The result is a governance class in all but name — not a political movement, not a regulatory agency, not a multinational corporation, but a class of practitioners whose hybrid expertise and shared methods give them the capacity to determine how AI actually behaves. This did not announce itself. As in earlier periods of institutional transformation, the shift emerged because the system required it — not by design, but by necessity. The world needed a translator, and it produced one.

B. The Mechanics of Influence — How the Translation Layer Governs Without Governing

There is no settled body of precedent to fall back on. The translation task is no longer iterative and local; it is concurrent with global diffusion. Templates must be built while the technology evolves, deployed while the rules remain incomplete, and relied upon with minimal margin for error — because the decisions they shape affect credit access, medical diagnosis, legal outcomes, employment, public safety, and the informational environment through which entire societies navigate daily life.

What distinguishes AI templates from every prior instrument that shares their name is not their existence but the consequence of their operation. In practice, these templates increasingly function as the real AI rulebooks for the entire world — the mechanisms organisations rely on to determine how models are designed, built, deployed, and controlled. No legislature drafted them. No electorate approved them. Yet they govern.

They occupy a distinctive position: embedded everywhere but owned nowhere, operating across the globe, across industries, across legal regimes, and across corporate hierarchies while belonging to none. They are accountable to clients yet constrained by law — too tied to legal regimes to ignore the law, too tied to operations to ignore reality, and too tied to markets to ignore incentives.

Their influence is not declared; it is procedural. It operates through three processes that together determine how AI governance functions in practice.

The first is control of the workflow. The translators design the processes that determine how data is collected, how models are trained, how risks are assessed, how audits are performed, and how decisions are escalated. In complex systems, the default is destiny. Once a workflow is embedded in a template, it becomes the path of least resistance for thousands of engineers, lawyers, auditors, and product teams.

The second is control of interpretation. Legal principles enter the process as abstractions — fair, transparent, safe, lawful, reasonable. The translators decide how those concepts become concrete practice: which fairness metric is acceptable, which transparency output satisfies scrutiny, which safety threshold is defensible, and which trade-offs are permitted under conflicting rules. Interpretation is not commentary. It is constraint-setting.

The third is control of propagation. Once a workflow is embedded, it travels — not merely as a template, but as a method. Partners rotate across countries and regions. Teams carry practices from one engagement to the next. What begins as an answer to a single client's question becomes the default approach for dozens of others. A method repeated across hundreds of systems becomes a norm. A norm repeated across the world becomes an expectation. An expectation encountered repeatedly by regulators becomes a baseline.

The translation class shapes outcomes without ever appearing to rule. It steers the system because the system routes its hardest problems through it. Yet its power has boundaries. It cannot manufacture legitimacy — it can operationalise the priorities others set, but it cannot pick the values themselves. It cannot override hard sovereignty — when a state draws a bright line backed by coercive force, no amount of workflow design can engineer around it. And it cannot resolve political conflict — some problems are not engineering problems but societal disputes dressed as technical puzzles.

The functional consequence is clear. The Translation Layer holds real influence, but influence is not democratic authority. Necessity does not create legitimacy, and being the only group capable of execution does not make it the group society has chosen to lead. This is the unresolved tension that carries the book forward.

When decision-making shifts to private firms that can implement solutions but were never designed to speak for the public, the question becomes: how can their influence be anchored to something the public can observe, hold accountable, and trust? That challenge is taken up directly in Part V.

VII. Conclusion — Where Law Fails and Governance Begins

The Rule-Execution Gap is the distance between the rules governments issue and the behaviour AI systems ultimately produce. It is not a malfunction. It is the predictable outcome of an era in which legal authority moves deliberately, models evolve at unprecedented speed, and global companies cannot wait for alignment. It is the void where real governance begins.

States legislate in broad principles; AI operates through precise procedures. One speaks in aspirations, the other in mechanisms. The Translation Layer converts one into the other, creating the operational form beneath statutory law. The statute declares the goal. The implementation determines the outcome.

The Gap persists because law is territorial but AI is borderless; because sovereign commands conflict but implementation must converge; because technical specificity outruns legislative capacity; and because AI time outruns democratic time.

Understanding the Gap completes the diagnosis of Part I, which has identified two structural failures: the Visibility Gap — the State cannot see the systems it governs — and the Rule-Execution Gap — the State cannot convert its commands into operational reality. Both describe what the State can no longer do.

Part II asks the deeper question: why can the State's own constitutive languages no longer express the systems it seeks to regulate? Section I turns to law as the language of power and introduces the Categorical Gap — the mismatch between legal categories and the objects they must now classify. Section II turns to accounting as the language of measurement and introduces the Measurement Gap — the instability of the metrics that mediate between economic reality and institutional recognition.

If Part I shows the State has lost control of the cockpit, Part II reveals the deeper truth: the cockpit's instruments were never built for this aircraft.

PART II — THE GRAMMAR OF GOVERNANCE

Protocol Note — v0.9 Release: Part II integrates the material drafted as Chapters 3 and 4; the full versions arrive in v1.0.

The Dual Languages of Governance

The failures catalogued in Part I — of sight and of execution — share a common root: governance is a language before it is an institution. Before the State can regulate, it must classify. Before it can enforce, it must measure. Before it can hold an actor accountable, it must render that actor legible in terms its institutions can process. The operational breakdowns of the Visibility Gap and the Rule-Execution Gap are, at bottom, failures of grammar — failures of the categories, classifications, and measurement structures through which institutional authority operates.

Law is that language — a syntax for classifying actors, actions, and obligations, whose authority rests not on moral aspiration but on its ability to make complex activity legible and enforceable. But for AI systems, law alone is insufficient. An AI model generates economic consequences — it allocates capital, prices risk, displaces labour, concentrates market power — yet none of these effects are self-evidently legal events until they are measured, classified, and rendered in financial terms. This is the work of accounting: to translate operational reality into the categories through which institutions recognise value, allocate cost, and assign responsibility.

Governing AI therefore requires two languages operating in concert: law as the language of authority — permission, prohibition, liability, and obligation — and accounting as the language of measurement — value, classification, and economic representation. Neither alone can make an AI system fully legible to the institutional order. Law without accounting cannot see what the system produces; accounting without law cannot say what the system owes. The dual-language thesis — that AI governance is incomplete unless both grammars operate together — is the foundation on which the rest of this manuscript builds.

This claim invites an obvious objection: governance of AI is already served by technical standards, ethics frameworks, and policy guidelines, each of which addresses some dimension of the problem. The objection is valid but incomplete. Technical standards specify how systems should behave but cannot allocate rights, assign liability, or determine what counts as a recognisable economic object. Ethics frameworks articulate values but lack the institutional machinery to create enforceable obligations or categories of measurement. Policy guidelines signal regulatory intent but do not generate the classifications, thresholds, or recognition rules that organisations must operationalise and auditors must verify.

Each contributes to the governance environment; none constitutes it. These instruments are complements, not substitutes — they operate within the space that law and accounting create, not in place of it. Without law to define the permission structure and accounting to define the measurement structure, AI remains legible only to its operators, not to the institutional order that claims authority over it.

Both languages were built for an industrial world. They assume human agency, fixed territory, stable time, and separable assets. AI strains that grammar. It operates through distributed computation, probabilistic outputs, evolving models, and intangible capabilities that resist stable classification.

Authorities have attempted to retrofit these languages — through European risk classifications, American liability regimes, and Chinese administrative supervision. These approaches narrow the gap but cannot close it, because the underlying syntaxes remain misaligned with the systems they are asked to describe.¹⁸

¹⁸ Regulation (EU) 2024/1689 (EU AI Act); Executive Order 14110 (2023); NIST, AI Risk Management Framework (2023); Restatement (Third) of Torts: Products Liability (1998); CAC, Provisions on Algorithmic Recommendation (2022); CAC et al., Interim Measures for Generative AI Services (2023).

The result is two new structural failures: a Categorical Gap, in which the State cannot classify the systems it seeks to govern, and a Measurement Gap, in which it cannot measure the systems that now drive economic value.¹⁹

The effect is constitutional rather than technical. As the State's languages fall out of alignment with the AI ecosystems they are meant to oversee, interpretive authority migrates to the Translation Layer — not by choice, but by structural necessity. Throughout this manuscript, Big Law and the Big Four are treated as a functional class because they perform this translation role. This functional cohesion does not imply organisational unity. The Translation Layer is internally fragmented: firms compete fiercely, operate under different regulatory exposures, and navigate distinct liability regimes. Their unity appears only at the level of constitutional effect, not corporate coordination.

I. Law As The Language Of Power

The First Function: Naming What Exists

Governance is, at its core, the power to classify — to render a fluid and indeterminate reality into the fixed categories of a legal order. Before a State can regulate, tax, or prohibit, it must first name what exists. Legal classification performs this foundational act: it stabilises meaning, assigns identities, and transforms unstructured activity into objects the State can recognise and act upon.

In the AI era, this act of construction determines whether a technical system becomes legible to the State at all. Law assigns the categories through which permission, prohibition, ownership, and responsibility attach. Without them, an AI model remains a technical artefact. With them, it becomes an object of attribution and regulatory control.

But the Westphalian legal order was built for land, factories, and human intent. Its categories assume territorial jurisdiction, identifiable agents, and deterministic causation.²⁰

AI operates through distributed computation, probabilistic outputs, and continuously evolving capabilities — none of which map cleanly onto the existing grammar. The result is not a system under strain but a classification framework encountering objects it was never designed to describe.

Jurisdictions have responded with AI-specific liability rules, regulatory sandboxes, and administrative controls. These narrow the distance but do not close it. The categories that once made economic life legible now capture only fragments of the systems they are meant to classify.

¹⁹ The Categorical Gap, Measurement Gap, and Translation Layer analysis are original conceptual contributions introduced in this manuscript, building on the Visibility Gap and Rule-Execution Gap established in Part I.

²⁰ L. Gross, 'The Peace of Westphalia, 1648-1948' (1948) 42 American Journal of International Law 20.

The Industrial Logic versus the AI Universe

Modern law was built for an industrial world. Its categories assume deterministic causation, human agency, territorial jurisdiction, and stable, identifiable objects. The mismatch with AI is not marginal; it is foundational.

Five assumptions embedded in the legal order illustrate the rupture. Law assigns liability through intent, yet AI systems cause harm without possessing intent at all. Jurisdiction follows physical presence, yet inference can occur across borders without any territorial anchor. Law governs through advance planning, yet models evolve continuously after deployment. Legal categories distinguish cleanly between person and thing, yet AI thins that distinction — operating simultaneously as product, service, tool, and agent. And property law presumes separability, yet the weights, data, and compute that constitute an AI system are deeply entangled.

These assumptions governed railroads and telegraphs because those technologies operated within the same world the law assumed — deterministic, territorial, and anchored in human agency. AI breaks that alignment. The result is a Categorical Gap: a thinning of meaning in which legal categories such as negligence or product defect cannot attach cleanly to technical realities such as stochastic hallucination or parameter drift. The translation machinery cannot map these inputs onto stable legal categories — not because the categories are poorly drafted, but because they were designed for a different class of objects.

The Categorical Gap in Practice

The Categorical Gap does not strike law uniformly. It propagates through distinct legal domains, each absorbing the mismatch between legal grammar and statistical behaviour in its own way.

Public law encounters the gap first. Jurisdiction presumes a discrete actor, a locus of intent, and a place where due process can attach. AI weakens these anchors by distributing agency across developers, deployers, datasets, and infrastructure. Duties still exist, but their attachment points become indirect and fragile. Responsibility migrates upstream to operators and infrastructure holders — a shift from direct control to oversight-based duties of care.

Private law absorbs the next shock. Contract assumes a stable subject matter and a knowable allocation of risk. AI does not stay still: it updates, drifts, and reconfigures itself. Warranties fail under temporal uncertainty, and parties retreat to outcome-based obligations and service-level agreements — a tacit admission that the object of the agreement does not remain fixed.

Tort law inherits the causation problem. Tort doctrine requires a legible causal chain. Non-linear behaviour obscures causation, making negligence harder to prove and pushing liability toward insurance, aggregation, and other probabilistic mechanisms. The doctrines remain, but their reach narrows as causal clarity erodes.

Infrastructure regulation marks the most architecturally significant shift. Platforms controlling compute, access, and model behaviour function as de facto regulatory nodes. They enforce constraints through technical design, API gating, and contractual terms. The State defines the boundaries; private firms execute the controls. Law has not disappeared; its point of execution has migrated.

Across these domains, the pattern is consistent: legal categories designed for deterministic, territorial, human-driven activity are being applied to systems that are none of these things. The categories still function — courts still adjudicate, contracts still bind — but the grammar no longer maps cleanly onto its objects. The gap between legal language and technical reality is not closing. It is being managed, provisionally, by the Translation Layer.

The Tripartite Failure and the Broken Translator

The Visibility, Rule-Execution, and Categorical Gaps form a tripartite failure — not of governance itself, but of its industrial-era instruments. Together, they define the inherent limits of governing AI through legal categories built for a different class of objects. The State has not lost authority; it has lost the capacity to exercise that authority directly.

Legal regimes manage this condition differently — the U.S. through ex-post litigation, the EU through ex-ante classification, China through administrative and technical supervision — but each approach mitigates symptoms while leaving the structural divergence intact. None closes the gap between Westphalian grammar and the systems it is now asked to govern.

The result is a broken translator. Legal commands pass through technical systems and institutional intermediaries never designed to carry them. The system continues to function, but law no longer operates through direct command. It operates through translation — and the Translation Layer, not the State, determines the fidelity of that translation.

This is the hinge between the two halves of Part II. The same structural mismatch that destabilises law's classificatory categories also undermines the State's measurement language. Accounting, like law, was built for a world of stable, separable, and human-directed activity. It now confronts systems that are none of these things — and it, too, must rely on intermediaries to make AI legible.

A. How Law Decides What Is AI

The State cannot regulate what it cannot name. Classification — the assignment of a legal type — is the first act of governance, and in the AI era it functions less like policy and more like a strict type system. Law receives source material — model weights, training data, architecture, deployment context — and attempts to cast it into categories built for discrete, stable, human-centred objects. When the cast succeeds, liability attaches, rights activate, duties bind. When it fails, the system becomes legally indeterminate: economically powerful but without a stable legal identity.

Under the Categorical Gap, these casts fail routinely. AI is simultaneously product, service, tool, and agent — and each classification activates a different legal universe. The consequences are not semantic but allocative. Classification determines who pays when a model hallucinates a diagnosis, misroutes an ambulance, or erases value from a balance sheet. As a product, it triggers strict liability; as a service, negligence; as an agent, responsibility shifts to the user; as speech, constitutional protections activate; as a tool, risk moves to the operator. The category is the business model.

This is why global firms engage in definitional rather than merely regulatory arbitrage. They are not shopping only for lenient regulators; they are shopping for more favourable meanings. The same model, deployed in different jurisdictions, can become a different legal object — a commercial tool in the U.S., a regulated risk in the EU, a strategic asset in China. The vocabulary is shared; the legal object is not.

The result is a linguistic fracture: the same word refers to different legal objects depending on the jurisdiction, the context, and the institutional interpreter. “Privacy,” “safety,” “harm,” and “data” are not universal concepts; they are local dialects. The State cannot fully reconcile these differences because each State treats its own grammar as the standard.

Into this fracture step the translators — the global ecosystem of elite law firms, accounting networks, technical compliance consultancies, and platform-governance teams. These actors are not merely advisors; they are implementers. They convert abstract law into operational reality. Fairness becomes bias testing; minimisation becomes data sharding; localisation becomes federated learning. They control the interpretive grammar, and in a world where the State’s own categories keep failing, they become the only actors capable of making the law execute at all.

B. The Five Functions of AI Legal Language

If law retains power in the AI economy, it derives not from what it commands but from what it classifies. Classification performs five core functions: it marks the limits of lawful action, assigns responsibility, defines ownership, creates visibility, and determines which sovereign can assert authority. Together, these functions form a Permission Framework — the grammar through which the State decides what may operate, on what terms, and at whose risk.

This is among the last acts of governance the State performs with relative confidence. Beyond it, the categories begin to fail — and the Translation Layer inherits what the State can no longer classify on its own.

The first function is permission — and permission is not the absence of prohibition. It is a legal state produced by classification. An actor is permitted only because the classifier has assigned it a recognisable legal type performing a recognisable legal act. If the classification fails, the action may occur factually but not legally — the system operates, but outside the boundaries of lawful recognition.

Jurisdictions apply different permission logics, each encoding a distinct theory of state authority. The U.S. operates on a residual basis: an activity is permitted unless specifically prohibited. The EU requires affirmative authorisation: an activity is prohibited unless it satisfies defined conditions. China applies a coordinated logic: an activity is permitted when aligned with state objectives. These are not different rules applied to a shared framework; they are different legal grammars.

The same model, performing the same function, can be lawful in one country and impermissible in another because the underlying permission logic diverges. Prohibition operates through the same classificatory machinery. When the EU AI Act prohibits social

scoring or real-time biometric identification in public spaces, it is not merely imposing a penalty; it is declaring the activity outside the set of legally recognisable acts.²¹

Between permission and prohibition lies a conditional zone: activities permitted only if specific requirements are met. This is where the Rule-Execution Gap widens most visibly. High-risk AI obligations under the EU AI Act — conformity assessments, technical documentation, human oversight, robustness thresholds — function as engineering specifications expressed in regulatory prose. Law can describe the requirement; it cannot execute it. Without translation, conditional permission becomes inoperable — formally available but practically unrealisable.²²

The second function is liability, and in the AI economy liability functions less as a moral judgement than as an allocation mechanism. Traditional liability rests on linear causation: drop a brick, break a toe, pay the bill. The doctrine presumes a person to blame, a duty to breach, a foreseeable harm, a traceable causal chain, and a mind capable of intending it.

AI strains every link in that chain. Duty becomes indeterminate — owed by the developer, the deployer, the data provider, or the platform, with no settled hierarchy. Breach loses its benchmark — the standard of care for a probabilistic model has no established precedent. Foreseeability collapses — if the model's creators cannot explain a hallucination, a court cannot assess whether the harm was predictable. Proximate cause fragments across the stack — the harm may originate in the prompt, the training data, the tuning process, or the deployment environment.

And mens rea evaporates entirely — an array of weights has no malice in any doctrinally cognisable sense. The doctrine does not break; it degrades into approximation. Responsibility does not disappear; it reroutes — upstream to operators and infrastructure holders, laterally to insurers and contractual counterparties, and downward to end users through terms of service. The result is not an absence of liability but a negotiation of it, mediated by the Translation Layer that absorbs the classificatory failures examined above.

The third function is ownership, and here the grammar strains most visibly. Property is not a relationship between a person and a thing. Property is a protocol — a system for defining exclusion, transfer, and control. In the industrial era, the objects of property were discrete: a factory, a book, a patent.

For AI, the object is a stack — weights, data, outputs, compute — each governed by a different legal dialect. Legal language does not recognise an AI model as a single ownable asset. It disassembles the stack into separate legal objects. The weights are governed through intellectual property, trade secret, or contract; the training data is fragmented across jurisdictions and rights regimes; the outputs resist clean ownership claims; the underlying algorithm is only partially protectable; and compute is treated as infrastructure rather than as a component of the model itself.

The result is a property regime in which the whole cannot be owned as a single legal object — only its fragments can be controlled, often by different parties under different legal systems. Markets require transferability, but the grammar of property cannot accommodate

²¹ Regulation (EU) 2024/1689, art. 5 (prohibited AI practices).

²² Regulation (EU) 2024/1689, arts. 9-15 (requirements for high-risk AI systems).

what it cannot cleanly name. A probability distribution cannot be liened. A continuously updating model resists foreclosure. Traditional security interests presume a stable asset with identifiable boundaries; AI satisfies neither condition.

Where formal legal categories fall short, the Translation Layer supplies its own: licensing frameworks, API-based access, and model-as-a-service arrangements that reproduce the economic effects of ownership without creating a legal property right. The gap between legal ownership and practical control becomes another domain in which the Translation Layer exercises constitutive authority.

The fourth function is disclosure — the mechanism through which AI becomes visible to capital markets. Markets cannot see code, weights, or gradients. They see disclosures — the legally mandated representations through which AI enters the market’s field of view. Disclosure does not illuminate; it constructs the market’s limited view of the system. Materiality thresholds, calibrated for deterministic assets, cannot capture the risk profile of probabilistic systems.

Firms respond by under-disclosing to limit fraud exposure or over-disclosing to inoculate against future claims — and the market ends up pricing the uncertainty of the disclosure regime as much as the system itself. The mismatch is foundational, not editorial. An SEC risk-factor disclosure may warn of “AI-related operational risks,” but it cannot express that the model’s failure rate is a moving target. A financial statement may recognise an AI-related intangible, but it cannot represent drift, context sensitivity, or emergent capability degradation.

The language of disclosure forces probabilistic reality into deterministic sentences — and that translation loss becomes part of the risk. The Translation Layer fills this interpretive space. Analysts, auditors, and risk consultancies become the interpreters through whom the market constructs its understanding of AI assets. Capital trades not on the model itself but on the representation produced by these intermediaries. The market does not price AI directly; it prices the translation — and whoever shapes the translation shapes the market’s perception of value.

The fifth function is sovereignty — and it is the one most visibly governed by the Translation Layer. Law operates in territorial space; AI operates across distributed computational infrastructure that respects no single boundary. When a model generates an output, it can trigger multiple jurisdictional claims simultaneously — server location, corporate domicile, data origin, data-subject residence, and regulatory effect — and these jurisdictional claims rarely align.

Each country asserts authority over some segment of the system’s lifecycle, yet none can observe or control the whole. Extraterritoriality becomes a form of grammatical reach: each legal system attempts to wrap global computation and machines intelligence in its own statutory language. The result is territorial conflict. U.S. AI disclosure frameworks, EU data-minimisation requirements, and China’s data-localisation mandates and outbound-transfer restrictions impose compliance instructions that cannot be satisfied simultaneously.

The Translation Layer engineers operational compatibility. Data sharding, federated learning, legal firewalls, and geofenced inference are not incidental workarounds; they are the organisational and technical mechanisms through which global AI systems function amid fragmented sovereignty. These tools do not resolve jurisdictional conflict; they route around

it — determining which sovereign commands are executed, which are emulated, and which are quietly bypassed.

Where permission, liability, ownership, and disclosure can each be handled within a single legal system, sovereignty is different: it demands reconciliation across multiple jurisdictions at once — and only the Translation Layer operates at that scale.

C. Limits of the Language of Law

The limits of legal power are not only political; they are linguistic. Law operates through a legacy representational system — open-textured principles such as fairness, transparency, accountability, and, as the EU AI Act illustrates, meaningful human oversight. AI systems operate through parameters, optimisation objectives, and high-dimensional representations. The result is a persistent translation problem: legal norms can guide AI, but they cannot be directly compiled into machine instructions without an intermediary. Normative language and computational form inhabit different representational grammars, and legislative precision can narrow, but not eliminate, that gap.

The EU AI Act illustrates the condition. Article 14 requires meaningful human oversight, including the ability for humans to understand relevant system behaviour, correctly interpret outputs, and override decisions in appropriate cases. These are significant commitments, but they remain under-specified at the point of execution.²³

The Act does not specify what a human must see to understand a large latent space, which interaction element constitutes a challenge, or what operational conditions make an override effective. The requirements assume a level of interpretability that the underlying systems — non-symbolic, high-dimensional, and continuously evolving — do not reliably provide.

This is the inherent limit of the language of law. A statute can articulate a normative objective; it cannot execute it. The Translation Layer fills this gap — converting legal requirements into operational proxies, determining which commands can be compiled into technical controls, which must be approximated, and which cannot be implemented at all.

The gap between normative language and technical execution produces a predictable institutional response. Firms translate oversight obligations into documentation, dashboards, and workflow gates that satisfy formal compliance without materially altering the underlying inference path. Early ISO/IEC 42001 audits may already show this pattern: lawyer–engineer negotiations resolved not through safety design but through interface-level substitutes — compliance artefacts that meet the standard while leaving the model’s dynamics substantially untouched.

The negotiation follows a recurring script. Consider a General Counsel and a Lead Engineer in a compliance session. The lawyer asks for meaningful human oversight; the engineer asks what the API specification for “meaningful” would be. The lawyer insists the human must be able to challenge the output; the engineer asks for the function signature. The lawyer finally says: build a button that says “I disagree.” The engineer asks what the button should do. The lawyer replies: make it compliant.

²³ Regulation (EU) 2024/1689, art. 14 (human oversight of high-risk AI systems).

By 5 PM, they have produced a blended template — a hybrid stitched from legal aspiration and technical constraint. “Meaningful oversight” becomes a prompt displaying the top three alternatives and a confidence slider. “Ability to challenge” becomes a counterfactual module. “Ability to override” becomes a manual review queue that rarely sees use.

The human is not in the loop; the human is an interface layer placed atop a process they cannot directly shape. Oversight becomes an interpretive surface rather than a point of control.

D. The Translation Imperative

Law can define categories, articulate principles, and assert authority — but norms do not execute themselves. Lessig observed that code operates as a regulatory instrument in its own right. The Translation Imperative inverts that insight: the question is not merely whether code constrains like law, but whether law can reach the operational layer at all without an intermediary capable of rendering its requirements in terms the system can enforce.²⁴

Before an AI system can act on a legal requirement, that requirement must be translated into operational controls. This is the foundational necessity of an intermediate layer — a translation architecture — capable of converting legal text into machine-readable constraints.

The conversion is specific and traceable. Privacy becomes encryption standards, access controls, and retention rules. Transparency becomes documentation protocols, inference logging, and audit interfaces. Safety becomes red-teaming regimes, robustness testing, and threshold filters. Accountability becomes audit trails and provenance tracking. Adequacy becomes data-residency requirements and cross-border transfer controls. In each case, an open-textured principle stated in legal grammar is rendered into engineering specifications that can be implemented, tested, and audited.

Legislatures generally cannot perform this conversion at the required level of technical specificity — not because they lack authority, but because the work requires forms of technical fluency and cross-jurisdictional observation that legislative institutions are not designed to provide. The Translation Layer is where this work occurs. It is not a convenience but a functional necessity: the only site in the governance chain where normative language and technical execution meet.

The belief that hyperscalers, frontier AI labs, and large technology firms can perform this translation internally — that enough lawyers and engineers can eliminate the need for an external interpretive layer — fails for three underlying reasons.

The first is epistemic. Internal teams see vertically: they understand their own systems in depth but lack the horizontal sightlines that translation requires — a cross-industry view of how dozens of organisations are interpreting the same clause, which voluntary frameworks regulators treat as de facto expectations, and what norms are emerging around audit depth, documentation granularity, and testing cadence. An internal team can benchmark itself against NIST or ISO. The Translation Layer sees how those benchmarks are being applied

²⁴ Lessig, *Code and Other Laws of Cyberspace*, (1999).

across the field, because it observes the convergence patterns that no single firm can access on its own.

The second is legitimacy. In regulatory environments, defensibility is strengthened by distance from the entity being assessed. Self-assessment plays a role, but its authority is limited when it rests solely on internal assertion. Regulators, insurers, counterparties, and auditors rely on convergence: standards gain weight when they are widely adopted, not when they are privately declared. The Translation Layer produces the templates, taxonomies, and compliance manuals that constitute this shared baseline — not because any one firm endorses them, but because they reflect collective practice.

The third is economic. Maintaining a live map of more than one hundred and fifty regulatory regimes is a fixed cost of unusual scale. For a single firm, it is a cost centre. For a professional services network, it is a product. The Translation Layer spreads this cost across thousands of clients, making compliance a shared infrastructure that the market sustains collectively because no single actor can sustain it alone.²⁵

The question then becomes: who performs this translation? The answer is the Translation Layer — the distributed ecosystem of legal, accounting, technical, and assurance institutions that operates between statutes and code. Its work spans several distinct functions. Engineering and platform compliance teams inside AI developers translate open-textured legal requirements into the control flows that models and infrastructure can enforce — “human oversight” becomes confidence thresholds, escalation logic, and structured testing pipelines.

Audit networks and adjacent assurance ecosystems stabilise expectations across the field, determining what counts as sufficient testing, what constitutes adequate documentation, and how evidence must be presented to regulators, insurers, and counterparties. Global law firms translate exposure, privilege, and regulatory posture into organisational and technical safeguards — “accountability” becomes disclosure regimes, review workflows, and liability-aware compliance frameworks.

The most consequential function is standard-setting through practice. The Translation Layer does more than convert individual requirements into individual controls; it determines which translations become the field norm. When an audit taxonomy is adopted across dozens of engagements, or a compliance template is replicated across jurisdictions, the Translation Layer is not merely implementing law — it is establishing the operational meaning of legal terms that statutes leave undefined. These are governance decisions expressed as institutional defaults.

E. The Adaptive Constitution: When Private Law Outpaces Public Law

A divide now runs through the legal system. Public law and private law operate on different clocks — and the divergence is constitutional in effect.

²⁵ OECD.AI Policy Observatory; IAPP, Global Privacy Law and DPA Tracker; DLA Piper, *Data Protection Laws of the World* (2024 ed.).

Public law legislates, regulates, and adjudicates on institutional time: slow, textual, and necessarily retrospective. This latency is not failure but design. Democratic legitimacy requires consultation, deliberation, and consensus — none of which can be compressed without weakening the authority they create. The EU AI Act took years to negotiate, during which the underlying technology passed through multiple capability generations. No democratic process of general applicability could have moved materially faster without altering the nature of the process itself.

Private law adapts at iteration speed. Each shift in AI capability — a new failure mode, a changed interface, a capability jump — triggers immediate contractual updates. Clause banks, indemnities, audit rights, data-lineage duties, and risk-allocation schedules revise at the same pace as model updates. When frontier models introduced reliable tool use, enterprise procurement terms adapted within weeks — not because private lawyers think faster, but because contracts are not subject to the same democratic process requirements.²⁶

The consequences compound. Because model builders and deployers hold real-time visibility into system behaviour, private instruments increasingly carry the day-to-day operational obligations that govern AI conduct: logging requirements, audit rights, safety thresholds, drift-detection triggers, and conditions under which models must be paused or retrained.

Public law can require that such controls exist. It rarely specifies them at operational granularity. The constraints that shape AI behaviour are often first articulated in contracts and platform terms, and only later — if at all — reflected in statute.

Until challenged, private ordering becomes the first effective governance layer in AI ecosystems. Litigation then acts as the ex-post interpreter through which private arrangements are translated into doctrine. Doctrine becomes downstream of contract rather than prior to it.

The adaptive centre of gravity shifts toward private law not because public institutions lack authority, but because their operating tempo cannot match AI's rate of change. The Translation Layer runs on private-law time — and as it does, the institutions best positioned to bridge the Rule-Execution Gap are those whose native instrument is contract, not statute. The institutional implications of this inversion, and the design required to restore public-law traction, are developed in Part V.²⁷

F. The Constitutional Consequence

The analysis yields a constitutional conclusion. The State supplies the grammar of authority — permission, liability, ownership, disclosure, sovereignty — but it cannot compile that grammar into operational controls at the speed and specificity that AI systems require. The Translation Layer performs this conversion, and in doing so shapes the conditions under which AI systems actually operate. Its templates, clause banks, audit methods, and evaluation protocols become the executable constraints that legal text alone cannot produce.

²⁶ Bütthe and Mattli, *The New Global Rulers* (2011); Halliday and Carruthers, *Bankrupt: Global Lawmaking and Systemic Financial Crises* (2009).

²⁷ Radin, *Boilerplate* (2013).

But law is only one of the languages through which AI becomes governable. It determines whether a system may operate lawfully; it does not determine whether that system becomes economically real. That work belongs to accounting — the State’s measurement language, built to track value, classify assets, and render activity legible to markets. The same institutions that translate “dignity” into bias tests also translate “model” into intangible asset, determining whether it is expensed, capitalised, or written off. This dual-language structure is the operating architecture of AI governance — and accounting is the second compiler.²⁸

II. Accounting As The Language Of Measurement

Accounting is the part of the governance stack most readers never see — yet it is one of the layers that most powerfully shapes what AI becomes in practice. Section I established law as the language of authority — the system that classifies actors, assigns responsibility, and sets the boundaries of AI action. Section II turns to the second compiler: accounting, the language through which AI becomes economically real.

This section is more technical than Section I because it addresses the machinery of valuation, classification, and measurement — the infrastructure that helps determine which AI systems are built, which are scaled, and which are abandoned. Without understanding this classification engine, the behaviour of AI firms and the economic outcomes they produce remains opaque. What follows is not an excursion into finance but an explanation of the incentive structures and constraints that shape the AI economy itself.

The distinction between law and accounting is not administrative. It is constitutive. Law determines whether an AI system may operate. Accounting determines what that system becomes once it does. Law grants permission. Accounting assigns economic identity.

Through accounting, an AI model is rendered into categories the economic system can process. It can be treated as an asset — capitalised, amortised, or collateralised. It can be treated as an expense that reduces current earnings. It can be treated as a liability that reflects future obligations. And beyond these formal categories, it can function as a source of economic value that shapes market capitalisation — a role that existing accounting standards do not fully capture or reflect. The choice among these treatments is not descriptive; it is constitutive. Each classification produces a different economic reality for the same underlying system.

The Industrial Logic and the AI Mismatch

The accounting standards that govern the global economy — U.S. Generally Accepted Accounting Principles (GAAP), International Financial Reporting Standards (IFRS), and Chinese Accounting Standards (CAS) — were built for the economics of the factory. They were designed to measure tangible inventory, physical plant, and linear financial instruments.

²⁸ The Dual-Compiler Thesis is an original contribution of this manuscript. R.L. Hale, 'Coercion and Distribution in a Supposedly Non-Coercive State,' 38 *Political Science Quarterly* 470–94 (1923); C.R. Sunstein, *After the Rights Revolution* (1990); A.G. Hopwood & P. Miller eds., *Accounting as Social and Institutional Practice* (1994); M. Power, *The Audit Society* (1997).

Their logic rests on implicit assumptions: that assets have stable and measurable value, that their components can be separated and individually assessed, that inputs yield predictable outputs, and that assets possess defined and exhaustible useful lives.

AI strains each of these assumptions. Models are entangled with the data, compute, and training dynamics that produced them. They generate probabilistic outputs rather than fixed values, making measurement possible but highly assumption-sensitive. They evolve through retraining, reinforcement, and iterative adaptation. And their competitive value can erode rapidly as new frontier systems emerge — a market dynamic that challenges traditional depreciation and impairment frameworks.²⁹

The resulting mismatch — the Measurement Gap — becomes visible once the five accounting operations are applied to AI systems in sequence.

The Rise of the Value Intermediaries

Just as with law, reconciling these divergences falls to the Translation Layer. As industrial-era accounting concepts strain under AI, the Big Four and specialist valuation firms expand from auditors into architects: they do not replace their traditional assurance role, but they increasingly shape the classifications through which probabilistic AI systems become legible as financial objects.

These firms are not the only participants in this process, but they are among the few with the institutional authority, methodological tooling, and professional legitimacy to produce classifications that are more likely to withstand regulatory, market, and audit scrutiny.

They do not merely interpret the rules. They translate probabilistic AI systems into deterministic balance sheets, and in doing so they operationalise the standards themselves. In the AI era, the Translation Layer becomes a de facto governance infrastructure for value — not by statute, but by ledgered classification.³⁰

A. The Mechanics of Recognition

The constitutive force of accounting is revealed not in the audit report but at the moment of recognition — the instant a physical expenditure is re-expressed as an economic object. Recognition does not mark the point where physics ends and economics begins; it is the point where a physical process becomes legible within the accounting system.

Consider a \$100 million training run. The physical sequence is unambiguous: GPUs convert electricity into heat; data moves across fibre networks; and hardware undergoes irreversible degradation. The physics are fixed. The economics are not.

The expenditure has no accounting meaning until a recognition rule is applied. Under different governing standards and factual contexts, the same industrial process can be

²⁹ Lev, *Intangibles* (2001); Lev & Gu, *The End of Accounting* (2016). IASB, *Discussion Paper: Goodwill and Impairment* (2020). The application to AI models whose capabilities evolve post-acquisition is this manuscript's extension.

³⁰ Ramanna, *Political Standards* (2015). See also Power, *Organized Uncertainty* (2007).

assigned to distinct economic categories — as an asset, as an expense, or, in a transaction setting, as a component of goodwill or an acquired intangible such as developed technology. The classification does not change the underlying physics, but it determines the economic world in which the expenditure will live.

The first treatment is expense — the erasure of value. When AI development is classified as research and development under ASC 730, the expenditure is expensed as incurred. Under IAS 38, development costs may be capitalised only when six demanding criteria are met — including technical feasibility, intention to complete, probable future economic benefits, and reliable measurement. In practice, much large-scale AI development does not satisfy these thresholds. Across both regimes, the prevailing outcome is that AI development costs are recognised as a current-period cost rather than an asset.³¹

Immediate expensing reduces net income and lowers taxable profit, but it also ensures that the resulting capability — the model, the infrastructure, the accumulated expertise — leaves no balance-sheet trace. The engineering reality compounds; the accounting record does not. Meta illustrates this dynamic. In its Form 10-K for the fiscal year ended December 31, 2025, Meta Platforms reported \$57.3 billion in company-wide research and development expenses, alongside \$19.2 billion in operating losses within its Reality Labs segment, which encompasses AR/VR and related AI platform work.³²

These investments — whether directed toward AI model development, AR/VR systems, or the infrastructure that supports them — are largely expensed under ASC 730. The result is a form of strategic invisibility: tens of billions of dollars in capability formation that never appear as assets, despite their centrality to the firm’s long-term trajectory. The investment occurs; the capability compounds; the asset does not.

The second treatment is capitalisation — the recognition of potential. If the project moves into a development phase and satisfies the recognition criteria under IAS 38, or the narrower capitalisation rules available under U.S. GAAP, the same expenditure may be recognised as an intangible asset rather than an immediate expense.³³

The cost shifts from the income statement to the balance sheet and is amortised over a designated useful life. The consequences are immediate: equity rises, common leverage ratios can improve, and borrowing capacity expands. What was once a reduction in current-period earnings becomes a long-lived asset that enhances the firm’s apparent financial strength.

The match between method and reality is fragile. Industrial-era amortisation schedules assume a useful life measured in years. Frontier models may lose competitive value on far shorter cycles as successor systems rapidly displace their predecessors. The accounting treatment is linear; the economic reality is often nonlinear. The result is a phantom asset — a balance-sheet object whose persistence reflects the standard more than the system it purports to represent.

³¹ FASB, ASC 730, Research and Development; IASB, IAS 38, Intangible Assets, paras. 54-67 (development-phase capitalisation criteria).

³² Meta Platforms, Inc., Annual Report (Form 10-K) for the fiscal year ended December 31, 2025.

³³ FASB, ASC 350-40, Internal-Use Software.

The third treatment is goodwill — the premium of uncertainty. If the model is acquired rather than built, the expenditure becomes subject to purchase-price allocation under ASC 805.³⁴

After all identifiable assets and liabilities are measured at fair value, any excess of the purchase price over the identifiable net assets is recognised as goodwill. Under both U.S. GAAP and IFRS, goodwill is not amortised. It is carried as an indefinite-lived balance-sheet asset, tested periodically for impairment. The mechanics matter: identifiable assets — whether tangible or intangible — are valued first; whatever cannot be separately recognised under the standards, including synergy expectations and other elements that fail the identifiability tests, is absorbed into the residual category of goodwill. The result is an asset whose magnitude reflects not only what is known but what cannot yet be known.

The scale of this residual can be substantial. Microsoft's acquisition of Activision Blizzard resulted in \$51 billion of goodwill allocated to the balance sheet from a total transaction cost of \$67.8 billion.³⁵

In AI-specific acquisitions, the challenge is often sharper: identifiable intangible assets may be harder to separate and measure, potentially increasing the proportion of consideration absorbed into goodwill. Uncertainty is not eliminated; it is formalised as an indefinite-lived balance-sheet asset.

Three treatments. One physical reality. Three economic classifications. The industrial process is identical. The economic reality is constructed by the classification applied.

The Limits of Industrial Logic

The three recognition treatments reveal a common structural limitation. Each attempts to contain AI within categories designed for deterministic, separable, depreciable assets. Expense erases value that compounds. Capitalisation imposes linear amortisation on value that decays nonlinearly. Goodwill formalises uncertainty as a residual rather than modelling it directly. These are not failures of application; they are the predictable consequences of a framework built for systems whose value is stable, isolable, and exhaustible.

AI systems operate under different informational and computational conditions — fluid, interdependent, and continuously reformed. Industrial accounting seeks to contain uncertainty through recognition and measurement rules; AI systems often operate through uncertainty as a functional property of their design.

The Translation Layer — Engineering Legibility

The recognition treatments do not resolve the mismatch; they translate it. A decaying, interdependent model is made to appear linear, separable, and durable not because the model changes but because the classification changes. The underlying system is unchanged — what shifts is the interpretive layer that renders it legible to the balance sheet. The Translation Layer constructs the representation that allows an unstable, rapidly evolving model to be

³⁴ FASB, ASC 805, Business Combinations; IASB, IFRS 3, Business Combinations.

³⁵ Microsoft Corporation, *Annual Report (Form 10-K)* for the fiscal year ended June 30, 2024, Note 8. The acquisition closed on 13 October 2023.

expressed within categories built for stable, exhaustible assets. It is not the model that adapts to accounting. It is accounting that adapts the model to itself.

B. The Recognition Selection

In quantum mechanics, a particle's state is unresolved until measured. In accounting, a frontier model's economic identity is similarly unresolved until recognised. Before recognition, the same system may be treated as a cost centre, a strategic asset, or a store of value depending on the governing framework and transaction context.

From that recognition choice follow distinct pathways under U.S. GAAP, IFRS, and Chinese CAS. Each framework applies different criteria for expensing, capitalisation, and the treatment of acquired intangibles. The result is not a single economic identity but a set of jurisdiction-specific representations of the same underlying model — three accounting regimes producing three different economic realities.

The divergence begins at the most fundamental level: the internal path, when a firm builds rather than buys. When a lab develops a model internally, it confronts the hardest hurdles in intangible recognition: separability, control, and reliable measurement. Under IAS 38 and ASC 730 — with a narrow capitalisation pathway available under ASC 350-40 for internal-use software — internally generated intangibles must satisfy stringent criteria before they can be recognised as assets.³⁶

A frontier model rarely meets these tests. It cannot be separated cleanly from the proprietary data, compute, and training infrastructure that produced it, and its technical feasibility cannot be demonstrated in isolation from those inputs. As the recognition mechanics in Section A demonstrated, the costs of internal model development are generally expensed rather than capitalised.

The consequences are immediate. No asset appears on the balance sheet. Tax treatment may accelerate deductions, but the benefit materialises only if the firm has sufficient taxable income to absorb them. Earnings appear suppressed and long-term investment appears penalised, because the accounting treatment forces the firm to recognise the expenditure upfront while the economic benefits emerge only over time.

The Mistral case illustrates the structural asymmetry. Mistral AI, operating under IAS 38, faces the same constraint as any European builder of frontier models: internally generated training costs encounter a high capitalisation bar. The model's components cannot be separated cleanly from the proprietary data, compute, and training infrastructure that produced them, and the economic benefits cannot be measured with the precision the standard demands. The result is a thin balance sheet relative to the scale of capability formation.

Now assume a U.S. company acquires Mistral at or near its most recently reported valuation of approximately €12 billion. Under ASC 805, the acquired business is remeasured at fair value: identifiable intangible assets are recognised separately where supportable, and any residual consideration is allocated to goodwill. The same underlying model that appeared as

³⁶ FASB, ASC 350-40, Intangibles – Goodwill and Other: Internal-Use Software.

an expense in the builder's accounts becomes part of the acquirer's asset base through purchase accounting. The mechanism is straightforward: acquisition bypasses the prohibitions on internally generated intangibles and triggers a step-up to fair value. Under these conditions, Europe funds the capability; the United States capitalises it.³⁷

This build-versus-buy asymmetry is not a technical curiosity. It reveals how recognition rules channel industrial outcomes. Under U.S. GAAP, business-combination accounting allows acquirers to recognise assets that builders were required to expense. Fair-value measurement and goodwill allocation turn external acquisition into a balance-sheet expansion strategy.

Under IFRS, internally developed AI capabilities face strict recognition hurdles under IAS 38. Most development costs remain off-balance-sheet until a qualifying transaction occurs. The result is a systematically thinner asset base for many European AI firms. Even when companies resist acquisition — as Mistral has thus far — the structural incentive remains: for a significant share of high-growth startups, migration or sale becomes a more viable monetisation path than remaining independent within the constraints of IFRS recognition.³⁸

Under China's CAS framework, recognition and disclosure operate within state-aligned financing structures. Policy banks, state-owned enterprises, and government procurement programmes absorb uncertainty that IFRS and GAAP push onto firm-level balance sheets. The economic representation of AI capability is therefore shaped not only by accounting rules but by the institutional framework that surrounds them.³⁹

Taken together, these regimes create a global arbitrage loop. Europe funds capability through expensed internal development; the U.S. capitalises capability through acquisition; and China embeds capability inside state-directed balance sheets. The result is a system in which innovation capacity, valuation, and strategic control migrate along the fault lines of accounting regimes rather than purely along lines of technological competence.

The question of who determines these classifications — whether an AI model is separable, feasible, or fairly valued — leads to the interpretive authority that governs the entire process. Engineers and regulators shape parts of the answer, but the decisive interpretive work is distributed across two groups: the Big Four accounting firms, which audit, advise, and influence the construction of accounting standards, and the specialist valuation firms engaged for purchase-price allocations and fair-value measurements. Together, they mediate how technical systems become legible within financial-reporting regimes.

Their influence begins upstream, in the formation of the rules themselves. The Big Four participate in debates at the IASB and FASB on separability, feasibility, and reliable measurement, helping define the judgement zones within which later decisions are made. Downstream, valuation specialists apply the instruments those standards authorise: fair-value

³⁷ Valuation figures for private companies are based on publicly reported funding-round disclosures, are inherently approximate, and may have changed since this manuscript's preparation.

³⁸ The build-versus-buy asymmetry as a structural bias favoring acquisition accounting over internal development is this manuscript's synthesis of IAS 38 and ASC 805. For broader context, see Lev, *Intangibles* (2001).

³⁹ Ministry of Finance of the PRC, Chinese Accounting Standards for Business Enterprises (as amended). The characterization of state-directed absorption of firm-level uncertainty is this manuscript's analysis.

hierarchies under IFRS 13 and ASC 820, relief-from-royalty models, multi-period excess-earnings methods, and cost-based approaches to internally developed intangibles.⁴⁰

These tools do more than measure. They shape outcomes. Scenario modelling, licensing hypotheticals, and probability-weighted cash flows influence whether an AI model appears accretive or burdensome, strategic or speculative. Purchase-price allocations determine how much value is assigned to goodwill versus identifiable intangibles, with material implications for future earnings, impairment risk, and the economic life assigned to frontier systems.

Within the boundaries of applicable standards — where professional judgement is permitted and outcomes may legitimately diverge — these actors perform a translation function central to how the governance stack operates. They shape which economic interpretations are available to the client and how technical systems are rendered into financial categories. They do not merely report what AI systems are worth; they materially influence how that worth is recognised, framed, and reflected in the financial system.

C. Accounting Authors Economic Reality

Law determines whether an AI system may enter the world; accounting determines what it becomes once it does. Law operates at the perimeter, deciding what may be admitted or excluded. Accounting operates at the centre, defining what can be built, scaled, and sustained within the economic system. Law grants permission. Accounting grants legibility and feasibility.

AI ecosystems cannot fully participate in the machinery of the State — taxation, finance, insurance, valuation, securitisation, disclosure — until they are rendered into accounting categories. They must be translated from probabilistic systems into defined economic objects.

This transformation is generative, not observational. Accounting does not merely discover economic facts; it helps constitute them. It does so through five irreducible operations that together determine how AI enters the economic system: recognition determines whether an item enters the accounting system; measurement establishes its magnitude; classification assigns its nature; disclosure shapes what becomes visible to markets and regulators; and consolidation defines the boundaries within which these elements are combined and reported. These are not administrative steps. They are the operating logic of economic life.⁴¹

The first operation is recognition — the binary gate that determines whether a phenomenon is admitted into the economic universe. The preceding section demonstrated this power: the same training expenditure can appear as expense, asset, or goodwill depending solely on which recognition rule fires. The accounting system does not merely record economic activity; it decides what counts as economic activity.

The deeper crisis of the AI era is that many of the most valuable components of AI ecosystems strain or fail the recognition tests that govern internally developed intangibles. Under the

⁴⁰ K. Ramanna, *Political Standards* (2015); M. Power, *The Audit Society* (1997).

⁴¹ P. Miller, in A.G. Hopwood & P. Miller eds., *Accounting as Social and Institutional Practice* (1994); M. Power, *The Audit Society* (1997); IASB, *Conceptual Framework for Financial Reporting* (2018), chs. 5–6. The five-operation framework as applied to AI systems is this manuscript's construction.

standards for internal development — IAS 38, ASC 730, and ASC 350-40 — an item enters the balance sheet only when it is identifiable, controlled, and capable of generating probable future economic benefits. Most internally generated AI components struggle to satisfy those requirements because their value is inseparable, emergent, or probabilistic in ways the current accounting system cannot classify.

Across the core elements of modern AI systems, the pattern repeats. Training compute functions as a development input but is typically expensed unless a narrow software-capitalisation rule applies. Data corpora operate as strategic moats yet rarely meet identifiability or reliable-measurement thresholds for separate recognition. Model weights — the central value driver — cannot be separated from the data, compute history, and training environment that produced them. Emergent capabilities confer competitive advantage but resist recognition because their value is probabilistic rather than contractually grounded.

And inference capacity, though revenue-generating, is treated as operating consumption rather than as a durable capability. The effect is a systematic erasure of the core. The balance sheet records the cost of the container — GPUs, cloud contracts, electricity — but treats the trained model itself, the accumulated capability, as an incidental byproduct. The model does not fail recognition because it lacks value. It fails because the current accounting system lacks a stable category for inseparable intangible value.

The second operation is measurement — the architect of magnitude. Measurement is not passive observation. It is mathematical compression: forcing a fluid, uncertain reality into a single scalar that accounting can recognise, compare, and report. In the industrial era, this compression was rarely contested. A factory had a cost. A machine had a salvage value. A bond had a yield. The underlying assets behaved in ways that were largely linear, and uncertainty could be contained within familiar estimation ranges.

AI does not fit that mode. These systems introduce two compounding layers of uncertainty: stochastic variability inherent in the model's own outputs, layered on top of epistemic uncertainty about the environment the model is attempting to represent. The result is a valuation object whose economic profile is neither stable nor easily parameterised. Training costs can vary by an order of magnitude depending on design, data quality, and optimisation strategy, yet accounting compresses that variance into a single historical figure.

Useful life, once a predictable function of physical wear, becomes a discontinuous curve: frontier models can be eclipsed overnight by algorithmic breakthroughs even as older models retain embedded or niche value. Emergent capabilities introduce nonlinear value interactions that traditional measurement frameworks treat as additive. And because there is no liquid market for frontier models, fair-value estimates rely on hypothetical exit scenarios rather than observable prices. Industrial accounting expects a ruler. AI hands it a probability cloud.

The measurement framework strains most visibly in the fair-value hierarchy. Under ASC 820, IFRS 13, and China's CAS 39, which adopts the same three-level structure, Level 1 inputs — quoted prices for identical assets in active markets — are almost never available for frontier models. Level 2 inputs offer only indirect analogues, requiring extensive adjustment and judgement. Level 3 inputs therefore become the default — and Level 3 is not valuation in the industrial sense. It is valuation by assumption. The hierarchy was built for assets whose economics could be triangulated from observable markets. AI offers no such anchor. The

deeper one descends into Level 3, the more the valuation becomes a constructed narrative rather than a discovered price.⁴²

The point-estimation problem sharpens the crisis. A model does not have a single value of \$50 million. It has a probability cloud: perhaps a ten per cent chance of being worth a billion dollars, a fifty per cent chance of being worth nothing, and an expected value somewhere in between. The balance sheet records one number and leaves the rest outside the frame.

Accounting collapses both layers of uncertainty — the variability inherent in the model's own outputs and the deeper uncertainty about the environment it represents — into a single scalar. The measurement operator fires. The distribution is flattened. The model is forced into one reporting state, not because economic reality demands it, but because the ledger cannot store superposition. The \$50 million figure is not a discovery. It is a negotiated estimate — a point chosen from a cloud because the system requires a point.

The third operation is classification — the architecture of capital. Once an object is recognised and measured, it must be sorted, and classification determines where the money goes. Capital expenditures are born onto the balance sheet and can be leveraged. Operating expenses hit earnings immediately, reducing taxes but disappearing from equity. Costs of goods sold are treated as inputs to production and compress gross margin. These categories are not neutral; they are the levers through which value is allocated, constrained, and redistributed.

The components already identified as straining recognition face a second structural constraint when they do enter the system. Their classification determines whether they build equity or compress earnings, and the default pathways consistently favour the latter. Investments in capability are routed into expense lines; revenue-enabling compute is routed into cost of goods sold; inseparable intellectual assets are routed into invisibility.

The system sorts AI not by economic function but by the legacy categories available to it. This produces a container bias: industrial accounting assumes value resides in the vessel, not the liquid. GPUs have mass and therefore enter property, plant, and equipment. Model weights have no mass and therefore fall into the category of internally generated intangibles — which means they are expensed. The map outranks the territory. The financial statements look efficient while the firm's actual intellectual capital base remains largely unbooked.

The result is the arbitrage of classification. Build internally and much of the cost is expensed: earnings before interest, taxes, depreciation, and amortisation (EBITDA) falls, and no asset is formed. Buy externally and the same capability enters through purchase accounting as an acquired intangible or goodwill, creating an instant asset base while leaving EBITDA untouched. Classification forces firms to choose between immediate earnings impact and balance-sheet strength — a design constraint imposed not by strategy but by the ledger itself.

The fourth operation is disclosure — the visibility operator that determines which phenomena enter the field of view of investors, regulators, and the State. Disclosure is not a window. It

⁴² FASB, ASC 820; IASB, IFRS 13; Ministry of Finance of the PRC, CAS 39. All three adopt comparable three-level hierarchies. The characterisation of Level 3 valuations as 'constructed narratives' in the AI context is this manuscript's framing.

is a spotlight, and the beam is shaped by materiality, not by physical magnitude. What is illuminated is what the system already knows how to see.⁴³

In AI, that spotlight falls unevenly. The current regime makes the container visible — the infrastructure, the inputs, the costs — while the content remains only partially illuminated. Firms may describe their training data in broad terms, but provenance, bias profile, and copyright exposure often remain opaque. Model architecture is typically treated as proprietary, which obscures the design constraints and scalability risks that determine long-term value.

And the economics of inference — the true marginal cost of deploying the model — are usually absorbed into aggregated operating lines, leaving investors with little sense of how capability translates into unit economics. The result is a disclosure field that captures the electricity bill and GPU depreciation but rarely reveals the internal strength or functional depth of the model itself.

This is a form of syntactic opacity. The accounting language lacks a precise category for a probabilistic reasoning engine. Recognition failure means the asset does not enter the system; disclosure failure means that even when related costs do enter, the system cannot adequately describe what they represent. The system can report risks, costs, and governance structures, but it cannot render the engine of value creation in a way that matches its economic significance.

The fifth operation is consolidation — and if recognition is the gatekeeper and measurement is the scale, consolidation is the wall. It defines the economic person: the entity that owns assets, bears liabilities, and interacts with the State. Consolidation is not arithmetic. It is constitutional. It determines who the firm is for legal, fiscal, and regulatory purposes.

In the industrial era, the boundary of the firm was physical. Ownership mapped cleanly onto factories, equipment, and subsidiaries. In the AI era, that boundary dissolves. The core of capability is distributed across cloud providers, data partners, and model-hosting environments that the firm may depend on operationally without ever owning.

Under IFRS 10, consolidation turns on power over relevant activities, exposure to variable returns, and the ability to affect those returns. Under ASC 810, a related but distinct framework applies, including the variable-interest-entity model.⁴⁴

AI strains both approaches by distributing economic substance across entities, jurisdictions, and infrastructure layers that the legal perimeter was never designed to contain. A firm may exercise operational control over a model through APIs, compute dependencies, and service-level agreements without owning the entity that hosts it; conversely, it may own a subsidiary while the value-generating assets reside elsewhere.

The doctrine of control presumes alignment between economic substance and legal boundary. AI destabilises that alignment by shifting substance into the network rather than the entity. The result is a structural hollow. The consolidated entity can appear legally coherent while

⁴³ S. Sunder, *Theory of Accounting and Control* (1997). The characterisation of disclosure as a visibility operator is this manuscript's framing.

⁴⁴ IASB, IFRS 10, Consolidated Financial Statements; FASB, ASC 810, Consolidation. The characterisation of consolidation as a 'constitutional' act defining the economic person is this manuscript's framing.

remaining economically incomplete. It reports revenue generated by capabilities it licences rather than controls. It carries balance-sheet assets whose functional substance depends on vendor infrastructure. It presents a full set of financial statements — revenue, cash flow, even capitalised intangibles — while the underlying technological capability sits outside the perimeter. The legal entity is intact; the economic reality is dispersed.

As AI models migrate value beyond the firm's legal boundary, consolidation becomes a forensic exercise. Auditors trace data and compute flows across jurisdictions. Regulators assess operational dependence and platform lock-in. Tax authorities map value creation to the data centres, training pipelines, and inference compute where the intelligence actually resides. The boundary of the firm no longer resembles a corporate chart; it resembles a network diagram.

This leads directly to a constitutive crisis. When the five accounting operations are applied to AI, they no longer align; they fracture the economic reality they are meant to represent. Recognition cannot capture the asset. Measurement cannot anchor its value. Classification routes intelligence into expense. Disclosure obscures the dependencies that matter. And consolidation assembles an entity whose substance resides outside its perimeter. The reporting system is formally complete yet economically unrecognisable.

The economic person constructed by the ledger no longer matches the intelligence deployed into the world. This is the constitutive limit of industrial accounting: it cannot enclose a decentralised, probabilistic, networked model within the rigid boundaries of a twentieth-century corporation. Once the logic of enclosure breaks down, the Translation Layer defines the boundary instead — and in defining it, determines what the State is able to perceive, tax, and regulate.

D. The Measurement Gap

The five accounting operations are not merely technical deficiencies within financial reporting. They cascade into a governance constraint: even with perfect sight and unquestioned authority, the State struggles at the point where command must become calculation. It can regulate AI behaviour, but it cannot reliably price its contribution.

This is the Measurement Gap — the rupture that forms when industrial-era fiscal architecture confronts probabilistic, self-updating capital. The State can govern a model's conduct, constrain its deployment, and shape its permissible uses, but it cannot consistently translate its evolving intelligence into a stable unit of account.⁴⁵

The fiscal system presumes assets with definable cost, predictable decay, and measurable output. AI strains those assumptions. Its costs are diffuse. Its useful life is contingent on retraining and drift. And its output emerges from distributions rather than points — from learning rather than depletion, from interactions that resist linear decomposition.

The result is a system that can command more easily than it can compute — a governance apparatus capable of directing the model yet unable to quantify with confidence the economic

⁴⁵ The Measurement Gap is an original concept introduced in this manuscript. B. Lev & F. Gu, *The End of Accounting* (2016).

force it exerts. The Measurement Gap is where authority meets uncertainty and loses its footing.

The first dimension of the gap is the collision between determinism and probability. As the five-operation analysis demonstrated, accounting compresses AI's probabilistic economics into deterministic point estimates. Fair-value regimes such as ASC 820 and IFRS 13 require a single reported figure even when the underlying economics are best described as a distribution with fat tails — sudden obsolescence, nonlinear upside, and performance drift driven by live data. Standards permit sensitivity analysis in the notes, but the primary statements remain scalar.⁴⁶

Once these compressed values enter the fiscal machinery, the distortion propagates. State metrics, tax bases, systemic-risk models, and capital-adequacy frameworks all inherit point estimates that were themselves the product of scalarisation. The result is a deterministic ledger imposed on a probabilistic economy — a system built to record fixed quantities attempting to capture assets whose value is contingent, contextual, and continuously reformed.

The second dimension is hyper-obsolescence. IAS 38 and ASC 350 require firms to estimate the period over which an intangible will generate benefits. Frontier models challenge this assumption: a model amortised over three years may become economically irrelevant in months. Straight-line and declining-balance methods impose continuous decay on assets that often behave more like binary options — generating value until they are abruptly overtaken. Impairment testing may not register this pace in time. The result is a balance sheet that can lag the technology cycle by an order of magnitude.⁴⁷

The sovereign inherits the illusion. Assets whose productive lives have effectively ended may continue to shape tax positions through depreciation schedules that no longer reflect economic reality. Carrying values can inflate collateral bases long after the underlying capability has expired. National accounts may incorporate book values that overstate the contribution of capital whose usefulness has already collapsed.

A linear measurement system confronting a nonlinear technology cycle overstates effective wealth — a distortion that travels from corporate reporting into the State's broader fiscal machinery. Industrial accounting assumes that time is linear, continuous, and predictable. AI operates on evolutionary time — long periods of stability punctuated by sudden extinction events.

The third dimension is emergence. Industrial accounting rests on methodological reductionism: the belief that value is the sum of its identifiable components. AI strains this logic. Its value is nonlinear, interaction-driven, and often materialises only at scale. Accounting's principal container for unexplained surplus is goodwill, and only in acquisitions.

⁴⁶ FASB, ASC 820, Fair Value Measurement; IASB, IFRS 13, Fair Value Measurement. A. Damodaran, *The Dark Side of Valuation* (3rd ed., 2018).

⁴⁷ IASB, IAS 38, Intangible Assets, paras. 88–96; FASB, ASC 350-30, General Intangibles Other Than Goodwill. The characterisation of frontier-model obsolescence cycles as outpacing standard amortisation periods is this manuscript's assessment.

Internally generated emergence is therefore rarely recognised as a separate asset. A firm may invest \$1 billion in compute and engineering and create economic potential that substantially exceeds the recorded cost. The balance sheet captures the expenditure; the surplus remains unclassified. The gap reflects a category mismatch, not merely a measurement error.⁴⁸

This gap extends beyond financial reporting into the State's capacity to perceive and capture value. When emergent capability expands faster than the ledger can register it, tax positions may be anchored to historical cost rather than current productive power. Competition authorities may underweight scale-dependent learning effects that generate durable advantage without corresponding increases in traditional market share.

National wealth accounts can reflect equity gains while missing the underlying capital formation that produced them. Emergent value accumulates in the economy but not in the machinery the State depends on to see it. The industrial ledger was designed to count bricks. It has no method for valuing the cathedral that emerges when they cohere.

The fourth dimension is the erosion of state capacity. The preceding subsections traced how recognition, measurement, and classification failures propagate from firm-level accounting into the State's fiscal machinery. Their cumulative effect is structural: a weakening of the informational foundations on which modern authority depends.

Supervisory regimes inherit this opacity most directly. Risk-weighted asset frameworks presume values that are stable, legible, and grounded in observable performance. AI-related assets often rest on assumption-heavy valuations whose economic durability is uncertain. When supervisors cannot reliably distinguish a robust capability from a latent failure mode, prudential oversight shifts from evidence to inference, and fragility can hide inside apparently compliant balance sheets. Attribution depends on measurement as well.⁴⁹

Tort and administrative liability require identifying who acted, how, and with what degree of control. Emergent behaviour resists clean decomposition, making responsibility harder to quantify and more contested across firms, regulators, and courts.

Macro-level steering confronts a parallel constraint. AI services may exhibit near-zero marginal cost, generate emergent value that never appears as capital, and absorb their output inside R&D expense. As a result, GDP, productivity, and inflation metrics become less informative about the structure of the economy they are meant to represent.⁵⁰

The fifth dimension is the void — and what fills it. The Measurement Gap creates a market for knowledge brokerage: specialists who translate between the probabilistic behaviour of AI systems and the deterministic grammar of financial reporting. Under ASC 820 and IFRS 13, frontier models often fall into Level 3, where observable market data are limited or absent and fair value depends on internal models, unobservable inputs, and management judgement.

In these conditions, valuation specialists and auditors become the primary translators of uncertainty into reportable numbers. Their role is not a monopoly over assets but an outsized

⁴⁸ W.B. Arthur, *Increasing Returns and Path Dependence in the Economy* (1994). The application to AI-system valuation is this manuscript's extension.

⁴⁹ R.S. Avi-Yonah, 'Globalisation, Tax Competition, and the Fiscal Crisis of the Welfare State,' 113 *Harv. L. Rev.* 1573 (2000). The extension to AI-specific measurement failures is this manuscript's analysis.

⁵⁰ J. Haskel & S. Westlake, *Capitalism without Capital* (2018).

influence over the judgement that determines how those assets appear in financial statements.⁵¹

The State shifts from measurer to mediator. It enforces process but cannot independently generate the numbers it enforces. When auditors adjust amortisation schedules for algorithmic obsolescence, national accounts may move. When consultants redefine useful life, the tax base can shift. Private methodologies begin to shape fiscal outcomes indirectly, through the practices they certify and the assumptions they normalise.

This is the culminating consequence of the Measurement Gap. When the State cannot independently generate the numbers on which its authority depends, the link between economic activity and institutional oversight weakens. The government sees the inputs and the outputs, but not the engine. A State that cannot measure value cannot fully govern value. Until measurement evolves to accommodate probabilistic, emergent, and discontinuous systems, the State risks governing a shadow — a representation of an economy that has already moved beyond the reach of the ledger.

E. Value Creation through Translation

Section D established that the Measurement Gap does not merely distort valuation — it creates a vacuum. That void is filled by accounting and audit networks, valuation specialists, and technical consultants who convert probabilistic AI systems into financial objects that institutions can act upon. Their work supplies the legibility that neither markets nor regulators can generate on their own.

In this environment, the Translation Layer functions as the mechanism through which AI's economic identity is constructed. The State retains the monopoly on law and enforcement, but increasingly depends on private intermediaries to surface AI-driven value creation and to signal its implications for economic development, financial stability, and systemic risk. This dependence is not a transfer of sovereignty but a shift in informational authority: where the State cannot measure, it must rely on those who can translate.

Because these translations operate through professional standards, valuation methodologies, and audit practices rather than democratic law, they function as a quiet but consequential source of economic power. The Translation Layer becomes the site where the boundaries of value, risk, and capital formation are first drawn — not discovered.

The mechanism through which this construction operates is procedural truth. As the preceding analysis demonstrated, AI systems resist measurement in the industrial sense — yet financial reporting, markets, and regulators all require a number. The Translation Layer supplies that number by constructing valuations that are not discovered facts but institutional outputs: estimates that satisfy the procedural demands of accounting even when the underlying asset resists precise description.

These outputs function as procedural truths. They are methodologically justified, auditor-defensible, and regulator-acceptable, and therefore become operationally binding. Their

⁵¹ M. Power, *Organised Uncertainty* (2007). The characterisation of valuation specialists as knowledge brokers is this manuscript's framing.

authority derives not from metaphysical accuracy but from conformity to standards, internal consistency, and the ability to withstand audit and regulatory scrutiny. A procedural truth is binding because institutions must act on it: it determines impairment charges, tax positions, capital allocation, and prudential ratios, regardless of the uncertainty that produced it.

Procedural truths also reveal the structure of economic authority. When GAAP and IFRS generate different valuations for the same model, the divergence is not an error but a demonstration that value is constructed through competing procedural regimes. Markets often treat these outputs as discovered facts — as if the number reveals something intrinsic about the asset — when in practice it reflects the assumptions, models, and methodological commitments embedded in the Translation Layer.

The number governs because it is accepted, not because it is inherently true. In this sense, accounting does not merely record AI value; it helps constitute the value that institutions can recognise, allocate, and regulate. Procedural truth is the mechanism through which probabilistic systems acquire economic identity, and through which uncertainty becomes actionable within the machinery of the State and the market.⁵²

The institutional space in which procedural truth operates has a specific location within the standards. As the fair-value analysis demonstrated, frontier AI models routinely fall into Level 3 of the hierarchy under ASC 820 and IFRS 13 — the level where unobservable inputs and internal models are explicitly permitted. What makes Level 3 distinctive is not merely that it accommodates uncertainty but that it authorises valuation through judgement when market evidence is limited or absent.

This authorisation is what enables the Translation Layer to operate. Level 3 provides the institutional footing on which procedural truths can be generated — valuations that are methodologically justified, auditor-defensible, and regulator-acceptable even when they remain approximations. Level 3 is not simply a category within the hierarchy; it is the zone in which AI's economic identity can be constructed in a manner that the State and the market are obliged to recognise.

The valuation tools deployed within this zone rely on accepted methods that render uncertain AI economics as auditable financial estimates. Whether through discounted cash flows, relief-from-royalty estimates, or cost-based reconstruction of inputs, the resulting figure is not an observation of underlying reality but a procedural output — a number validated by method, assumptions, and documentation rather than by market transaction. These techniques do not eliminate uncertainty; they formalise it.

Each approach embeds a different theory of value: income methods depend on projected cash flows, terminal-value assumptions, and the weighted-average cost of capital; relief-from-royalty blends benchmark royalty rates with projected revenue; cost approaches reconstruct the resources required to recreate or replace the system. What unites them is not their epistemic accuracy but their institutional acceptability.⁵³

⁵² P. Miller, in A.G. Hopwood & P. Miller eds., *Accounting as Social and Institutional Practice* (1994). The concept of procedural truth as applied to AI valuation is this manuscript's contribution.

⁵³ R.F. Reilly & R.P. Schweihs, *Valuing Intangible Assets* (1999); A. Damodaran, *The Dark Side of Valuation* (3rd ed., 2018). The observation that these methods formalise rather than eliminate uncertainty is this manuscript's framing.

In this regime, the criterion of success is not accuracy but defensibility. A valuation is acceptable if it can withstand regulatory challenge, satisfy the audit committee, and be disclosed in a footnote without triggering a restatement or qualified opinion. Professional services firms do not measure the value of AI models; they negotiate it.

Because the value of a probabilistic asset is often unobservable, the number that appears on the balance sheet is the negotiated equilibrium of management advocacy, auditor scrutiny, and methodological brokerage. The process converges through proposal, challenge, and synthesis — a structured bargaining cycle that yields a figure no party has incentive to contest.⁵⁴

Truth in this environment is procedural. A valuation is “true” if it is documented, consistent, comparable, and disclosed — not because it is numerically exact, but because it satisfies the evidentiary and methodological requirements of the standards. The balance sheet becomes a negotiated contract: in the AI economy, value is not measured; it is co-authored.

The Translation Layer controls this negotiation because it possesses a form of visibility that neither markets nor regulators can match. Regulators operate under ex-post disclosure; markets operate under voluntary signalling; the Translation Layer operates under direct engagement with the operational underlay — training curves, GPU logs, data-provenance chains, red-team failures, draft impairment memos, and cross-border IP structures.⁵⁵

This is not omniscience, but a structurally authorised vantage point produced by modern corporate governance. The power at issue is exclusive, institutionalised visibility.

Visibility becomes a mode of influence. Because the Translation Layer sees the internal mechanics of the AI ecosystem, it does not merely report on it; it shapes it. The State, unable to observe the computational interior directly, relies on the audit opinion as a proxy for economic truth. The Measurement Gap does not merely weaken state capacity; it reconfigures it. The State governs the representation framework, while the Translation Layer shapes the numbers through which the economy becomes legible.

AI capabilities are compounding faster than the State’s capacity to measure them, while the global accounting apparatus remains largely static. The distance between AI economic activity and the accounting vocabulary used to represent it is no longer a conceptual mismatch but a structural fault line.

Risk is mispriced because the instruments designed to detect it cannot represent probabilistic value. Capital is misallocated because the categories that direct it were built for separable, depreciable assets. And solvency increasingly depends on assumptions that no longer map to the underlying dynamics of the systems they purport to describe. The accounting system accumulates the very risks its own instruments cannot detect.

The Translation Layer is the primary channel through which AI value enters the financial system. It is under-scrutinised, structurally opaque, and privately controlled. This cannot be remedied with footnotes or incremental standard revisions. What is required is a post-

⁵⁴ K. Ramanna, *Political Standards* (2015). The application to AI valuation as a structured bargaining cycle is this manuscript's framing.

⁵⁵ M. Power, *Organised Uncertainty* (2007). The characterisation of the Translation Layer's vantage point as structurally authorised visibility is this manuscript's analysis.

industrial accounting framework capable of representing probabilistic, emergent, and discontinuous value — one in which translation becomes a public institutional function rather than a private compliance workaround. Until that framework exists, the State governs an economy it cannot fully see, through numbers it did not produce, on terms it did not set.

F. The Invisible Charter

Law and accounting were built to govern different objects — and for most of the industrial era, that division shaped not only what each discipline did but how each was taught, licensed, and practised. Law defined rights, obligations, and institutional authority. Accounting measured economic activity through valuation, disclosure, and financial representation. The lawyer asked: *what is permitted?* The accountant asked: *what is it worth?*⁵⁶

These were parallel inquiries, not intersecting ones — sustained by separate graduate faculties, separate licensing bodies, separate regulators, and separate theories of expertise.⁵⁶

The institutional separation ran deeper than subject matter. A lawyer advising on a corporate acquisition and an auditor certifying the same transaction's financial statements operated from different professional mandates, applied different professional standards, and owed duties to different constituencies.⁵⁷

They shared a client; they did not share a conceptual vocabulary. The lawyer interpreted contractual language against a body of precedent. The accountant interpreted economic substance against a body of measurement standards. Where their analyses touched — as they inevitably did in questions of asset classification, contingent liability, or regulatory compliance — the contact was mediated through the client, not through any shared analytical framework. Neither profession needed the other's answer to produce its own.

That separation held because the objects of governance stayed within their respective domains. Economic life was organised around human decision-makers, identifiable assets, and transactions whose legal character and economic substance could be determined independently. Law dealt with actors and their conduct. Accounting measured outcomes. Classification and measurement proceeded on parallel tracks because neither discipline's conclusions depended on the other's methodology.⁵⁸

AI breaks that separation.

The line between tool, asset, and agent blurs. A large language model fine-tuned for autonomous financial analysis is simultaneously a capital asset under IAS 38, a potential decision-maker whose outputs may trigger legal liability, and an operational system whose risk profile shifts with every retraining cycle. A single model can accumulate capabilities through fine-tuning, reinforcement learning, or retrieval-augmented generation and alter an organisation's economic position faster than the legal and accounting categories used to

⁵⁶ R.L. Abel, *American Lawyers* (1989); R. Stevens, *Law School* (1983). S.A. Zeff, 'How the U.S. Accounting Profession Got Where It Is Today,' 17 *Accounting Horizons* 189 (2003). This manuscript frames these as structurally parallel but non-intersecting credentialling systems.

⁵⁷ IAASB, ISA 200, Overall Objectives of the Independent Auditor, paras. 3–5.

⁵⁸ IASB, Conceptual Framework for Financial Reporting (2018), chs. 5–6. The observation that legal and accounting methodologies operated independently until AI forced their convergence is this manuscript's claim.

classify it. What functions as a tool at deployment becomes, through operational drift, something closer to an agent — and the reclassification problem is not sequential but simultaneous. The legal status and the economic measurement must change together, or both lose fidelity to the object they purport to describe.

This simultaneity forces convergence where none previously existed. To classify a model's legal status, the lawyer increasingly needs economic context — the model's position on the balance sheet, the capital committed to its development, the revenue streams it generates or displaces. To measure a model's economic value, the accountant increasingly needs legal and operational context — the regulatory regime governing its outputs, the liability exposure created by its deployment, the contractual frameworks within which it operates. Each domain now depends on the same underlying technical facts: model architecture, training regimes, update frequency, deployment environment, and patterns of autonomous behaviour. As those facts shift — and in AI systems, they shift between audit periods — both systems must shift with them or lose fidelity to the object they purport to describe.

Law and accounting therefore no longer operate as fully separate disciplines in the governance of AI systems. They function jointly, as a fused evaluative grammar, determining whether an AI system is recognised, how its value is constructed, and how its risks are allocated. The lawyer's question and the accountant's question have become the same question — and neither profession was trained to answer it alone.

Together they do not merely describe the AI economy. They help constitute it.⁵⁹

This is the invisible charter of the AI economy: a fused grammar of legitimacy and value that determines which systems may operate, how they are recognised, and on what terms they enter the economy. It is not yet a constitution — that architecture is developed in Part V — but it is the precondition for economic existence. Law sets the boundaries of authority. Accounting sets the boundaries of value. Only through both — operating together as a single evaluative language — does an AI system become economically real.

No university teaches this convergence as a unified field. No professional body governs it as such. No regulator oversees the joint product. The institutional vacancy is remarkable: the most consequential act of AI governance — the act that precedes regulation, precedes valuation, and precedes every market transaction involving an intelligent system — has no disciplinary home, no curriculum, and no licensing regime. The charter is invisible because neither profession recognises it as its own. But it governs nonetheless.

G. The Delegated Interpreter

Once law and accounting converge into a single evaluative system, a new problem appears: no institution is equipped to operate it alone.

The technical facts that now shape legal classification and economic measurement — model architectures, training regimes, update cycles, deployment patterns — are not directly

⁵⁹ J. Searle, *The Construction of Social Reality* (Free Press, 1995), chs. 1–2; N. MacCormick, *Institutions of Law* (Oxford University Press, 2007), ch. 2; M. Power, *The Audit Society* (Oxford University Press, 1997).

observable by the State. It cannot interpret their technical grammar. And it cannot value what it cannot observe.

Into this vacuum steps a familiar set of professionals — lawyers, auditors, valuation specialists, and technical consultants. None is new. But together they form the Translation Layer: the only collective with enough partial visibility into legal grammar, accounting standards, and the technical substrate to render AI economically legible. They lack formal regulatory power, yet they perform regulatory functions. They are not legislators, yet they shape the categories through which AI enters the economy. They are not economists, yet they materially influence what counts as capital.

Law creates the categories. Accounting creates the values. Interpretation constructs the reality.

These intermediaries have become the de facto executive arm of economic legibility in the AI era — not through enforcement, but through concentrated control over translation. Their collective function is not optional; it is required whenever legal classification, economic measurement, and technical interpretation must be resolved at the same time. They help decide whether a model is an asset or R&D, whether a hallucination is a defect or a disclosed limitation, and whether a system is legible enough for an audit opinion or too opaque to recognise on a balance sheet.

They are not consulting a map. They are drawing it.

The constitutional analysis of Part II is now complete. The evaluative system has converged, the interpretive function has migrated, and the Translation Layer has assumed practical control of economic legibility. But a system governed through private interpretation cannot remain stable across borders. Once interpretation becomes the mechanism of governance, divergence becomes the mechanism of geopolitics. That is where Part III begins.

PART III — THE ARCHITECTURE OF THE GLOBAL AI STACK

Protocol Note — v0.9 Release: Part III is presented in compressed form; full-length development of each component will appear in v1.0.

Part II established that law and accounting — the State’s dual governance languages — can no longer classify or measure AI with the precision their institutional authority requires. The Translation Layer has filled that gap by default. But the Translation Layer does not operate in isolation. It operates within a material structure — the AI governance stack — through which legal texts, economic grammars, institutional intermediaries, and computational infrastructure interact to produce operational outcomes.

Part III maps that structure. It traces the shift from horizontal power, organised through territory, to vertical power, organised through the stack — and examines how that vertical architecture concentrates around three blocs whose definitions of economic reality are diverging.

I. From Architecture to Implementation

The industrial world was organised horizontally. Power radiated outward from capitals toward borders. Jurisdiction was the unit of competition, and law and accounting were the paired languages through which states coordinated that competition.

AI shifts that geometry. A model can be trained in one country, deployed in another, and enforced through infrastructure that spans several more. A semiconductor does not register borders; a data centre routes workloads according to latency, not allegiance. A content-moderation rule written in Brussels shapes deployments in Oregon, Abu Dhabi, and Jakarta — its effects propagating vertically through the stack rather than horizontally across territory.

Authority now operates through this vertical architecture — the AI governance stack — a layered system in which legal texts, economic grammars, translation institutions, and computational infrastructure interact to produce operational results. Territorial law still matters, but its practical effect depends on how it is compiled into procedures, protocols, and technical controls.

Power no longer radiates outward; it descends. A requirement becomes a protocol, a protocol becomes a process, and a process becomes a constraint executed in silicon. In this architecture, authority is exercised not through enactment alone but through the design of the systems that implement it.

II. The AI Governance Machine

The stack has four layers. Law defines the permission language. Accounting defines the value grammar. Translation institutions turn both into protocols, templates, and controls. Compute executes the result in silicon.

In the industrial era, these layers were loosely coupled. In the AI era, they interlock. A change in one layer propagates through the others. The layers no longer operate as parallel domains; they function as sequential stages in a single compilation pathway.

A statute, standard, or policy does not bind AI directly. Its practical force emerges when language is compiled into structure — when a legal text becomes an audit protocol, a disclosure rule becomes a reporting requirement, and a liability regime becomes a runtime constraint. The transformation is compilational: ambiguity is progressively reduced into operational determinacy, even as each layer introduces its own discretionary choices. The overall movement is downward — toward fewer degrees of freedom.

Symbols become form, form becomes process, and process becomes constraints enforced by the machinery itself.

A conformity requirement under the EU AI Act illustrates the pathway. Statutory text becomes an audit checklist; the checklist becomes a compliance template produced by major assurance firms; the template becomes an automated documentation pipeline; the pipeline becomes a model-level control. The pattern generalises: regulation becomes decisive at the point of execution — when the constraint reaches silicon.

This is the AI governance stack — a vertical system in which law, accounting, institutions, and compute operate as a layered compilation process.

III. The Arrival of Agents

AI is shifting from systems that generate predictions to systems that initiate actions. This is the agentic shift — and it is determinative, because agency reallocates decision rights.

A predictive model can be governed through disclosure, licensing, and risk controls. An AI agent requires constraints embedded in the system's tools, permissions, and execution pathways — controls that operate at runtime rather than oversight applied after the fact.

As agents proliferate, models become actors, outputs become decisions, errors become harms, and capabilities become power. A demand-forecasting system becomes a supply-chain negotiator. A transaction classifier becomes an executor. A route optimiser becomes a logistics controller. These are not predictions for human review; they are actions that reshape economic life directly.

The agentic shift drives governance downward into the stack. Statutes cannot bind an agent at runtime. Effective control depends on the protocols, templates, and machine-level constraints that monitor and shape an agent's behaviour — regulation compiled into the same stack through which the agent operates.

IV. The Geopolitics of Vertical Integration

Vertical integration in AI produces a new geopolitical hierarchy — defined not by ideology but by control over the technical and institutional layers through which intelligence operates. A small number of states and firms now command the pathway from data extraction to model execution, from semiconductor fabrication to the rulemaking machinery that governs deployment. Two countries operate the most complete vertically integrated ecosystems at global scale: the U.S. and China. The EU occupies a structurally different position.⁶⁰

The U.S. leads through gravitational pull. Its ecosystem combines permissionless innovation, private-sector primacy, capital-driven scaling, and a Translation Layer largely outsourced to the market. The vertical stack spans domestic chip design, hyperscale compute, frontier model and agentic system development, and the deepest capital markets financing global AI diffusion. The U.S. does not govern AI through a single comprehensive regime; it shapes the global ecosystem through the structural dominance of its platforms, capital markets, and infrastructure — creating a default environment that foreign firms absorb simply by building on it.⁶¹

China leads through a different mechanism: policy coherence. Its ecosystem combines state-directed integration, vertically aligned industrial policy, and a regulatory framework fused

⁶⁰ U.S. CHIPS and Science Act (2022); EU Chips Act (2023). China's semiconductor self-sufficiency and national compute-network initiatives reflect parallel strategic responses.

⁶¹ Empirical data on U.S. cloud-market concentration, GPU design leadership, and AI venture-capital flows will be documented in the v1.0 edition.

with political authority. The vertical stack spans a growing domestic semiconductor supply chain, a national energy and compute network, state-aligned data regimes, and coordinated deployment pipelines. China aligns energy, technology, industry, and infrastructure into a single strategic direction — creating a system in which AI governance is not layered onto the economy but embedded within it.⁶²

The EU operates through neither gravitational pull nor policy coherence. It leads through rule projection. It does not operate a fully integrated AI stack; it operates a law-led ecosystem built from regulation outward rather than silicon downward. Its leverage derives from the ability to define the constraints under which AI built elsewhere must operate — through standards, liability regimes, and disclosure mandates that propagate extraterritorially across supply chains. This creates a distinctive structural position: significant normative authority that depends, in part, on continued access to the compute and models produced by the other two blocs.⁶³

Europe is not without sovereign compute ambitions. EuroHPC, national exascale programmes, and emerging public–private compute alliances represent material investments in domestic capacity and reduced dependence on external hyperscalers. Europe nonetheless remains less concentrated in scale compute and frontier model clusters than the United States or China — a position that continues to evolve as new investments come online. Europe’s comparative strength lies less in compute concentration than in regulatory projection and institutional coherence. The G3 divergence is one of structural orientation, not static deficit.⁶⁴

These are operational distinctions, not moral ones. Each bloc exercises authority through a different mode — gravitational pull, policy coherence, and rule projection — and each produces a different form of dependency for non-G3 countries. A country building on U.S. infrastructure absorbs market-driven norms. A country integrated into Chinese supply chains absorbs state-directed requirements. A country trading with the EU absorbs compliance obligations. No non-G3 country can engage with AI without navigating all three — and the terms of engagement are set by the blocs, not by the countries that depend on them.

V. The Vertical Geometry of Global Influence

The G3 do not merely compete over AI — they define the terms under which the rest of the world engages with it. Compute becomes a strategic asset, standards become instruments of projection, translation institutions become geopolitical actors, and the stack becomes the terrain on which global competition is conducted.

This creates a new dependency structure. Third-country firms building on U.S. or Chinese AI foundations inherit constraints they did not author — terms of access that can be redefined at any moment through export controls, sanctions, or platform-policy changes. Countries that

⁶² New Generation Artificial Intelligence Development Plan (2017); 15th Five-Year Plan for National Economic and Social Development (2026–2030), adopted March 2026. Subsequent compute-infrastructure and industrial-policy initiatives will be documented in the v1.0 edition.

⁶³ A. Bradford, *The Brussels Effect* (2020); Regulation (EU) 2024/1689 (EU AI Act).

⁶⁴ Council Regulation (EU) 2021/1173, establishing the European High Performance Computing Joint Undertaking (EuroHPC JU).

control none of the stack's critical layers rely on those who do. Territory still matters, but it no longer alone determines who sets the terms. The stack's control points do.

National autonomy in the AI era is not declared but constructed — encoded in technical standards, operationalised through private firms, and anchored in the computational physics that shape access and control. The question for every non-G3 country is not whether to engage with AI, but on whose terms — and at what cost to sovereign discretion. Any viable response must operate at the level of machine execution, where behaviour is actually produced, constrained, and steered. Part V returns to this challenge directly, asking whether a governance architecture adequate to these constraints can be built.

This diagnosis sets the stage for the fracture that follows: the point at which the G3 no longer share a common definition of economic reality.

VI. The G3 Fracture

Vertical AI power now intersects with the legal, accounting, and classificatory fractures examined in Part II. The U.S., China, and the EU no longer interpret AI through a shared conceptual frame. Each operates a distinct definition of economic reality, encoded in a different AI worldview.⁶⁵

The U.S. projects a GAAP-aligned, market-driven valuation framework: venture-capitalised, Level 3 fair-value oriented, and optimised for capital formation. Foreign firms building on U.S. AI foundations absorb this financial worldview by default — a translation that shapes how their models are priced, deployed, and scaled.

China embeds a CAS-aligned, state-directed industrial model: integrated into national planning and oriented toward long-horizon capability accumulation. Foreign firms leveraging Chinese AI supply chains face different requirements: localised data control, security obligations, and algorithmic-transparency standards that may conflict with commitments under other G3 regimes.

The EU codifies an IFRS-aligned, law-led governance regime: prudence-constrained, disclosure-intensive, and normatively anchored. Its compliance obligations reach foreign AI deployments extraterritorially, even as its regulatory authority rests on access to compute and models it does not produce.

The same model. The same weights. The same compute. Three distinct realities — because value, risk, and agency are stack-relative properties. A frontier model trained in California, fine-tuned in Shenzhen, and deployed in Frankfurt is not a single asset moving across borders. It becomes three different objects: a financial instrument, a strategic capability, and a regulated risk — depending on the interpretive frame through which it is read.

What is emerging is not a policy divergence but a fracture in the underlying logics that define value, risk, and agency. As these foundations separate, the consequences compound: supply

⁶⁵ The G3 fracture analysis is an original analytical contribution of this manuscript, building on the G2/G3 categories introduced in Part I.

chains that cannot be synchronised, liability regimes that resist reconciliation, measurement systems that refuse normalisation, and constraints on AI agents that defy interoperability.

These are not disagreements over regulatory detail. They are collisions between incompatible accounts of what AI is — three ecosystems offering different answers to what machine intelligence constitutes, how it should be evaluated, and who bears responsibility for the actions of AI agents. The fracture widens not because nations disagree on the need for AI governance, but because they begin from different assumptions about the nature of AI itself.

VII. The Offramp from Dependency

The fracture mapped in this Part forces a structural choice on every non-G3 jurisdiction. No country can engage with AI without inheriting dependencies — on U.S. platforms, Chinese supply chains, or EU compliance regimes, and often on all three simultaneously. The question is not whether to accept external constraints but how to manage them: align with one stack, borrow selectively from several, or construct an independent hybrid.

Part IV turns to the mechanisms that mediate across this divide. It begins with the institutional infrastructure of translation — examining Hong Kong as a reference implementation of how incompatible governance stacks currently communicate — and then develops the Incorporation Heuristic: an analytical model describing which rules propagate into global practice and why, as a function of visibility, workability, and necessity. Part IV also introduces the Sovereignty Paradox, the Gateway Rules framework, and the AI Power Hierarchy as structural consequences of translation under conditions of opacity and national competition. Part III identified where power resides in the AI stack. Part IV examines how that power is exercised — and where agency remains for countries navigating between the dominant blocs.

PART IV — THE GOVERNANCE ENGINE

Protocol Note — v0.9 Release: Part IV is presented in compressed form; full development will appear in v1.0.

Design, Propagation, and the Incorporation Heuristic

Part III mapped a world organized around three incompatible governance stacks — internally coherent, externally misaligned, and dependent on intermediaries to function across borders. Part IV turns to those intermediaries.

The Translation Layer is not a metaphor. It is an institutional, legal, financial, and operational architecture through which the G3 stacks make contact — the apparatus that determines which rules travel, which are filtered, and which are blocked at the border. This Part examines how that apparatus works, where it concentrates, and why its design determines who governs AI in practice.

The Translation Layer operates as a structured pipeline that converts legal directives into machine-executable form. The full sequence — from rule articulation through interpretation, specification, constraint encoding, system integration, operational enforcement, and evidentiary-artifact generation — varies in ordering across institutions, but the functional stages remain consistent.

In practice, the pipeline consolidates into three phases: interpretation, in which legal and regulatory language is converted into operational categories and compliance specifications; encoding, in which those specifications are embedded as constraints within systems, workflows, and technical architectures; and verification, in which resulting behavior is measured, documented, and attested through audit logs, system-generated activity traces, and exception reports. These phases frequently overlap or recur, but together they describe the dominant structure through which rules become operational.

Law firms anchor the interpretive phase, converting statutory obligations into the categories that organizes downstream compliance. Consultancies and technical integrators execute the encoding phase, embedding constraints into the systems through which AI operates. Accounting networks provide the measurement and recognition frameworks that determine what counts as compliant behavior and how it is reported. Together, these actors form a continuous pipeline — from rule to constraint to evidence — producing the operational records on which governance visibility depends.

I. Why the World Needs a Place Where Systems Can Talk

For most of the twentieth century, translation between governance systems was informal. London handled finance, Geneva handled diplomacy, New York handled capital, Singapore handled trade — a stylized geography reflected in the historical scholarship on international financial centers and global command nodes. The system worked because the underlying infrastructures were similar enough that informal mediation could bridge them.⁶⁶

AI removes that buffer. A model trained in California can influence a factory in Guangdong, trigger a regulatory response in Brussels, and move a pension fund in Oslo — all within a single decision loop. When incompatible stacks are forced into continuous contact at machine speed, translation becomes infrastructure rather than improvisation.

The new scarcity is optionality. If the leading stacks cannot interoperate natively, every other jurisdiction is pushed toward compatibility with one or more of them. The world needs sites where multiple stacks can be observed, compared, priced, and reconciled without collapsing into any single one.

Hong Kong functions as a reference implementation of this capability — a jurisdiction positioned at the intersection of the U.S., Chinese, and IFRS-aligned European stacks while maintaining its own legal, financial, and regulatory architecture. Its evolving political conditions have placed pressure on the dual-system framework on which that intermediary role historically depended, but the underlying function — intermediation between incompatible legal and economic regimes — remains analytically distinct from the

⁶⁶ Y. Cassis, *Capitals of Capital* (2006); S. Sassen, *The Global City* (2001).

jurisdiction that currently performs it. Other countries could develop comparable capabilities; the conditions for doing so emerge from the analysis that follows.

II. The Architecture of the Translation Layer

Hong Kong's Translation Layer operates through three structural capabilities that interact as a system.

The first is legal duality. Hong Kong's common-law tier governs contracts, commercial reasoning, and dispute resolution. Its PRC-law tier governs constitutional authority and national security. The two legal grammars coexist with domain-specific boundaries rather than clean separation — enabling rights, obligations, and enforcement pathways to move across otherwise incompatible legal systems.⁶⁷

The second is tri-grammar financial mediation. HKFRS — Hong Kong's implementation of IFRS — functions as an operational bridge between the G3 accounting regimes. CAS industrial realism becomes legible to global capital; GAAP intangible logic becomes interpretable to Mainland regulators. Firms can present coherent financial identities across divergent accounting systems, reducing the translation loss that occurs when moving between incompatible measurement grammars.⁶⁸

The third is a buffered topology for data and compute. Hong Kong is a jurisdiction where Western cloud services, Mainland cloud services, and global routing paths converge under a legal regime that permits controlled interaction between data sovereignties. Model weights can move without raw data crossing borders; inference can run on Western compute against Chinese-origin data; and federated training can occur with mutual verification — subject to evolving export-control and licensing constraints.⁶⁹

This combination is difficult to replicate because the barriers are foundational rather than incremental. Singapore has deep administrative capacity but lacks Hong Kong's RMB-native financial infrastructure and reciprocal civil-judgment enforcement with Mainland courts. Dubai International Financial Centre (DIFC) can scale common-law adjudication but has no PRC-facing enforcement pathways or CAS–HKFRS accounting corridor. London retains IFRS/GAAP depth and dispute-resolution leadership but lacks offshore RMB infrastructure and PRC enforcement reciprocity. New York has unmatched GAAP capacity and capital-market scale but faces structural barriers — sanctions, national-security architecture, data-sovereignty constraints — to any PRC-facing integration. Shanghai and Shenzhen anchor CAS industrial logic but cannot host dual-grammar legal execution within the current constitutional framework.

These barriers are real, but not permanent. Network-level alternatives — decentralized identity systems, cross-border regulatory sandboxes, distributed compliance mechanisms — are emerging, but they remain instruments rather than autonomous governance systems.

⁶⁷ Arrangement on Reciprocal Recognition and Enforcement of Judgments in Civil and Commercial Matters (2019).

⁶⁸ HKICPA, IFRS Convergence Framework.

⁶⁹ AWS Hong Kong Region; Alibaba Cloud Hong Kong Region.

Substitution is possible in principle; in practice, it requires rebuilding an interface that Hong Kong already operates.

III. The Valuation Translator — How HKFRS Bridges GAAP and CAS

A GAAP investor reads a CAS-based AI firm as undercapitalized; a Mainland regulator reads a GAAP-narrated technology firm as over-financialized. The same enterprise becomes a different economic object depending on the accounting grammar applied — and AI-driven firms, whose value depends on data, software, and forward-looking development capacity, expose this incompatibility most acutely.

HKFRS — Hong Kong's implementation of IFRS — functions as a voltage converter between these systems. It renders CAS-grounded industrial performance legible to global capital, allows GAAP-familiar intangible narratives to be interpreted within a prudential framework, and enables fair-value logic to coexist with conservative balance-sheet discipline. Hong Kong completed full IFRS convergence in 2005; China introduced the modern Accounting Standards for Business Enterprises beginning in 2006 as part of a broader alignment with international standards. CAS remains distinct in several areas, but the reform created substantial technical compatibility — and with it, a narrow but durable reconciliation corridor.

Hong Kong's professional ecosystem operates within that corridor. Drawing on the tri-grammar mediation described above, it reconciles CAS-native financials into HKFRS without violating Mainland prudential expectations, while presenting disclosures in forms that GAAP-trained global investors can interpret without systematic mispricing. The result is a jurisdiction with a uniquely interoperable financial language — one that bridges measurement grammars rather than choosing between them.

Valuation, however, is only half the translation. Financial statements encode rights — and those rights must be executable across legal systems. That is the work of the legal translator.

IV. The Legal Translator — Where Rights and Obligations Compile Across the G2 Divide

Cross-border capital structures require a jurisdiction where common-law and Mainland administrative-law grammars can both function without displacing one another. Hong Kong provides that environment. Its legal order preserves private-law autonomy, commercial adjudication, and international arbitration under common law while remaining embedded within the PRC's sovereign constitutional structure under the Basic Law. The layers operate in parallel, creating a hybrid jurisdiction in which commercial rights formed under one grammar can interface with enforcement institutions governed by the other.

The enforcement mechanisms extend beyond court judgments. Arbitral awards issued through the Hong Kong International Arbitration Centre (HKIAC) can be enforced in Mainland courts under the 1999 Arrangement on Mutual Enforcement of Arbitral Awards and

its 2020 Supplementary Arrangement. Both corridors — judgment recognition and arbitral enforcement — contain statutory scope limitations and public-policy safeguards, but together they provide the cross-border enforcement framework on which the Translation Layer depends.⁷⁰

A simplified example illustrates the pathway. A U.S. venture fund invests in a Shenzhen robotics company through a Hong Kong holding structure; the shareholders' agreement is governed by Hong Kong law; disputes are resolved under HKIAC rules; and any resulting award is enforceable in Mainland courts under the cross-border arbitration arrangements. The investor operates within a common-law contractual framework while the operating company remains fully subject to PRC sovereign regulations.

Capital moves between the systems without forcing either to abandon its own legal grammar. The same logic applies to derivatives, structured products, and cross-border financing instruments — each traversing the corridor without requiring either legal order to rewrite its internal logic. Contracts and rights can cross the divide; the data and computational infrastructure they govern must be able to do the same. That is the work of the physical translator.

V. The Physical Translator — Hong Kong as Data Port in the AI Era

Earlier eras of globalization required translation between commercial systems — trade through logistics, contracts through jurisdiction, valuations through accounting. Data introduces a different constraint. The storage and processing of data are bound to specific hardware and legal regimes, yet AI development routinely requires computational workflows that span them. The major digital powers — the U.S. through export controls and sector-specific governance, China through the Cybersecurity Law, Data Security Law, and PIPL, and the European Union through the GDPR — operate data-sovereignty frameworks that are not designed to interoperate.⁷¹

Hong Kong functions as a physical interface where these frameworks converge without collapsing. Western and Mainland cloud providers maintain regional infrastructure in close proximity. Model weights and trained parameters can move even when the underlying datasets cannot. Inference workloads can run on external compute while sensitive training data remains within Mainland borders. Federated training architectures allow collaborative model development across borders with verification conducted through Hong Kong's legal and institutional framework. This topology is reinforced by ongoing infrastructure investment, including the AI Supercomputing Centre at Cyberport, and policy initiatives facilitating controlled cross-boundary data flows.

⁷⁰ Arrangement on Mutual Enforcement of Arbitral Awards (1999); Supplementary Arrangement (2020).

⁷¹ PRC Cybersecurity Law (2017); Data Security Law (2021); Personal Information Protection Law (2021); Regulation (EU) 2016/679 (GDPR); U.S. Export Control Reform Act (2018) and BIS semiconductor-export rules.

China is positioned to generate some of the world's richest industrial datasets; the U.S. continues to dominate global model-deployment ecosystems. Hong Kong sits between them as one of the few jurisdictions where both systems can interact without collapsing into either one's regulatory gravity.

A simplified example illustrates the pathway. A model trained on Mainland industrial data exports its trained weights to Hong Kong following security review under China's data-export framework. Those weights are then deployed on external compute for global inference — without the underlying raw data leaving the Mainland. The pathway remains subject to evolving export-control and sector-specific conditions, but it demonstrates how AI pipelines can satisfy multiple sovereignty regimes simultaneously.

This is the Data Port — an emerging institutional form in which cross-sovereign computation occurs within a jurisdiction capable of hosting it physically, legally, and infrastructurally. As AI systems require continuous interaction between incompatible data regimes, jurisdictions that can provide this buffered interface become critical nodes in the global AI economy.⁷²

VI. The Systemic Risk of Losing the Interface

The interface depends in part on both technological ecosystems continuing to treat Hong Kong as a distinct operational environment. A shift in classification — through export-control categories, cloud-deployment policies, capital-routing rules, sanctions frameworks, or audit-recognition regimes — can narrow this function even when no physical infrastructure changes. Disconnection in such systems is primarily logical rather than physical.

These classifications are embedded in the compliance systems of private actors: banks, exchanges, cloud providers, auditors, and index compilers. They operationalize risk through internal models, standardized templates, and automated methodologies. When those models shift, the interface contracts independently of formal political decisions.

A neutral scenario illustrates the mechanism. A global cloud provider updates its internal risk model and downgrades Hong Kong's interoperability profile. Cross-region replication is reduced; high-performance workloads face new restrictions. Enterprise compliance teams update templates governing Hong Kong-hosted infrastructure. Financial institutions adjust operational pathways. Auditors revisit exposure classifications. Index providers reconsider weighting methodologies. Capital-allocation systems begin to reprice the jurisdiction. No law changes. No infrastructure disappears. Yet a classification adjustment propagates through interdependent systems, tightening constraints incrementally.

If the interface degrades materially, the effects can cascade outward. China faces higher friction in accessing global risk capital. Global investors lose visibility into Mainland industrial performance. Emerging markets lose one of the few neutral pathways into both stacks. Cross-border M&A becomes more expensive as translation premiums accumulate.

⁷² PRC Measures for Security Assessment of Outbound Data Transfers (2022). Hong Kong SAR Government, AI Supercomputing Centre at Cyberport (policy announcement). The concept of the Data Port is this manuscript's contribution.

For jurisdictions that control neither stack, the interface preserves strategic optionality — the ability to operate across systems without forced alignment.

The Translation Layer preserves systemic coherence under divergence. But coherence does not determine governance. Interfaces enable systems to interact; they do not decide which rules ultimately govern those interactions. To understand why some sovereign rules propagate globally while others remain confined to their origins, the analysis turns from the mechanics of translation to the logic of incorporation.

VII. The Sovereignty Paradox: Maps versus Networks

Legal sovereignty and operational control no longer coincide. Every state retains the formal authority to regulate AI, but the systems it seeks to govern operate across the globe simultaneously, following a network architecture that no single country controls. In this gap, authority migrates toward the operational compliance frameworks maintained by global firms — composite rule-sets embedded in the internal templates of hyperscalers, financial institutions, and multinational platforms, continuously updated to reconcile hundreds of overlapping regulatory regimes into a single executable system.

A routine AI-mediated commercial system makes the divergence concrete. A model trained in Seattle on data drawn from dozens of countries is executed on weights distributed across cloud regions in Oregon, Ireland, and Singapore; invoked by a customer in Lagos; routed through an API gateway in Frankfurt; and used to determine whether a Brazilian manufacturer receives financing. Multiple jurisdictions can plausibly assert authority, yet none — by itself — governs the system in full. The system’s behavior follows its network construct rather than any single territorial perimeter.

This dynamic becomes clearer when examining how global compliance systems absorb — or ignore — regulatory demands. A global privacy template illustrates the logic. It incorporates the GDPR because enforcement visibility and penalties are high; the California Consumer Privacy Act because the U.S. market cannot be ignored; and the PIPL because access to China’s ecosystem requires it. But it omits the data-protection frameworks of more than one hundred other nations. Each rule enters the template for a different reason — enforcement visibility, operational feasibility, or market gravity — but the logic is consistent: a rule is incorporated only when the cost of ignoring it exceeds the cost of implementing it across global systems.

The result is structural asymmetry. A citizen in South Africa or Indonesia is effectively governed by a hybrid of European, American, and Chinese rule-sets — the gravitational centers whose norms propagate through platforms, markets, and supply chains. Their own parliament may pass an AI Act, but unless that Act is absorbed into the global compliance template, it remains symbolic sovereignty: law that exists on paper but never reaches operational force.

This is the Sovereignty Paradox. The political map recognizes nearly two hundred sovereign nations. The AI economy organizes itself around a far smaller number of regulatory gravity centers whose rules propagate unevenly through global systems. The difference is not rhetorical; it reflects structural properties that determine whether a rule becomes visible,

executable, and unavoidable. The Incorporation Heuristic formalizes those properties — explaining why certain rules enter the global runtime and why most do not.

VIII. The Incorporation Heuristic

The System-Selection Logic of Global Power

The Translation Layer describes how rules travel across incompatible systems. The Incorporation Heuristic explains which rules survive the journey.

The heuristic can be formalized as a decision function. Incorporation likelihood is a function of three variables — Visibility, Workability, and Necessity — each exerting directional, threshold, and interaction effects. Visibility increases incorporation probability because rules cannot operate on activity the State cannot reliably detect. Workability increases it because rules that cannot be operationalized cannot take effect. Necessity increases it because rules that carry no meaningful consequence for omission are unlikely to be incorporated even when visible and workable.⁷³

Each variable functions as a gate: below a practical threshold of visibility, workability, or necessity, incorporation becomes unlikely. The variables also interact: visibility amplifies the effect of workability; necessity increases the incentive to improve visibility; and high necessity can drive redesign to make otherwise unworkable rules workable. The Incorporation Heuristic is not a descriptive law but an analytical model of how rules compete for incorporation in a dynamic, multi-jurisdictional environment.

This competition operates beneath the diplomatic surface of the international system. It emerges not through treaties but through the internal compliance engineering of global firms — the operational triage performed by lawyers, auditors, and risk architects whose mandate is to minimise exposure while maintaining functionality across the widest possible set of markets. The heuristic formalizes the logic through which these compliance architects determine whether a national rule must be compiled into the global baseline or can remain locally bounded.

For any rule that enters the cross-border compliance environment, incorporation likelihood can be expressed as:

$$I = V \times W \times N$$

The multiplicative structure is deliberate. A national rule enters the global template only when it satisfies three conditions simultaneously. It must be visible — detectable by institutions managing cross-border liability, with credible enforcement behind the detection. It must be workable — capable of being operationalized as a procedure across heterogeneous regimes. And it must be necessary — such that omission would trigger unacceptable legal or market risk.

⁷³ The Incorporation Heuristic — including the multiplicative structure $I = V \times W \times N$ and the sequential-gate interpretation — is an original formulation of this manuscript, drawing on traditions in systems reliability, institutional decision theory, and global compliance practice.

The variables are continuous, not binary. Visibility ranges from near zero — where regulators lack monitoring capacity or credible sanction authority — to near one, where enforcement is instrumented, specialized, and existentially credible. Visibility itself decomposes into a paired subsystem:

$$V = f(\text{detection, sanction})$$

Detection without sanction produces weak signals; sanction without detection produces random ones. Visibility registers only when both operate together with sufficient credibility to influence institutional risk calculations. Workability and necessity are similarly graduated: rules may be fully compiled into global templates, partially parameterized for specific markets, or isolated within local sandboxes.

Because the variables are multiplicative, failure at any stage is dispositive. A visible but unworkable rule stalls at the execution gate — for example, when a jurisdiction can detect non-compliance but lacks the technical specification needed to tell firms what compliance requires. A visible and workable but dispensable rule fails at the necessity gate. Only rules that remain non-zero across all three variables propagate into global application.

Other factors—enforcement severity, reputational spillovers, geopolitical alignment — operate by modifying one of the three variables rather than constituting independent considerations. Enforcement severity amplifies visibility by raising expected penalties. Reputational spillovers increase necessity by adding non-market exit costs. Geopolitical alignment can shift all three simultaneously by altering coordination expectations. The three-variable formulation captures the operative variance in observed compliance practice while remaining analytically tractable.

The heuristic does not claim that public authority disappears. States continue to legislate, regulate, and enforce within their territories. What it formalizes is the mechanism through which certain national rules extend beyond those territories and become embedded in the operational templates that govern cross-border AI. Regulatory power is exercised not only through formal jurisdiction but through the probability that a rule becomes incorporable — and the Incorporation Heuristic describes the arithmetic of that probability.

A. Visibility (V): The Detection Test

Visibility measures whether a rule’s violation will be detected and whether detection will trigger consequences that alter incentives. The first question a compliance architect asks is not *“Is this illegal?”* but *“If this rule is broken, will anyone see it — and will it matter?”*

Many countries have enacted AI-specific legislation — prohibiting algorithmic discrimination, mandating explainability, requiring human oversight — without building the capacity to observe violations. Where regulators lack model-audit capability, access to inference logs, inspectors with machine-learning expertise, and enforcement budgets sufficient to sustain investigations, the rules exist formally but produce no credible compliance risk. To the compliance architect, such rules are effectively invisible, and invisible rules are unlikely to enter global templates.

Contrast this with regulatory regimes that invest directly in observation. The EU’s Digital Services Act (DSA) combines standardized transparency reporting, vetted researcher access

to platform data, algorithmic-audit requirements, and systemic-risk assessment obligations — backed by fines reaching six percent of global turnover. The result is a regulatory environment where opacity is not a viable defense. Compliance architects cannot safely ignore rules enforced through this level of instrumentation. The telescope is pointed directly at them.⁷⁴

Visibility is not a function of sovereign size; it is a function of institutional instrumentation. Several smaller jurisdictions have demonstrated this. Singapore’s Monetary Authority of Singapore (MAS) deploys machine-readable reporting and conducts on-site technical supervision. Israel’s National Cyber Directorate (INCD) operates real-time cyber-telemetry pipelines. Norway’s Data Protection Authority (Datatilsynet) conducts high-resolution technical investigations. The UK’s Information Commissioner’s Office (ICO) maintains specialized algorithmic-auditing teams.⁷⁵

Individually, these efforts are modest. Collectively, these examples show that observation can be engineered and institutionalized — even by countries with comparatively small populations — when the mandate is clear and technical capacity is built deliberately.

Visibility is therefore the first gate. Only rules that produce credible detection signals enter the evaluation set. The remaining variables — workability and necessity — determine whether those observable rules can be operationalized and sustained within global compliance systems.

B. Workability (W): The Implementation Test

If visibility determines whether a rule enters the evaluation set, workability determines whether it can be operationalized. The second question a compliance architect asks is technical: *Can this rule be converted into machine-testable controls, repeatable audit procedures, and scalable compliance workflows?*

Many countries publish AI principles that are normatively meaningful but operationally indeterminate. A statute declaring that AI must be “fair,” “just,” or “aligned with cultural values” cannot be implemented as a control variable. A compliance system cannot audit “respect for human dignity” without defined metrics, nor evaluate cultural alignment without measurable criteria. When rules lack operational specificity, institutions respond predictably: the rule is documented, localized, and treated as a jurisdiction-specific exception rather than incorporated into the global template.

By contrast, rules with explicit technical thresholds translate directly into compliance systems. The U.S. Department of Commerce, Bureau of Industry and Security (BIS) advanced-computing export controls specify performance and interconnect-bandwidth thresholds for controlled semiconductors, requiring a license when those thresholds are exceeded. The compliance instruction is therefore relatively unambiguous and can be integrated directly into screening systems, workflows, and audit procedures, even though the

⁷⁴ Regulation (EU) 2022/2065 (Digital Services Act).

⁷⁵ MAS, Technology Risk Management Guidelines (revised 2021); INCD, Annual Report (2025); Datatilsynet, Investigation Report: Grindr Case (2021); ICO, AI Auditing Framework: Draft Guidance for Consultation (2020).

thresholds, licensing policies, and country frameworks have been revised repeatedly since the original 2022 rule.⁷⁶

As with visibility, workability does not depend on national size. It depends on whether legal text is drafted in a form that can be translated into operational controls. Japan's Personal Information Protection Commission (PPC) issues detailed guidance on cross-border data transfers under the APPI. Singapore's Personal Data Protection Commission (PDPC) provides structured guidance on data-protection compliance for AI systems. Switzerland's Financial Market Supervisory Authority (FINMA) mandates standardised liquidity-reporting requirements. Chile's Financial Market Commission (CMF) publishes open-banking standards. The United Arab Emirates' Telecommunications and Digital Government Regulatory Authority (TDRA) establishes defined information-security requirements. These jurisdictions exert regulatory influence not through economic scale but through drafting clarity. Their regulations can travel.⁷⁷

Workability is the second gate. A rule that is visible but cannot be operationalized stalls at the implementation stage. Only rules that pass both the detection test and the implementation test reach the final variable — necessity.

C. Necessity (N): The Market Gravity Test

If visibility determines whether a rule can be observed and workability determines whether it can be operationalized, necessity determines whether it must be obeyed. Necessity is economic rather than technical. It measures the cost of exiting a market relative to the cost of complying with its regulations — a market's exit elasticity. When exit is inexpensive, compliance treats a rule as optional; when exit is prohibitively costly, the rule becomes unavoidable.

Consider a small island state that mandates local data centers for all cloud providers. With a population of roughly 100,000 and a GDP near \$100 million, the cost of constructing even a modest Tier II facility — often \$12–15 million before operating expenses — would exceed any plausible revenue. The rule remains valid domestic law, but its incorporation weight within global compliance approaches zero. The outcome is not punitive; it is arithmetic.

The opposite dynamic emerges in markets whose scale or strategic position makes exit economically irrational. China's PIPL imposes data-localization requirements, restrictions on outbound transfers, granular consent obligations, mandatory security assessments, and requirements for local legal representation. These obligations are specific, operational, and enforceable — conditioning access to China's domestic AI ecosystem on compliance with an institutionally backed regulatory perimeter. For firms operating at global scale, abandoning

⁷⁶ U.S. Department of Commerce, Bureau of Industry and Security (BIS), Implementation of Additional Export Controls: Certain Advanced Computing and Semiconductor Manufacturing Items, 87 Fed. Reg. 62186 (2022) (subsequently revised).

⁷⁷ PPC, Guidelines on Cross-Border Transfer of Personal Data under the APPI (as amended 2022); PDPC, Advisory Guidelines on the Use of Personal Data in AI Recommendation and Decision Systems (2024); FINMA, Liquidity Risks — Banks (Circular 2015/2, as amended); CMF, Open Banking Standards (2022); TDRA, UAE Information Assurance Standards (2020).

the market would impose losses far exceeding the cost of implementation. Compliance teams absorb the costs and execute the rule-set.⁷⁸

The Necessity variable is not binary; it operates on a gradient shaped by enforcement credibility, penalty magnitude, and reputational exposure. Jurisdictions that combine high penalties with active enforcement generate stronger Necessity signals. Those with formal rules but weaker enforcement generate weaker ones. The following regimes illustrate the gradient: China's Personal Information Protection Law, India's Digital Personal Data Protection Act, Brazil's LGPD, Indonesia's Personal Data Protection Law, Nigeria's Data Protection Act, and California's consumer-privacy statutes. Each regime generates Necessity pressure, but the intensity depends on whether the enforcement machinery behind the rule is credible, resourced, and active.

Necessity emerges not from GDP alone but also from market gravity: the degree to which a market's economic opportunity or strategic position makes exit irrational. India's Digital Personal Data Protection Act (DPDPA) shapes access to one of the world's largest digital populations. Brazil's Lei Geral de Proteção de Dados (LGPD) influences compliance across Latin America's payments and fintech ecosystems. Indonesia's Personal Data Protection Law (PDP Law) governs a rapidly expanding cloud-services market. Nigeria's Data Protection Act (NDPA) affects the trajectory of African digital finance. California's Consumer Privacy Act (CCPA) and Privacy Rights Act (CPRA) shape national data-handling practices through regulatory spillover. These jurisdictions do not compel compliance through force alone; they compel it through significance.⁷⁹

Necessity is the third gate. A rule that is visible and workable but governs a market with low exit costs remains optional. Rules backed by sufficient market gravity are far more likely to become binding. Together, visibility, workability, and necessity determine which national rules enter the global compliance template and which remain confined to their originating jurisdiction.

D. The Result: The Framework Sort

Applying the Incorporation Heuristic across the world's regulatory environments produces a distribution that is directionally decisive. The pattern emerges not from political preference but from structural convergence: global firms gravitate toward rule-sets that minimise ambiguity, reduce operational friction, and limit exposure to enforcement risk.

The GDPR illustrates full incorporation. It scores highly on all three variables—strong enforcement visibility, machine-testable operational requirements, and substantial market-exit costs — and functions as a de facto baseline for many multinational firms. U.S. GAAP demonstrates the same dynamics in capital-market contexts: maximal enforcement

⁷⁸ Personal Information Protection Law of the People's Republic of China (2021).

⁷⁹ On the global proliferation of data-protection regimes generating market-gravity Necessity, see Digital Personal Data Protection Act (India, 2023), as implemented by subsequent rules; Lei Geral de Proteção de Dados (Brazil, 2018, as amended); Personal Data Protection Law (Indonesia, 2022); Nigeria Data Protection Act (2023); California Consumer Privacy Act (2018), as amended by the California Privacy Rights Act (2020).

visibility, precisely defined requirements, and indispensable market access produce dominant incorporation.

Singapore's Model AI Governance Framework shows a different trajectory. Its guidance is operationally clear, and tools such as AI Verify support implementation. Yet because the framework is voluntary and the domestic market comparatively small, necessity remains low. Singapore shapes global discourse more than it determines global practice.⁸⁰

Brazil's LGPD illustrates partial incorporation. Operational clarity and a strengthening enforcement regime give the LGPD meaningful regional influence, particularly across Latin America's digital and financial sectors. Yet Brazil's market gravity does not compel universal adoption. The result is selective incorporation by firms that handle Brazilian data.

Geopolitical constraints can override the heuristic entirely. Following comprehensive sanctions on Russia, more than a thousand multinational firms were tracked by Yale CELI as having curtailed or withdrawn operations, regardless of local regulatory quality. When firms are legally or politically compelled to withdraw, necessity collapses to zero and the heuristic loses force irrespective of visibility or workability.⁸¹

Across the global landscape, most national rules sit far from the upper bound of incorporability. The distribution reflects an evident reality: only rules that achieve sufficient visibility, operational clarity, and market gravity travel beyond their home jurisdiction. The Incorporation Heuristic does not evaluate the merits of any country's choices. It identifies which rules global firms actually embed in operations — and which remain legally valid but practically inert.

E. Raising Your Score — Strategies for Increasing Incorporation Weight

The three variables of the Incorporation Heuristic are not independent levers. Institutional changes that raise one score often impose costs on the others. But because incorporation weight is multiplicative, improvements in any single variable can materially increase the probability that a rule enters the global template. Across jurisdictions that succeed, five design strategies recur.

The first is instrumentation. Visibility increases when regulatory obligations become observable to external evaluators. Countries that invest in supervisory technology, standardized disclosures, monitoring pipelines, and machine-readable reporting create rules that external systems can detect. Instrumentation turns abstract obligations into operational signals.

The second is precision. Workability improves when statutes can be decomposed into specific checks, measurable actions, and processes that firms can implement at scale. Clear

⁸⁰ IMDA and PDPC, Model AI Governance Framework (2d ed. 2020); IMDA and PDPC, AI Verify Testing Framework and Toolkit (2022).

⁸¹ Yale Chief Executive Leadership Institute (CELI), List of Companies Leaving and Staying in Russia (maintained continuously since March 2022). The database grades more than 1,000 companies across a spectrum of disengagement and continued operation.

definitions, stable interfaces, and executable compliance steps allow external actors to translate legal requirements into operational controls.

The third is control over coordination bottlenecks. Necessity increases when a country controls a scarce coordination point — a licensing gate, a settlement platform, a data chokepoint, or a certification regime on which external actors depend. Market gravity is not solely a function of GDP; it is a function of indispensability.

The fourth is alignment with a legal center. Countries that anchor new rules in a dominant legal tradition — common law, civil law, or administrative law — reduce interpretive overhead and increase incorporability. Coherence lowers translation costs and allows external actors to apply familiar analytical tools rather than constructing new interpretive frameworks.

The fifth, and most consequential for smaller nations, is routing through a translation node: a hub whose legal, accounting, and professional-services ecosystem functions as a visibility amplifier, a workability converter, and a normative validator. The mechanism is concrete.

Vietnam's AI Law begins with a low incorporability profile. The statute exists only in Vietnamese and lacks an official English translation. its definitions are framed within domestic administrative concepts, and its terminology reflects Vietnam-specific regulatory priorities. Without authoritative translation, the law has limited legibility for global compliance systems.⁸²

Now imagine routing the same obligations through Hong Kong's professional-services ecosystem. None of the substance changes; what changes is the interface. Once rewritten in common-law language, cross-referenced to HKFRS-aligned disclosure requirements, and embedded in documentation that international evaluators already understand, the rule becomes detectable, interpretable, and structurally compatible with existing compliance pipelines.

Hong Kong often renders PRC obligations legible to global markets without altering their content — a pattern observed across finance, arbitration, and cross-border commercial law. The translation node does not replace domestic authority; it extends its operational reach.

These strategies do not transform Fiji into Brussels — and they do not need to. They move a nation from peripheral rule production to incorporated rule production: from rules that remain locally legible but externally inert to rules that external systems can interpret, execute, and rely upon. Incorporation is not granted by larger states. It is built — through institutional choices that allow domestic rules to enter the systems global firms depend on.

IX. The Power Hierarchy — Root Access, Admin Access, and User Access

Applied across the world's regulatory environments, the Incorporation Heuristic reveals a hierarchy of operational influence that does not map onto the political geography of sovereign equality. Every state remains sovereign in the formal sense; but within the stack that governs

⁸² Vietnam, Law on Artificial Intelligence (Law No. 134/2025/QH15, enacted 10 December 2025, effective 1 March 2026).

global AI, sovereignty behaves as a gradient of access privileges — a spectrum determining how far a state’s rules travel beyond its borders.

The gradient resolves into three tiers. Root-access jurisdictions write the defaults that propagate across multiple system-critical domains — data governance, capital markets, and frontier-model deployment. Their incorporation weight is highest globally, and their rules function as the baseline against which others are measured.

Admin-access jurisdictions shape or constrain system behaviour within particular sectors or at specific chokepoints. Their incorporation weight is above zero but domain-specific — a semiconductor licensing gate, a financial clearing corridor, a certification regime on which external actors depend. The Taiwan Semiconductor Manufacturing Company (TSMC) is not a geopolitical superpower, yet it is a single point of failure in advanced chip fabrication. Advanced Semiconductor Materials Lithography (ASML) is the sole supplier of extreme-ultraviolet lithography systems. Admin access is also a form of power.⁸³

User-access nations — the majority of the world’s states — retain full legal authority within their borders but lack the visibility, workability, or market gravity to project their rules into global AI systems. Their sovereignty is complete; their incorporation weight is low.

The AI economy therefore operates in two geometries simultaneously: the territorial map of formal authority and a permission structure layered on top of it, determining which rules can steer the global machinery of data, capital, and frontier models. The heuristic does not diminish sovereignty. It identifies which nations possess the operational permissions to shape AI governance beyond their own borders — and which do not.

A. The Three Global Centers (Root Access)

Three jurisdictions currently operate at Root level — each projecting rules along a distinct propagation axis.

The U.S. anchors the valuation core. Its influence projects through U.S. GAAP, SEC disclosure regimes, and the extraterritorial reach of export controls. AI firms seeking access to U.S. equity or debt markets must present growth, risk, and performance in terms legible to American capital markets, because the deepest capital pool sets the grammar in which value is measured. The propagation channel is necessity, reinforced by the workability of a detailed, rules-based system that global investors can reliably interpret.⁸⁴

China occupies the infrastructure core. Its influence channels through data-localisation mandates, cybersecurity classifications, and a state-defined network perimeter that governs what data may cross the border and what computation may occur within it. Foundation-model

⁸³ C. Miller, *Chip War* (2022); S.M. Khan, A. Mann & D. Peterson, 'The Semiconductor Supply Chain: Assessing National Competitiveness,' CSET (2021). The characterisation of TSMC and ASML as single points of failure is this manuscript's analytical description.

⁸⁴ SEC Form 20-F, Items 8, 17, and 18; SEC Release No. 33-8879 (2007).

licensing rules and inference-server controls extend this perimeter into AI development itself. Access requires accommodation; the alternative is exclusion.⁸⁵

The EU defines the standards core. Its influence is transmitted through the GDPR, the EU AI Act, and a procedural apparatus that treats classification, documentation, conformity assessment, and product design as sequential requirements. For frontier AI labs, exclusion from a 450-million-person and affluent market is rarely viable — and maintaining a single EU-compliant system is typically more efficient than engineering regulatory divergences.

Together, these three cores define the default operating environment of the AI-era global order. Valuation rules determine how firms measure success and allocate capital. Infrastructure constraints determine where AI can be deployed and where data may reside. Compliance regimes determine how deployment must be documented, evaluated, and justified. The global operating paradigm emerges from the interplay of these forces — not from territorial authority but from the structural conditions that determine which rules propagate.

B. The Power Users (Admin Access)

Admin-access countries do not usually set global defaults. They shape how those defaults are implemented, interpreted, or routed — reaching a localized peak of incorporability within a specific domain. In software terms, they resemble power users: they cannot rewrite the core, but they can configure how the system behaves. In global AI, the same logic applies. These states do not define the baseline rules, but they meaningfully shape how those rules operate in practice.

Three pathways tend to produce this form of influence.

The first is workability as power. Singapore's regulatory institutions do not author global standards, but they implement them with exceptional precision. MAS converts frameworks such as Basel III and Financial Action Task Force (FATF) recommendations into high-clarity, low-ambiguity guidance that multinational banks use as operational benchmarks. Singapore does not write the world's financial rules; it produces a highly executable version of them, and executability becomes a form of influence.

The first is workability as power. Singapore's regulatory institutions do not author global standards, but they implement them with exceptional precision. MAS converts frameworks such as Basel III and FATF recommendations into high-clarity, low-ambiguity guidance that multinational banks use as operational benchmarks. Singapore does not write the world's financial rules; it produces the most executable version of them — and executability becomes a form of influence.⁸⁶

The second is translation as power. The UK cannot match Brussels in regulatory scale or Silicon Valley in AI investment, but it retains a deep commercial-law tradition, centuries of

⁸⁵ CAC et al., Interim Measures for the Management of Generative Artificial Intelligence Services (2023); Cybersecurity Law of the People's Republic of China (2017); Data Security Law (2021); Personal Information Protection Law (2021).

⁸⁶ MAS Notice 637, Risk-Based Capital Adequacy Requirements (as revised); MAS Notice 626, Prevention of Money Laundering and Countering the Financing of Terrorism. FATF, Mutual Evaluation of Singapore (2016).

precedent, and the global language of contracts. English law remains the predominant governing law for International Swaps and Derivatives Association (ISDA) Master Agreements in cross-border transactions, and many commercial disputes are heard in London. The UK provides a stable legal interface through which global actors coordinate, reducing interpretive uncertainty between jurisdictions.⁸⁷

A Swiss variant exists: Switzerland's institutional neutrality, reinforced by the Bank for International Settlements (BIS) headquarters in Basel, offers a trusted venue for financial coordination where competing powers can cooperate without ceding political advantage. In this model, neutrality itself becomes infrastructure.⁸⁸

The third is necessity as power — influence created when a jurisdiction controls a function the global system cannot easily replace. TSMC in advanced semiconductor fabrication and ASML in extreme-ultraviolet lithography are the clearest examples; Chile and the Democratic Republic of the Congo occupy similar positions in critical minerals, Kenya in regional fiber-optic landing stations, and Iceland in stable, low-cost compute and energy.

These actors hold admin-level influence because they control non-substitutable capacity. But bottleneck power is unstable: concentration invites diversification. TSMC's dominance has already triggered fabrication initiatives under the U.S. CHIPS and Science Act, the EU Chips Act, and Japan's Rapidus program.⁸⁹

Admin access matters because most states will not reach Root level — but they need not remain at User level either. Meaningful influence does not require writing global rules. It requires controlling a domain where others depend on you: for operational clarity, for legal translation, or for irreplaceable capacity. In the AI economy, systemic criticality is not a function of GDP; it is a function of placement within the stack.

The remaining tier — User access — comprises the majority of the world's states. These countries retain full sovereign authority within their borders, and their legal systems govern domestic life completely. But their rules rarely propagate beyond their territory. They lack the enforcement instrumentation to generate visibility, the drafting precision to achieve workability, or the market gravity to compel necessity — and in the multiplicative logic of the heuristic, weakness in any single variable is sufficient to prevent incorporation.

These nations are locally sovereign but globally uncompiled. The Incorporation Heuristic does not diminish their authority. It identifies the structural conditions they would need to change — conditions the strategies for raising incorporation weight, examined earlier in this Part, are designed to address.

⁸⁷ ISDA, 2018 Choice of Court and Governing Law Guide; ISDA Master Agreement, Section 13 (governing law and jurisdiction).

⁸⁸ BIS, About the BIS (institutional overview).

⁸⁹ U.S. CHIPS and Science Act (2022); EU Chips Act (2023); Rapidus Corporation (est. 2022), a Japanese government-backed initiative targeting 2nm fabrication capacity.

C. The Permission Matrix — A Hierarchy That Is Real, but Not Closed

The hierarchy is real but not closed. Root access is rare and slow to form — the product of decades of institutional accumulation, market gravity, and positional centrality. Admin access is attainable through specialization, regulatory precision, or control of non-substitutable capacity. And User-access jurisdictions are not permanently confined to that tier. The strategies for raising incorporation weight — instrumentation, precision, bottleneck control, legal-tradition alignment, and routing through translation nodes — describe a structured pathway from local sovereignty to global compilation.

In AI governance, authority follows critical functions rather than territorial lines. Formal sovereignty remains universal, but operational influence concentrates around the actors who manage the permissions through which AI systems execute. The topology can be mapped; it can also be changed. That is the hierarchy's most consequential feature — not that it exists, but that it admits movement.

X. Gateway Rules — How Global Norms Move Across the World

Root access determines who writes rules that travel. A Gateway Rule is a regulatory requirement whose compliance forces changes to foundational design — data schemas, reporting formats, documentation pipelines, network configurations, or contractual defaults. Ordinary regulations function as local parameters, shaping conduct within the enacting jurisdiction. Gateway Rules alter the underlying frameworks on which other systems depend. When those frameworks shift, the effects radiate outward through supply chains, software stacks, and operational workflows, transforming practices in countries that never adopted the rule directly.⁹⁰

Gateway Rules travel through the same three channels mapped in the Root-access analysis. Reporting-format dependencies propagate through the U.S. valuation core, altering how systems measure. Network-layer dependencies move through the Chinese infrastructure core, altering what systems connect to. Compliance-process dependencies transmit through the European standards core, altering how systems demonstrate conformity.

Each channel operates on a different tier of the global system — measurement, infrastructure, or governance — and each generates a compliance cascade: a chain of secondary obligations that firms must satisfy together to maintain market access. Impact assessments, disclosure procedures, consent management, localization requirements, and audit documentation cluster into bundles that become embedded in product design. Once embedded, these bundles resist unwinding; removing one component often requires dismantling the entire compliance framework or exiting the market that created the rule.

⁹⁰ "Gateway Rules" is an original analytic category developed in this manuscript. It extends the Incorporation Heuristic by isolating the mechanism through which certain regulatory requirements propagate across borders via system design rather than formal national adoption.

Not all countries absorb Gateway Rules passively. Most adapt domestic systems to external dependencies, incorporating compliance cascades as they arrive. A smaller group takes a different approach — anticipating where new dependencies will land and building the institutional, technical, or legal frameworks through which those rules are executed. Singapore and the UK already occupy this position: they do not author Gateway Rules, but they control the channels through which Gateway Rules are implemented. India is building the institutional capacity to do the same.

The result is a governance framework in which a small number of nations shape the behavior of AI systems operating across sovereign borders — not through diplomatic consensus but through the technical and institutional patterns that global systems must adopt in order to function.

A. Why Gateway Rules Spread — The Economics of Convergent Design

Gateway Rules spread because large-scale AI systems converge toward minimal divergence. Uniformity is not a preference; it is the operational equilibrium that reduces maintenance overhead, testing permutations, and failure modes.

The GDPR’s Right to Be Forgotten is a useful illustration of the dynamic. To satisfy a single deletion request from a German user, a global AI system may need to purge personal data from active storage, scrub multi-year backups, notify downstream processors, reconcile the deletion with countries that mandate retention, remove references from indices and ranking systems, and generate an auditable record of the sequence. The exact steps vary by platform, but deletion in every case requires a system-wide engineering capability.⁹¹

Once that capability exists, the platform faces a design-level choice: maintain fragmented, country-specific branches or unify around a single framework calibrated to the strictest requirement. Fragmentation compounds operational complexity with each additional jurisdiction. Unification produces the opposite cost curve — once a capability is built, extending it across additional regions approaches near-zero marginal cost.

At scale, the gap becomes unsustainable. The deletion capability becomes a global feature. Users in Brazil, Indonesia, or India receive the functional equivalent of the Right to Be Forgotten — not because their legislatures enacted it, but because the platform will not maintain parallel systems. The EU legislated for 450 million; the effect extends to billions.

The pattern generalizes. Once components such as data-protection impact assessments, data-subject access portals, consent-tracking modules, or transfer-assessment workflows are embedded in cloud platforms and enterprise software, they become default modules in the global stack — and once embedded, they exhibit strong path dependence. Removing them is more expensive than maintaining them. Engineering changes undertaken to satisfy one

⁹¹ Regulation (EU) 2016/679, art. 17. First established in Case C-131/12, Google Spain (CJEU, 13 May 2014).

jurisdiction's rule become woven into global infrastructure, extending that rule's practical effects far beyond its formal intent and scope.

This dynamic creates a strategic opening for mid-power nations. Countries that understand where these compliance capabilities must run can position themselves along the pathways — building the services, standards, and institutions that convergent designs depend on. Influence need not come from writing the rule; it can come from becoming indispensable to its implementation.

B. The Golden Image — How Templates Spread the Rule

Gateway Rules spread not only through engineering design but through the global advisory ecosystem — the major law firms, audit networks, and multinational consultancies that build and maintain compliance systems for global firms.

The internal economics of these networks push toward standardization. Every deviation increases liability exposure, complicates audit defensibility, and raises delivery cost across distributed client portfolios. The result is what systems administrators call a Golden Image: a preconfigured baseline — master service agreements, audit protocols, data-protection addenda, cross-border transfer mechanisms — replicated across deployments to ensure consistency and reduce configuration risk. These templates anchor an organization's canonical compliance regime, the baseline against which all local deviations must be justified.⁹²

Embedded within these templates are controls drawn from multiple regulatory cores. U.S. anticorruption provisions from the Foreign Corrupt Practices Act sit alongside European privacy obligations from the GDPR, Chinese data-governance requirements from the Cybersecurity Law, and U.K. financial-crime controls shaped by the Bribery Act. The resulting image is composite, portable, and enforceable — assembled from the regulatory requirements of multiple jurisdictions.⁹³

When a bank in Nigeria hires a London law firm to update its cloud contracts, it receives this Golden Image — complete with FCPA clauses and GDPR-aligned processing terms it never explicitly requested. The controls arrive because the law firm deploys a common baseline across its portfolio. Removing them would require abandoning the template entirely — an option that increases legal cost, operational complexity, and professional liability.

The template persists because deviation is expensive. Once embedded in clause banks, contract libraries, audit frameworks, and regulatory checklists, these standards exhibit the same path dependence as the engineering capabilities described above. They travel with every contract, audit protocol, and governance framework deployed across multinational

⁹² The application of the Golden Image concept to global compliance template propagation is an original formulation in this manuscript. The term is borrowed from systems administration.

⁹³ U.S. Foreign Corrupt Practices Act (1977, as amended); EU General Data Protection Regulation (Regulation 2016/679); PRC Cybersecurity Law (2017); U.K. Bribery Act (2010).

systems.

C. Three Ways Rules Travel

Gateway Rules reshape systems through three distinct mechanisms, each operating at a different layer of the stack.

The first is measurement-format dependency. The converged revenue-recognition standards — ASC 606 under U.S. GAAP and IFRS 15 under international standards — require firms to record income when performance obligations are satisfied rather than when payments are received. Implementing this standard often forces redesign across sales-contract structures, commission frameworks, CRM workflows, billing logic, and financial-reporting pipelines — and in AI-enabled firms, it reshapes the enterprise data, labels, and model inputs used in forecasting, risk-scoring, and automated decision systems.⁹⁴

Once a parent company standardises these structures, subsidiaries inherit the same configurations, contract templates, reporting logic, and AI-model inputs. A pair of converged standards issued in Norwalk, Connecticut and London can therefore reshape sales, reporting, and prediction workflows across global organisations, even where the legal obligation originates in only two jurisdictions.

The second is network-layer dependency. China's data-governance framework requires certain categories of personal or sensitive data collected domestically to remain within national borders. For a global AI platform, compliance begins with the most expensive step: building a China-specific data stack — dedicated storage, separate compliance workflows, local operational teams, and independent audit controls.

That first partition often pushes the system toward federation. Once federation exists, the economics shift. Each subsequent localisation mandate — Russia, Saudi Arabia, Vietnam, Indonesia — arrives at dramatically lower marginal cost. The system has already been engineered for partition; adding a new country becomes an incremental extension. A structure that was once unified begins to resemble a constellation of regional enclaves.

The third is process-layer dependency. European AI regulation requires certain high-risk systems to undergo conformity assessment before market entry, including documentation of training-data provenance, testing procedures, and risk-management protocols. These obligations create what can be described analytically as build-time irreversibility: the relevant compliance record is most reliably generated during model construction, and the compliance window narrows once training is complete.⁹⁵

⁹⁴ FASB, ASC 606, Revenue from Contracts with Customers; IASB, IFRS 15. Both effective from 1 January 2018.

⁹⁵ "Build-time irreversibility" is an original analytic term introduced in this manuscript to describe compliance obligations that are most reliably satisfied during AI model construction and whose evidentiary assurance diminishes once training is complete.

Build-time irreversibility does not mean that models are fixed or behaviourally static after training. Post-training alignment techniques — including RLHF, constitutional tuning, and preference optimisation — can meaningfully steer surface-level behaviour, and inference-time controls can constrain or redirect outputs.

But these mechanisms operate on top of a pre-established representational structure. They shape expression more than provenance, and they modulate tendencies more than they reconstruct the training history embedded in the model. The foundational properties of a system — including the statistical imprint of its training corpus, the structure of its learned representations, and capabilities that emerge with scale — are set during the building phase and only partially alterable afterward.

The regulatory effect operates as a consequence of this structure. Because much of the required compliance documentation is most reliably generated during model construction, the rule shapes design long before a product reaches the EU and before any audit or supervisory exposure arises. A startup that implements data lineage from day one may be effectively aligning with the EU AI Act before it has a single European user, because retrofitting later often requires retraining models or rebuilding workflows. Though drafted for the European market, the rule's effects extend wherever standardised systems are deployed.

D. From Description to Design

The dynamic described in the preceding sections — where legal text becomes system constraint and then propagates as standardised infrastructure — is not only a feature of the current ecosystem. It is the mechanical pathway that any viable governance model must co-opt. The framework proposed in this manuscript relies on this same causal chain.

The sequence begins with the regulation itself, which defines the substantive obligation. It specifies the requirements that frontier-scale models must satisfy: auditability, provenance, interruption, network identity, and safety-case disclosure. These obligations exist prior to any technical implementation and function as legal and constitutional commitments.

The Gateway Rule converts the obligation into an systematic requirement, operationalising the regulation by making access to shared compute, credentialled identity, or interoperability conditional on compliance. The Gateway Rule does not create new duties; it enforces existing ones by tying them to the chokepoints through which high-risk systems must pass.

The Golden Image propagates the resulting compliance pattern — a standardised implementation template that satisfies the Gateway Rule and can be adopted, forked, or extended by operators. The Golden Image does not originate or enforce the obligations; it distributes them, ensuring that the same compliance logic is applied across multiple operators without bespoke engineering at each site.

The causal chain — Regulation → Gateway Rule → Golden Image — is the mechanism through which governance decisions made at the legal layer shape the long-term architecture of AI systems.

E. Strategies for Non-G3 Countries — Designing for Incorporation

For countries outside the G3, Gateway Rules define a design space rather than a constraint. The dominant regulatory cores set the system foundations — valuation frameworks, infrastructure boundaries, compliance procedures — but the cascade still depends on operational layers that no single core controls: audit capacity, verification infrastructure, interoperability mechanisms, and compliance tooling. Countries that position themselves within these pathways gain leverage. Three strategic patterns have emerged as especially salient.

The first is the early-adopter strategy. When a major jurisdiction signals that it will require algorithmic audits or system certification, a country can move early — enacting aligned requirements and building the institutional capacity to implement them before the mandate spreads. When it does, the early adopter is not absorbing compliance cost; it is supplying verification capacity.

Certification ecosystems exhibit economies of scale: once an audit framework exists, marginal assessments become cheaper than establishing new capacity elsewhere. The UAE's positioning around AI assurance frameworks reflects this dynamic. India's engagement with AI standard-setting through the Bureau of Indian Standards — including the work of the Artificial Intelligence Sectional Committee (LITD 30) and its adoption of ISO-aligned AI management-system and AI-application standards — illustrates the same trajectory. This strategy carries an inherent risk: credibility must match ambition. Early adoption without institutional depth erodes trust rather than building advantage.⁹⁶

The second is the safe-harbor strategy. When regulatory environments demand stronger verification of counterparties, countries can position themselves as trusted intermediaries — supplying registries, certification services, or governance frameworks that reduce compliance friction for global transactions. Financial markets offer precedent: jurisdictions such as Luxembourg and the Cayman Islands became structurally embedded as global financial hubs not by writing the rules, but by supplying the legal, custody, due-diligence, and intermediation structures those rules required.⁹⁷

Canada's emerging role in AI model evaluation and safety benchmarking, anchored in the Canadian Artificial Intelligence Safety Institute (CAISI), reflects the same logic —

⁹⁶ ISO/IEC 42001:2023, Artificial Intelligence Management System — Requirements. Bureau of Indian Standards, Artificial Intelligence Sectional Committee (LITD 30), including Indian Standards IS/ISO/IEC 42001:2023 and IS/ISO/IEC 5339:2024.

⁹⁷ R. Palan, R. Murphy & C. Chavagneux, *Tax Havens: How Globalization Really Works* (Cornell University Press, 2010).

credibility as a service. The risk is proportional to the role: verification hubs attract scrutiny, as financial hubs have learned through sustained OECD and FATF pressure. Their durability depends on whether the verification function provides enough systemic value to justify the attention.⁹⁸

The third is the federation strategy. As data-localisation requirements and infrastructure partitioning increase, global AI systems are becoming progressively federated — operating as multiple country-specific deployments that must interconnect. Countries that provide the legal, technical, or institutional capabilities enabling partitioned systems to communicate gain structural importance.

Hong Kong has historically performed this function for cross-boundary financial and data flows between mainland China and international markets. Singapore's trusted-data-intermediary frameworks pursue a similar objective. The risk is inherent to the role: interoperability layers add complexity, and their value depends on whether they reduce friction elsewhere in the system.

These strategies do not require contesting the dominant cores. They require identifying where the cascade depends on capabilities the cores do not supply — and building them. Rulemaking authority in AI governance is concentrated; operational influence is distributed. The gap between the two is where non-G3 countries find agency.

XI. Conclusion — The Invisible Constitution

A. The Age of Unseen Power

The pattern that emerges from Part IV is structural. Three regulatory cores write the defaults that govern global AI. Gateway Rules propagate those defaults through measurement formats, network architectures, and compliance procedures. The Golden Image distributes them through the translation ecosystems of Big Law, Big Four accounting networks, and major advisory firms. The result is an engineered order — a functional rule-set produced not by legislatures alone but by compliance engineers, cloud platforms, and the propagation logic of convergent design.

This is the Invisible Constitution: the operational rulebook that governs global AI regardless of what formal law provides. It is not fully democratic, not fully negotiated, and not equal in its distribution of influence. But it functions — and it functions at global scale, because the alternative — a fully fragmented regulatory landscape — would impose coordination costs that frontier AI systems cannot sustain.

The Invisible Constitution resolves a coordination challenge. It does not solve a legitimacy problem. Execution capacity has become a source of governing authority, but execution capacity is not the same as democratic mandate. That gap — between functional governance

⁹⁸ Government of Canada, Canadian Artificial Intelligence Safety Institute (CAISI).

and legitimate governance — is the subject of Part V.

B. The Strategic Design Space

For countries outside the G3, the Invisible Constitution is a constraint — but it is also a design space. The cores are already established, yet the global system depends on operational infrastructure that no single core supplies. Countries that build capacity in those layers become part of the implementation machinery through which AI governance actually operates.

The Incorporation Heuristic provides the design logic: build verifiable regulatory capacity, implement rules that integrate cleanly across systems, and position yourself where coordination becomes unavoidable. The preceding sections have shown what this looks like in practice — early adoption of assurance standards, construction of verification infrastructure, development of federation and interoperability capabilities, and alignment with the advisory templates through which global firms operate. These are not peripheral activities. They are the operational layers on which the Invisible Constitution depends.

Influence in this framework does not come from rewriting the cores. It comes from supplying components the system cannot operate without. For most nations, that is not a consolation. It is an invitation — one that requires institutional depth, technical capacity, and sustained investment, but one that remains structurally available to any country prepared to build what the system needs.

C. The Rules That Execute

Public authority in the AI era is shaped less by what is declared than by what is compiled into the systems through which global AI runs. The Invisible Constitution, the Incorporation Heuristic, the Gateway Rules, and the Golden Image are not abstractions. They are the machinery through which governance now functions — and they operate whether or not any legislature has endorsed them.

The Translation Layer cannot, on its own, constitute a governance order capable of managing the material foundations of AI. Its function is descriptive and mediating: it renders model behaviour legible to existing institutional forms, but it does not supply the institutional authority or enforcement capacity required to govern machine intelligence at scale.

Translation stabilises visibility; it does not generate the normative commitments, auditability conditions, or procedural guarantees that a full governance framework requires. And rules, however well they travel, govern behaviour — they do not directly allocate the compute, data, and models on which AI power ultimately depends, even where they shape access to those resources indirectly through procurement constraints or export controls.

Without deliberate State intervention — including the creation of an Audit Tier, the specification of ex ante obligations, and the establishment of enforceable oversight pathways — the Translation Layer remains a conduit rather than a constitutional foundation.

Part V moves from description to construction: from how rules travel to whether a governance architecture adequate to the material foundations of AI power can be built.

PART V — THE MATERIAL CONSTITUTION OF THE AI ERA

Protocol Note — v0.9 Release: Part V is presented in compressed form; full development will appear in v1.0.

Compute, Capacity, and the Constitutional Limits of Governance

Part IV mapped how rules travel and who shapes their movement. Part V shifts from the linguistics of governance to its physics — from the frameworks through which rules spread to the material foundations that determine what can be built, who can build it, and how far their authority can reach.

Traditional authority rests on rule primacy: whoever controls the rules controls the system. Frontier AI introduces a different structural logic — capacity primacy — in which control of computational capacity conditions which rules can be made effective and, increasingly, who can set the technical defaults that become embedded throughout the stack. A state without sufficient compute cannot train frontier models; a state that cannot train models cannot shape the defaults that propagate through the mechanisms described in Part IV.⁹⁹

This shift completes the manuscript's analytical arc. Law provides the language of permission. Accounting provides the grammar of value. The Incorporation Heuristic identifies which rules travel. Part V identifies the boundary condition — the physical capacity without which none of them execute. The framework is now complete: from grammar to design to physics.

⁹⁹ "Capacity primacy" is an original analytic term introduced in this manuscript to describe the structural condition in which control of computational resources conditions the effectiveness of legal and regulatory authority.

The Physics, Made Explicit

The physics of frontier AI are material, measurable, and increasingly constraining. Training a state-of-the-art model can require more than 100 megawatts of continuous power — the electrical load of a small city — with some forecasts projecting multi-gigawatt-scale clusters before the decade closes. Frontier-scale training runs can cost tens of millions to billions of dollars in compute alone.¹⁰⁰

The semiconductor supply chain that makes such runs possible passes through a small number of critical nodes worldwide, with leading-edge fabrication concentrated in a single cluster of facilities in Taiwan.¹⁰¹

The global talent capable of designing and operating these frontier models remains a narrow pool — a subset of the broader AI workforce with direct experience at frontier scale. And the capital required to sustain a frontier programme exceeds the annual R&D budgets of many nation-states.¹⁰²

This is not software. It is an industrial ecosystem — one whose resource requirements many countries cannot meet independently, and whose constraints increasingly define the outer limits of what governance frameworks can reach.

The Capacity Challenge — Especially Outside Major Powers

The capacity demands of frontier AI impose constraints that most states cannot overcome through regulation alone. Unlike earlier waves of the internet and digital infrastructure — which could be distributed and adopted by almost all countries — frontier AI depends on concentrated high-performance computation with material limits: training clusters, inference hosting, data-centre scale, and the capital required to build and maintain them. For countries outside the primary compute centres, the constraint is real.

A country without meaningful training capacity has limited ability to shape the underlying assumptions embedded in the models it depends on. A country unable to host inference at scale cannot consistently apply its own rules to the decisions those models produce. Authority shifts toward those who control the underlying capacity — not by design, but by material necessity.

The Incorporation Heuristic provides pathways for participation within this asymmetry. But incorporation does not dissolve it. The resource intensity of frontier AI remains a binding

¹⁰⁰ Frontier-model energy and cost estimates evolve rapidly. Figures cited here reflect conditions as of early 2026. See IEA, *Electricity 2026: Analysis and Forecast to 2030* (February 2026); Epoch AI, *Data on Machine Learning Hardware* (March 2026); Epoch AI, *Trends in Artificial Intelligence* (accessed April 2026).

¹⁰¹ Leading-edge semiconductor fabrication remains concentrated in Taiwan, particularly at the most advanced process nodes. SIA, *State of the U.S. Semiconductor Industry* (most recent edition).

¹⁰² Epoch AI, *Data on Machine Learning Hardware* (March 2026); Epoch AI, *Trends in Artificial Intelligence* (accessed April 2026); CSET, *Global AI Talent Tracker* (latest available edition).

constraint on national agency — one that governance frameworks must accommodate rather than assume away.

AI Physics — The New Seat of Power

Cloud providers, AI labs, semiconductor ecosystems, and energy networks provide not just services but foundations that shape rulemaking outcomes. Their design decisions set the practical boundaries within which public authority operates, as constitutional provisions set the limits of legislative power.

When a hyperscaler decides which regions receive GPU allocation, it shapes, in practice, which countries can exercise AI autonomy. When a chip designer sets export thresholds, it draws compute boundaries as consequential as many treaties. These are not commercial decisions with incidental policy effects. They are governance decisions expressed as engineering choices — accumulating into the material constitution of the AI era.

From Capacity to Legitimacy — The Convergence of All Prior Gaps

The legitimacy crisis of the AI era is not merely a new category alongside the Visibility Gap, the Rule-Execution Gap, the Categorical Gap, and the Measurement Gap identified in Parts I and II. It rests on their material foundation.

Each of those gaps — the inability to see, to act, to classify, or to measure AI systems — is ultimately a capacity gap. States that cannot support machine intelligence and cognition at scale cannot readily close any of them. The crisis is therefore not only ideological. It is the point where prior failures converge on a single physical constraint.

The Three Constitutional Questions of Part V

Part V is organised around three constitutional questions.

First, what are the new assets of power? Compute, data, and models are becoming the physical foundations on which authority is exercised — and their distribution helps determine which countries possess meaningful AI autonomy and which do not.

Second, what are the hard constraints? Five material boundaries — energy, semiconductors, capital, talent, and time — limit what any governance system can achieve, regardless of institutional design.

Third, what comes next? If control is bounded by physics and mediated largely by private firms, what institutional forms can build or restore democratic accountability for systems that currently operate beyond democratic reach?

These questions define the constitutional terrain of the AI era.

The Stakes — The Reality Every Country Must Confront

AI is more than a policy challenge. It is a general-purpose cognitive infrastructure with material resource requirements that cannot be wished away. Its diffusion across law, finance, medicine, logistics, defence, education, and identity systems will continue even when institutional readiness lags.

The asymmetry is material and time-sensitive. Countries that internalise capacity constraints will design within them. Those that treat AI as a purely regulatory challenge will encounter those constraints only in retrospect, as their rules collide with physical limits they failed to anticipate. And countries that cannot power AI will encounter machine decisions they cannot shape, supervise, or constrain.

The window for institutional design remains open — but narrowing under the same material constraints this book has identified. Governance without capacity is commentary. Capacity without governance is power unchecked. The challenge of the AI era is to hold both in tension — and to build institutions adequate to constraints no previous generation of political designers has confronted at this scale.

I. The New Assets: Compute, Data, Models

Frontier machine intelligence is defined by three properties that no previous industrial system has combined. It shows extreme capital indivisibility — you cannot train half a frontier model. It exhibits radical supply-chain concentration — leading-edge chip fabrication capacity remains confined to a handful of global facilities. And its progress is cumulative and path-dependent — each completed training run generates advantages that cannot easily be unwound.

From these properties emerge the three foundational assets of the AI era: compute as the engine, data as the fuel, and models as the territory.

A. Compute — The Engine

Compute is the capacity to transform energy into probabilistic inference at scale. It is not merely a resource. It is a national capacity.

The mechanism is straightforward but carries profound implications. A hyperscaler, following a regulatory dispute, reprioritises GPU allocation away from a region. The decision is commercial in origin; its consequences are structural. Countries lose access to training

clusters. Research labs suspend experiments. Startups miss product deadlines. Public agencies dependent on cloud inference see throughput collapse. What appears as operational reprioritisation is, in practice, a reallocation of national capacity.

Compute shapes the full lifecycle of AI. Training capacity influences the foundational assumptions of models; inference capacity determines who can operate them; alignment capacity determines who can control their behaviour. A country without compute faces real limits on its ability to control the AI systems operating within its borders — not merely to regulate them on paper, but to exercise operational authority over the decisions machine intelligence produces.

B. Data — The Fuel

Data is generated — by people, by institutions, by societies, and by civilisation itself. It is path-dependent, nationally situated, and politically charged. And the fundamental truth is unavoidable: many countries cannot independently generate the volume and quality of data required to train or align domestic models.

The asymmetry runs deeper than population size. Industrial economies generate substantially more high-value data than developing economies — not only in volume but in structural composition, because industrial applications, medical archives, legal documentation, and financial flows carry far more signal than consumer clicks or social-media traces.

And data carries norm encoding: it shapes a model's behavioural priors — what it learns to treat as normal, acceptable, risky, or harmful. This is why data is not the new oil. Oil is fungible. Data is non-fungible capital — the memory of a civilisation and the foundation of its future cognition.

Data asymmetry therefore produces cognitive disenfranchisement: the systematic exclusion of a society's norms, harms, and institutional patterns from the decision outcomes of AI models that increasingly mediate its economic and institutional life. Unlike political disenfranchisement, which can be remedied by extending the franchise, cognitive disenfranchisement requires the prior existence of the data that makes a society legible to machine learning. Societies absent from training corpora are not underrepresented — they are unweighted. Their harms are invisible. Their citizens become statistical outliers to the systems that increasingly determine access to credit, healthcare, mobility, and public services.

Data is not simply fuel. It is representation — the pathway by which a civilisation appears, or fails to appear, in the intelligence that its models generate and enact.

C. Models — The Territory

Frontier models embed decision boundaries shaped by mathematical constraints that determine, for any given input, what outputs are possible, probable, or prohibited. Whoever shapes those constraints shapes the universe of decisions available to every downstream user. Models are not merely products. They are territory — cognitive landmass.

Deployment of a foreign-trained model entails downstream inheritance. The foundational training decisions — data selection, alignment regimes, safety constraints — are made once by an upstream actor and inherited by every downstream user. Countries that deploy foreign-trained models accept machine behaviour shaped by upstream priors, alignment choices, and design constraints unless they possess the capacity to meaningfully influence or adapt those defaults. This is more than technical dependency. It is the adoption of behavioural boundaries encoded elsewhere.

The policy consequence is direct. At training time, models absorb rules — the behavioural constraints embedded through data selection, alignment, and safety design. At inference time, they enforce them — defining the acceptable boundaries for institutions, the decision logic of markets, and the informational landscape through which societies navigate. Once deployed, models become the operative carriers of these constraints within their deployment domains. Every downstream user operates within boundaries set upstream.

D. The Triad in Practice — A Single Scenario

Consider Nigeria — a country attempting to build a national model trained on African languages, legal corpora, and institutional data. The ambition is real; the constraints are equally real.

On the compute side, a competitive training run at frontier scale demands GPU capacity on the order of thousands of accelerators — a scale that exceeds the training-cluster capacity currently available within the country. Procuring such hardware requires navigating supply chains that run through Taiwan, Malaysia, South Korea, and U.S. hyperscalers — a dependency structure that most African states cannot meaningfully reroute independently.

On the data side, Nigerian legal texts, Hausa and Yoruba corpora, and financial records exist, but not in the digitised, structured, and cleaned form that frontier-scale training requires. The raw material is present; the industrial refinement pipeline is not.

On the model side, even if Nigeria were to secure compute and refine its data, the resulting system would still inherit architectural priors from upstream frameworks, toolchains, and alignment methods developed elsewhere. Without the capacity to adapt or retrain these components at scale, the country would deploy a model whose behavioural boundaries — what it treats as plausible, safe, risky, or prohibited — are shaped materially by external design choices. The cognitive territory would be national in aspiration but foreign in origin.

E. Transition — From Assets to Constraints

Compute, data, and models form the material foundation of AI. But recognising their importance does not answer the harder question: how many countries can assemble these assets at scale — and what happens to legitimacy when access to machine intelligence is globally uneven.

These assets are not evenly distributed, and they are not independently obtainable. Each rests on supply chains, energy systems, capital flows, talent pools, and time horizons that most countries do not control. National autonomy meets physics; ambition meets industrial reality.

The next section examines these dependencies — and the hard constraints they impose on states navigating AI industrialisation under conditions they did not choose.

II. The Legitimacy Crisis — The Widening Gap Between Statute and AI Governance

Industrial-era legitimacy rested on four stabilisers: consent, representation, procedure, and rights. These held because public authority was executed by human agents, institutions controlled the means of enforcement, and rules could be applied through established processes.¹⁰³

The AI era introduces a fifth — and materially prior — stabiliser: operational capacity. This is the ability of states to convert resources into machine-executed decision power at scale, and to maintain meaningful control over how that power is applied across institutions and society.

Without sufficient capacity, the classical stabilisers remain formally intact but lose practical force. AI systems increasingly execute decisions that were previously made by human agents within institutional frameworks — and states that cannot govern how those decisions are made, applied, and corrected find traditional instruments of accountability operating at a growing distance from the systems they are meant to constrain.

The pace of AI advancement is outstripping the adaptation cycles of most existing institutions, widening the gap between authority written in statute and authority realised through the ability to direct and control AI systems in practice.

¹⁰³ J. Rawls, *A Theory of Justice* (Harvard University Press, 1971); J. Habermas, *Between Facts and Norms* (MIT Press, 1996); L. Fuller, *The Morality of Law* (Yale University Press, 1964).

A. The Shift of Authority — From Institutional Control to AI-Driven Execution

Public authority that once flowed through institutions designed for democratic accountability is increasingly exercised through AI-mediated decisions that shape rulemaking and rule keeping in real time. The actors who control frontier models and agents gain influence over which rules can be enforced, how they are interpreted, and what forms of oversight remain possible at scale — not through elections or public mandate, but through the operational fact that AI now mediates how rules are executed.

Alignment and safety teams inside a handful of companies make design choices that determine how AI systems behave across vast numbers of interactions — which outputs are encouraged, restricted, or blocked. In practice, these choices function as large-scale norm-setting at a reach and pace that few legislatures or courts can match. The processes through which these decisions are made remain only partially subject to parliamentary debate, judicial review, or public accountability.¹⁰⁴

Operational authority is migrating into systems built for speed, scale, and computational advantage. A state whose instruments cannot extend to monitoring and guiding machine-executed decisions does not lose formal authority — but it loses the capacity to govern how power is exercised within its borders. The State continues to hold formal power on paper; AI systems increasingly shape daily life in practice.

B. The Five Non-Substitutable Constraints

The AI era imposes five non-substitutable constraints that determine whether a country can govern AI on its own terms: energy, semiconductors, capital, talent, and time. Each operates independently but locks together with the others — and they do not substitute for one another. Money cannot compensate for missing chip access; skilled engineers cannot overcome a lack of power; and time lost to earlier capability cycles cannot be recovered once others have moved ahead. A shortfall in any single constraint limits a state's ability to build, train, or control advanced AI systems, regardless of its strengths elsewhere.

Crucially, these constraints do not determine whether a country can use AI. Most nations can — and should — adopt AI through imported models, cloud services, or API access. What the constraints determine is whether a state can exercise control at the machine level, where decisions are actually executed. A country facing binding limits will rely on external AI systems — and in doing so cedes influence over how rules are interpreted, enforced, and adjusted in real time.

¹⁰⁴ R. Gorwa, R. Binns & C. Katzenbach, 'Algorithmic Governance: A Modes of Governance Approach' (2019) *Regulation & Governance*.

The choice is not between AI and no AI. It is between different modes of participation — each with its own trade-offs, dependencies, and opportunities for agency. The material constraints of the AI era define the boundaries of that choice: when key constraints are missing, countries are pushed toward more dependent forms of operation, while only those with all five at sufficient scale can exercise machine-level control. These conditions do not eliminate agency, but they shape the strategic space within which it can be exercised.

C. The Gradients of Operational Agency

Scarcity in the core inputs of AI does more than limit what countries can build. It produces the same permission gradient identified in Part IV — but grounded in material capacity rather than regulatory propagation. Those who build and run frontier models shape system behaviour at its foundations. Those who deploy or regulate these systems operate within boundaries set by the builders. Most governments and institutions apply systems whose deeper logic is defined elsewhere.

A parallel gradient runs through the rules inside these systems. Public rules — laws, regulations, and court decisions — move at institutional pace. Private rules — contracts, terms of service, audit requirements — evolve iteratively, revised by the same actors who control the underlying models and compute. The rules that most directly shape AI behaviour are increasingly written and updated outside the public sphere.

Reconnecting democratic authority to these systems requires more than access to models. It requires access to the legal and operational mechanisms that determine how those models behave in practice — and that access is currently concentrated among a small number of actors whose authority derives from capacity, not mandate.

Countries without these inputs are not excluded from AI — but they navigate a landscape where influence is unevenly distributed and the cost of delayed action compounds. The urgency lies not only in adopting AI, but in securing the mechanisms that determine how it behaves.

D. Legitimacy Debt — The Silent Failure Mode

The permission hierarchy rarely announces itself. Democratic institutions continue to function — legislatures legislate, regulators regulate, courts adjudicate. Yet the gap between what institutions believe they govern and what they can actually influence widens with each new model generation, each expansion in agent capability, each increase in compute deployment.

Legitimacy Debt accumulates when formal responsibility and real decision authority drift apart. Like technical debt, it is not recorded when it forms — it remains hidden until the system is under strain. Institutions are held accountable for outcomes they cannot reliably

shape at the machine-execution layer. The gap between attributed responsibility and operational control grows silently.¹⁰⁵

This produces a Misleading Dashboard Effect. Monitoring tools continue to signal compliance and control even as actual system behaviour moves beyond institutional reach — a dynamic foreshadowed by the cockpit analogy in Part I, where surface indicators remain calm while underlying misalignment deepens. The dashboard stays green; the execution layer moves on. When failure arrives, it is often sudden and nonlinear — not because institutions are inattentive, but because the operational mechanisms have already shifted beyond their grasp.

E. What Legitimacy Collapse Looks Like

Legitimacy collapse in the AI era is unlikely to resemble revolution. It is more likely to take the form of institutional irrelevance. A ministry publishes an AI ethics framework. It circulates through international forums. It appears in the country's Universal Periodic Review. Yet few global firms operationalise it, and no dominant compliance template incorporates it. Deployed systems continue unchanged.

The framework enters the documentary archive but never reaches the machine. The ministry continues to regulate. The minister continues to speak. Yet citizens interact daily with AI models and agents their government cannot shape, audit, or align.

Institutional formalities persist while decision authority migrates. The throne remains. The occupant has vacated.

F. The New Test of Legitimacy — Consent versus Execution

In the industrial era, legitimacy centred on consent: do the people agree? In the AI era, legitimacy increasingly also depends on execution: can the State execute? Both principles remain essential. The crisis emerges not because one displaces the other, but because system stability now depends on variables the two optimise differently.

Democratic legitimacy prioritises consent — deliberative, procedural, inclusive. Operational legitimacy prioritises execution — fast, reliable, technically coherent. No widely adopted institutional framework has yet reconciled both at system scale.

¹⁰⁵ 'Legitimacy Debt' is an original analytic term introduced in this manuscript, modelled by analogy to technical debt. W. Cunningham, 'The WyCash Portfolio Management System' (1992), OOPSLA Experience Report.

Institutions optimised for consent struggle under inference-speed decision cycles. Institutions optimised for execution concentrate authority beyond classical consent mechanisms. This is the defining constitutional tension of the AI era.

G. The Five Design Conditions for AI Legitimacy

Any governance framework capable of sustaining legitimacy in the AI era must satisfy five design conditions simultaneously. It must preserve output, since legitimacy cannot survive persistent performance degradation. It must restore input, ensuring that democratic mandates reach executable decisions rather than remaining declarative policy. It must bridge the velocity gap, allowing institutional processes to interoperate with machine-speed execution. It must enable accountability, keeping decision authority observable, attributable, and contestable at system speed. And it must enable exit, allowing individuals and institutions to leave a regulatory arrangement without prohibitive costs, informational lock-in, or dependence on a single model's reasoning system.

As with the Incorporation Heuristic, these conditions operate multiplicatively rather than additively. A framework that preserves output but cannot restore input accumulates Legitimacy Debt regardless of performance. An arrangement that enables accountability but cannot bridge the velocity gap induces regulatory delay incompatible with machine-speed systems. A design that enables exit but cannot preserve output fragments the very system it seeks to govern.

In multiplicative systems, a single zero collapses the product. Legitimacy is no exception. All five conditions must be satisfied simultaneously — a standard that existing institutions were not designed to meet. The next section — Code After — sets out the architecture proposed to meet it.

III. Code After: The Operating Manual for the AI Era

Part V concludes with the architecture itself — not a vision or a manifesto, but a constitutional design proposed for the AI era, grounded in material constraints and derived from the analytical framework developed across the preceding Parts.

The manuscript has traced a continuous arc. It began with the State's inability to see or execute the AI economy it governs, moved through the dual governance languages that define economic reality and the Translation Layer that now mediates between them, and mapped the vertical stack and the G3 fracture it produces. It then followed the Incorporation Heuristic and the propagation mechanisms that determine which rules reach the systems that actually execute, before arriving at the material constraints and legitimacy conditions that bound what any institutional response can achieve. What remains is the institutional design those constraints imply.

The Third Architecture of Governance introduces three mechanisms — the Sovereign API, Parametric Democracy, and the Protocol of Federation — that together form a separation-of-powers arrangement for machine-executed AI ecosystems.

The Sovereign API provides an institutional interface linking democratic authority to machine-executed systems. Parametric Democracy calibrates democratic input into executable constraints. The Protocol of Federation enables incompatible normative foundations to coexist without requiring internal value harmonisation.¹⁰⁶

Together, these mechanisms satisfy the five multiplicative design conditions — preserving output, restoring input, bridging the velocity gap, enabling accountability, and enabling exit — not through trade-offs but through functional separation across distinct levels of authority. The design does not prescribe political preference. It derives institutional necessity.

A. The Sovereign API

The Sovereign API is an institutional interface specification. It defines the channels through which democratically legitimate authority shapes machine-executed outcomes. On the input side, it receives statutes, regulations, thresholds, and jurisdiction-specific obligations. On the output side, it renders them as constraint schemas, parameter ranges, prohibited and required behaviours, and documentation requirements that can be embedded directly into models, workflows, and systems. The Sovereign API does not execute rules; it translates them into a form that technical systems can consume. It does not replace the Incorporation Heuristic; it reshapes the inputs that determine its results.

A statutory non-discrimination norm — "no automated decision may consider protected attributes except where expressly permitted by law" — illustrates the mechanism. The API would not adjudicate fairness or resolve the underlying substantive inquiry; instead, it would enforce an interface requirement that any model used for an eligibility decision declare whether protected attributes or recognised proxies are present in its inputs, and block execution unless a lawful basis or statutory exception is supplied.

In this sense, the Sovereign API is not a technical instrument but a constitutional one: a structured pathway through which legal authority is expressed as operational constraint, ensuring that the rules governing AI systems originate in public law rather than private construct.

This does not reduce anti-discrimination law to feature disclosure alone, but it does translate one component of the legal rule into an ex ante constraint that the deployment environment must satisfy. In this way, a normative legal value becomes a binding precondition for

¹⁰⁶ The Third Architecture is an original institutional design proposed in this manuscript. "Third Architecture" distinguishes the proposed framework from the Westphalian territorial model (first) and governance through private intermediaries (second, the Invisible Constitution).

technical execution, without requiring the State to modify the underlying model. This logic generalises.

The adjustment the Sovereign API imposes is visibility: deployed AI systems must make their behaviour legible — observable, verifiable, and auditable through institutionally authorised channels. Under this specification, visibility becomes a constitutional requirement rather than an operational preference.

The first adjustment establishes visibility as institutional design. By requiring legibility as a condition of execution, the Sovereign API ensures that downstream normative constraints — whether grounded in non-discrimination, due process, or sector-specific mandates — operate on systems whose behaviour can be inspected and justified. Visibility is not an ancillary safeguard; it is the precondition that allows democratically legitimate authority to govern machine-executed outcomes without intervening in the model's internal mechanics.

The second adjustment establishes workability. The Sovereign API translates democratic input into constraints that AI systems can execute — rendering legislative intent in machine-interpretable form. By reducing normative ambiguity, it increases the workability of public rules across the national AI ecosystem.

The third adjustment establishes necessity. Through federated interface standards — developed further in the Protocol of Federation — the Sovereign API allows national rules to become embedded in the economic network itself. When multiple sovereigns adopt compatible interface specifications, compliance ceases to be a bilateral regulatory demand and becomes a condition of network participation.

Firms that align with common interfaces maintain uninterrupted access; those that diverge encounter increasing friction through the economics of network position rather than diplomatic sanction. Under these conditions, necessity shifts from market size to network membership.

The Sovereign API therefore operates at the interface — the constitutional hinge linking democratic consent to machine execution. Rules that once failed to influence machine behaviour gain operational weight when they are institutionally observable, expressed as executable parameters, and reinforced through network-level necessity.

B. Parametric Democracy — Governing at Machine Speed

Citizens cannot directly govern AI systems operating at neural-network scale. Frontier models operate at a complexity, speed, and scale that text-based legislation cannot specify in detail. Parametric Democracy preserves democratic legitimacy by shifting regulation from rule authorship to constraint calibration. Citizens do not vote on executable logic; they vote on the trade-offs that executable systems must respect.

The Thermostat Model illustrates the principle. A citizen does not need to understand thermodynamic engineering to set a temperature, just as they do not need to understand neural networks to express calibrated preferences across familiar tensions — privacy and convenience, innovation and safety, efficiency and employment, autonomy and oversight. The analogy is not about simplification; it is about constructing the right abstraction boundary. The interface translates value preferences into parameter constraints, while the model manages the execution dynamics beneath them.

This raises the central institutional question: who controls the interface? The legitimacy of Parametric Democracy depends on which parameters are exposed, how they are framed, and what decision ranges are made available. If the design of those parameters is controlled by the actors the process is meant to constrain, the arrangement reproduces the legitimacy deficit it was created to resolve. The safeguard is institutional pluralism in interface design — competing parameterisations, transparent selection procedures, and independent oversight of the constraint-definition framework.

Parametric governance is not speculative. Its functional components already operate in isolation. Quadratic voting demonstrates how preference aggregation can be expressed as calibrated intensities rather than binary choices.¹⁰⁷

Reinforcement learning from human feedback and Constitutional AI show how value constraints can be encoded into model behaviour.¹⁰⁸

Central banks illustrate how mandates can be translated into parameterised policy instruments. Platforms such as vTaiwan demonstrate how digital deliberation can structure public input on contested questions.¹⁰⁹

Parametric Democracy integrates these mechanisms into a coherent design capable of sustaining democratic oversight at machine speed.

C. The Protocol of Federation — Interoperability Without Harmonisation

The Third Architecture does not require cultural convergence. It requires protocol-level compliance. The Protocol of Federation enables interaction between incompatible normative systems without demanding alignment of their internal values.

Its foundational principle is straightforward: national rules remain authoritative within their own borders, while cross-system interoperability depends on adherence to shared protocol specifications. When a country departs from those specifications, it encounters network-level

¹⁰⁷ E.A. Posner and E.G. Weyl, *Radical Markets* (Princeton University Press, 2018), ch. 2.

¹⁰⁸ P.F. Christiano et al., 'Deep Reinforcement Learning from Human Preferences' (2017) NeurIPS 30; Y. Bai et al., 'Constitutional AI' (2022), arXiv:2212.08073.

¹⁰⁹ C. Hsiao et al., 'vTaiwan: An Empirical Study of Open Consultation Process in Taiwan' (2018); vtaiwan.tw.

consequences rather than diplomatic sanction. Isolation emerges as the structural outcome of non-compliance, not as a political punishment.

Isolation does not prevent the existence of rogue or state-sponsored shadow networks; it prevents their integration. The Protocol of Federation is not a universal security perimeter but a constitutional boundary: it ensures that systems operating outside the compliance regime cannot draw on, contaminate, or impersonate actors within it. By tying access to shared identity services, auditability, and cross-operator interoperability to compliance, the protocol makes the benefits of federation unavailable to non-aligned networks.

Rogue ecosystems may persist, but they remain informationally and operationally segregated, unable to exploit the trust, identity, or compute pathways that define the primary network. In this sense, isolation is not a tool for suppressing shadow networks; it is a mechanism for preventing their expansion into the institutional core.

Hong Kong offers a transitional illustration. Its dual-system configuration — a common-law framework operating within Chinese national authority — demonstrates how divergent legal regimes can coexist within a shared economic environment. The arrangement also illustrates a structural truth about federated governance: it is contingent on the stability of the underlying political compact, not a permanent constitutional settlement.

The Protocol of Federation marks a shift in regulatory logic — from harmonisation to compatibility. Nations need not agree on the meaning of liberty, privacy, order, or rights; they need only maintain protocol-level interoperability. The framework achieves coexistence not by resolving value disagreements but by rendering them operationally manageable at the machine interface.

The Protocol federates constraint schemas, compliance parameters, audit outputs, and risk classifications across jurisdictions, enabling systems to operate across borders without collapsing into the lowest common denominator. Its function is not to erase legal differences but to structure them — to allow divergent regimes to express their authority without forcing technical systems into mutually exclusive configurations.

Where jurisdictional constraints conflict, the Protocol resolves them through three mechanisms. Precedence rules determine which country's requirements govern in cases of overlap, subject to the applicable legal test. Arbitration mechanisms provide defined channels for negotiated reconciliation where precedence alone cannot resolve the tension. Fallback defaults ensure that systems remain operable even when substantive disagreements persist, reducing the risk of system-level failure. The output is a set of interoperable rule bundles that allow AI systems to operate under multiple legal regimes simultaneously.

In this sense, the Protocol of Federation is neither a treaty nor a standard. It is an institutional mechanism for managing legal pluralism at the machine interface — a way for nations to remain sovereign while participating in a shared computational environment.

D. The Four-Tier Constitutional Design and the Role of Licensed Operators

The Third Architecture operates through a separation of powers not between branches but between tiers. Each performs a distinct constitutional function, and none can substitute for the others.

The Value Tier is where societies form preferences, express norms, and generate legitimacy. The Protocol Tier translates those preferences into machine-readable constraints. The Operations Tier executes those constraints within the limits of physics — compute, energy, bandwidth, and latency. The Audit Tier measures and verifies AI behaviour, ensuring that systems remain observable and contestable over time. Together, these four tiers constitute the minimum constitutional structure capable of stabilising governance in the AI era.¹¹⁰

The Operations Tier is constrained not only by physics but by limits on authority. AI systems at frontier scale — with their tight interconnections, cross-domain effects, and autonomous capabilities — require operation by Licensed Operators: entities authorised to run advanced AI under domestic rules and within protocol-compatible licensing regimes. Licensing is not centralisation; it is functional delegation.

A Licensed Operator functions as a legally constrained environment: an entity required to comply with the Sovereign API, maintain real-time and retrospective audit and interruption mechanisms, and operate under transparent public oversight consistent with alignment and safety obligations. Licensing thresholds are tied to capability levels rather than fixed technical specifications, allowing the framework to evolve as systems advance.

This structure enables frontier-scale AI to be deployed without compromising the integrity of the domestic constitutional order. It permits multiple entities to operate advanced systems within the same country while maintaining uniform conditions for accountability, contestability, and national control.

A further consideration is economic feasibility. If interface compliance becomes the core of authority, then the cost of maintaining compliant interfaces must not become a de facto barrier to entry. Frontier-scale systems require secure logging, identity verification, auditability, and interruption mechanisms, all of which impose fixed costs that can be prohibitive for smaller firms.

Ensuring a contestable ecosystem therefore requires complementary economic supports: open-compute subsidies that lower the baseline cost of compliant operation, shared public infrastructure for identity and audit services, and standardised compliance modules that reduce duplication across operators. These measures do not prescribe a particular market

¹¹⁰ The four-tier constitutional design — Value, Protocol, Operations, Audit — is an original framework introduced in this manuscript, drawing on traditions in constitutional separation of powers, systems architecture, and audit theory.

structure; they ensure that the operational overhead of compliance does not, by itself, determine who is allowed to participate.

Preventing capture requires more than technical neutrality; it requires ensuring that no single actor can monopolise the compliance stack. Antitrust scrutiny is necessary to prevent dominant firms from controlling the interfaces through which all operators must pass, and to ensure that compliance tooling does not become vertically integrated into a single platform.

Open-compute subsidies and shared infrastructure serve a structural purpose: they prevent compliance costs from concentrating power in incumbents and preserve the possibility of new entrants operating frontier-level systems under the same constitutional constraints. The Licensed Operator model is not only a governance design but an economic design — one that distributes the cost of safety in a way that keeps the space contestable.

This does not eliminate the material advantages of scale — capital, compute, and expertise remain significant barriers — but it ensures that those advantages operate within a framework of public constraint rather than private incumbency. The purpose of the licence is not to entrench hyperscalers but to condition continued participation on adherence to public rules that are equally available to all. Firms that fail to meet these constraints lose access not through discretionary licensing, but through the predictable consequences of non-alignment within a federated network. The Licensed Operator model preserves contestability while anchoring frontier-scale AI within a constitutional order that remains legible, accountable, and open to new entrants.

The Audit Tier measures and verifies key aspects of AI system behaviour. It evaluates model outputs, data flows, control effectiveness, and compliance states, producing attestations, exception reports, and behavioural traces. Unlike financial audit, which is periodic, document-based, and ledger-bound, this function is continuous, instrumentation-based, and system-level. It does not interpret rules; it generates an evidentiary record against which compliance and non-compliance can be determined.

Measurement warrants particular emphasis. As the manuscript has argued throughout, it is not auxiliary — it is constitutive. Without persistent measurement, continuous adaptation produces opacity rather than oversight. In the terms of the Incorporation Heuristic, the Audit Tier sustains the visibility variable — and without sustained visibility, no rule incorporates regardless of its workability or necessity. Accounting disciplines, built to track change over time, supply continuity that legal doctrine alone cannot maintain at machine speed.

This tier exists because the Measurement Gap identified in Part II cannot be resolved within existing accounting frameworks. Level 3 fair value — designed as a last-resort valuation method when market evidence is weak — is structurally mismatched to frontier AI systems whose behaviour is probabilistic, path-dependent, and continuously updated. Instead of estimating price, the Audit Tier generates continuous visibility into system operations, provenance, and effects through secure logging, identity verification, traceability, and interruption mechanisms. It tracks behaviour rather than financial position — a functional alternative to Level 3 valuation when valuation ceases to be informative.

Because existing accounting standards remain anchored to fair-value frameworks that cannot express model provenance, drift, or operational behaviour, the Audit Tier operates as a parallel evidentiary system. Its inputs are evaluation traces — safety-case documentation, training-data attestations, capability disclosures — rather than valuation estimates. The Audit Tier does not repair the Measurement Gap; it bypasses it, creating a verifiable visibility regime for systems whose economic and constitutional relevance cannot be rendered through conventional accounting categories.

E. Calibration Events — Democratic Revision at Machine Speed

Continuous systems require periodic legitimacy inputs. Calibration Events are proposed as the adaptive mechanisms through which citizens adjust the parameters that guide AI behaviour without interrupting ongoing operation. They provide structured, recurring signals of democratic judgement to systems that learn and adapt at machine speed.

Calibration Events occur at predictable but adjustable intervals, creating a stable rhythm for citizen input. They rely on deliberative platforms that support informed public reasoning rather than reactive preference expression. Their outputs take the form of parameter updates rather than statutes. Collective judgement is translated into machine-readable adjustments that enter the Sovereign API as regular revisions. Revised preferences then flow through the Protocol Tier and refresh the normative inputs that guide the system — without disrupting continuity in the Operations Tier.

Democratic governance adapted to machine-speed systems requires this regular update. Calibration Events allow societies to revise goals, adjust risk tolerances, and update normative baselines while keeping AI systems running — turning legitimacy from a one-time mandate into a recurring control signal. They complete the constitutional cycle: from Value to Protocol to Operations to Audit — and back to Value.

F. The Hard Problems — Inherent Tensions That Remain

Four underlying tensions define the research frontier of the Third Architecture.

The first is enforcement: how the Sovereign API can bind actors whose computational capacity, operational reach, or technical autonomy may exceed that of any single state.

The second is coordination: how multiple Sovereign APIs interoperate across national boundaries without drifting into incompatible standards or collapsing toward the lowest common denominator. This tension is especially acute for non-G3 countries, which must navigate competing AI ecosystems while preserving their own regulatory voice.

The third is transition: how the world moves from today's Invisible Constitution to the Third Architecture without destabilising the digital and economic foundations on which societies

depend. For countries reliant on foreign cloud providers or imported AI capabilities, the transition path is as consequential as the destination.

The fourth is scale: elements of parametric democracy have been demonstrated in small, digitally cohesive societies such as Estonia and Taiwan. It remains an open question whether these approaches can operate in much larger countries where trust, identity, and digital infrastructure are unevenly distributed.¹¹¹

What works for a population of five or twenty-three million — high-bandwidth participation, stable digital identity, and broadly shared civic baselines — may not translate to India, Nigeria, Indonesia, or the U.S., where the parameters themselves are contested and where administrative capacity varies by orders of magnitude.

Scaling a deliberative system requires not only more participants but shared agreement on the common civic baselines that make deliberation possible: identity, authenticity, eligibility, and the legitimacy of the parameters being tuned. In nations where these baselines are fractured — because identity systems are contested, digital infrastructure is uneven, or trust is thin — parametric democracy may face structural limits that are not visible in smaller, more cohesive societies.

These tensions are not deficiencies. They define the frontier of constitutional engineering — the unresolved problems that will shape the research agenda and institutional debates of the AI era. They mark where further work is needed, and where the next generation of scholars, engineers, and policymakers must focus to ensure that the proposed design serves all societies, not only those with the most resources.

G. The Interdisciplinary Necessity

No single discipline contains the concepts or tools required to design constitutional frameworks for AI systems and the ecosystems of autonomous agents they are producing. Each tier of the Third Architecture draws on a different intellectual tradition because each addresses a distinct category of problem.

The Value Tier, expressed through the Sovereign API, depends on legal definition and interpretation — the ability to express obligations, permissions, and constraints in forms that institutions and machines can both process. The Protocol Tier, realised through the Protocol of Federation, draws on geopolitical reasoning — an understanding of how states, blocs, and cultural systems negotiate standards, assert authority, and manage interdependence. The Operations Tier rests on computational engineering — the knowledge required to understand what frontier models and agentic systems can and cannot do within physical limits. The Audit

¹¹¹ R. Kattel and I. Mergel, 'Estonia's Digital Transformation: Mission Mystique and the Hiding Hand', in *Great Policy Successes* (Oxford University Press, 2019).

Tier relies on accounting and measurement — disciplines built to verify claims, detect inconsistencies, and create records that can be trusted and contested.

This interdisciplinary foundation — law as form, geopolitics as coordination, engineering as constraint, accounting as evidence — reflects the structural realities of the AI era. Legal rules become machine-readable. Geopolitical dynamics shape technical standards. System architecture determines which forms of governance are possible. Economic measurement becomes continuous. No single field addresses these interactions alone. The Third Architecture requires their integration because the systems it governs operate across all four domains simultaneously.

H. Closing the Loop — The Visibility Gap Revisited

This work began with two linked impulses: a parent's question about the professional world two daughters — one studying law, one studying accounting — would enter, and the recognition that no single State can fully see the AI universe it now governs. The first impulse led to the second. The second became this manuscript.

The Visibility Gap, the Rule-Execution Gap, the Categorical Gap, and the Measurement Gap are not failures of political will. They are structural features of an AI-driven universe that has outpaced the observational and operational capacity of the institutions built to govern it. Closing these gaps requires not better regulations but better design — institutions capable of translation, mechanisms matched to the speed of computation, and democratic processes built for complexity.

The Translation Layer already exists. The Incorporation Heuristic already compiles. The Material Constitution already operates. The physics will advance regardless.

What remains absent is the democratic framework that legitimises these mechanisms and connects them to the consent of the governed. That framework is not inevitable. It must be designed, built, and defended — by the institutions that claim authority and the societies that grant it.

The architecture remains a choice. The choice remains ours.

Code After.

GLOSSARY OF TERMS

Terms are listed alphabetically. Cross-references appear in *italics*. Part of first significant appearance noted in parentheses.

Agentic Shift The transition from AI systems that generate predictions to systems that initiate actions — reallocating decision rights and driving governance downward into the stack. (Part III)

AI Era The period in which AI systems shift from passive tools to autonomous actors capable of executing decisions, initiating transactions, and modifying their environments — forcing governance to operate at machine speed rather than institutional time. (Part I)

AI Governance Stack A four-layer architecture through which sovereign authority is compiled into operational reality: law (permission), accounting (measurement), professional services (translation), and compute (execution). In the AI era, these layers are tightly coupled. (Part III)

Blended Template A unified operational compliance document synthesised from multiple, often conflicting, regulatory regimes — the mechanism through which the *Translation Layer* converts legal plurality into executable order. (Part I, Ch. 2)

Build-Time Irreversibility The condition in which compliance documentation and design constraints are most reliably generated during system construction, and whose evidentiary assurance diminishes once training is complete. A flow-on effect of European AI regulation. (Part IV)

Calibration Event A structured, periodic democratic process proposed in this manuscript through which citizens adjust the parameters that guide AI behaviour — producing parameter updates rather than statutes. Feeds into the *Sovereign API* through the *Value Tier*. (Part V)

Capacity Primacy The governance logic of the AI era: control of computational capacity conditions which rules can be made effective and who can set the technical defaults that become embedded throughout the stack. Introduced as a complement to classical rule primacy, not a replacement. (Part V)

Categorical Gap The structural mismatch between legal categories designed for deterministic, territorial, human-driven activity and AI systems that are none of these things — producing a thinning of meaning in which legal classifications cannot attach cleanly to technical realities. (Part II)

Cognitive Disenfranchisement The systematic exclusion of a society's norms, harms, and institutional patterns from the decision outcomes of AI models — arising from the absence of that society's data in training corpora. (Part V)

Compliance Cascade A chain of secondary obligations that firms must satisfy together to maintain market access — impact assessments, disclosure procedures, consent management, localisation requirements, and audit documentation that cluster into bundles and become embedded in product design. Generated by *Gateway Rules*. (Part IV)

Compliance Gridlock The condition in which incompatible commands from multiple sovereigns must all be obeyed but none can be reconciled by the states that issued them. (Part I, Ch. 2)

Data Port An emerging institutional form in which cross-sovereign computation occurs within a jurisdiction capable of hosting it physically, legally, and infrastructurally. (Part IV)

Dual-Compiler Thesis The argument that law and accounting function as the two constitutive languages through which AI becomes governable — law as the compiler of permission, accounting as the compiler of value. Neither alone is sufficient; both are required. (Part II)

Fairness Trap The condition in which firms attempting to operationalise fairness across incompatible legal regimes produce not harmonisation but balkanisation — because mathematically rigorous fairness definitions are mutually exclusive when base rates differ across groups. (Part I, Ch. 2)

Five Design Conditions The simultaneous requirements for legitimate AI governance: preserve output, restore input, bridge the velocity gap, enable accountability, and enable exit. These conditions operate multiplicatively. (Part V)

Five Functions of AI Legal Language The core classificatory operations through which law governs AI: permission, liability, ownership, disclosure, and sovereignty. Together they form the *Permission Framework*. (Part II)

Four-Tier Constitutional Design The minimum constitutional structure proposed in this manuscript for stabilising AI governance, comprising the *Value Tier* (preference formation and legitimacy), the *Protocol Tier* (translation into machine-readable constraints), the *Operations Tier* (execution within physical limits), and the *Audit Tier* (measurement and verification). An original framework drawing on traditions in constitutional separation of powers, systems architecture, and audit theory. (Part V)

G2 An analytical category referring to state-level AI ecosystems that meet the threshold of vertically integrated capability across hardware, compute, model development, talent concentration, and capital scaling. The United States and China are the clearest G2 states. The term is descriptive rather than geopolitical. (Part I)

G3 States whose regulatory frameworks propagate globally through market size, legal-harmonisation capacity, and institutional legitimacy, even without full integration across the AI stack. The EU is the clearest G3 actor alongside the G2 states. (Part I)

G3 Fracture The condition in which the United States, China, and the European Union no longer interpret AI through a shared conceptual frame — each operating a distinct definition of economic reality encoded in a different accounting, legal, and institutional worldview. (Part III)

Gateway Rule A regulatory requirement whose compliance forces changes to foundational system design — data schemas, reporting formats, documentation pipelines, network configurations, or contractual defaults — causing effects that propagate across borders via system design rather than formal national adoption. Extends the *Incorporation Heuristic* by isolating the mechanism of cross-border regulatory reach. (Part IV)

Golden Image A preconfigured compliance baseline — master service agreements, audit protocols, data-protection addenda, cross-border transfer mechanisms — replicated across deployments by professional services firms to ensure consistency and reduce configuration risk. The term is borrowed from systems administration and applied to global compliance template propagation as an original formulation in this manuscript. (Part IV)

Governance Vacuum A structural space in which the State retains formal authority but loses the practical capability to exercise it — filled by familiar intermediaries whose proximity to AI systems gives them the operational vantage required to interpret, evaluate, and constrain behaviour in practice. (Part I)

Incorporation Heuristic ($I = V \times W \times N$) An analytical model explaining why some sovereign rules propagate into global practice and others do not, as a function of three variables: Visibility (detectability and credible enforcement), Workability (operational implementability), and Necessity (market gravity making exit irrational). The multiplicative structure means failure at any variable is dispositive. (Part IV)

Incorporation Priority Hierarchy The selection logic through which *Blended Templates* determine which rules become the global baseline: market access (Tier 1), enforcement credibility (Tier 2), litigation exposure (Tier 3), and reputational stakes (Tier 4). (Part I, Ch. 2)

Invisible Charter The evaluative framework produced by the combined practice of lawyers and accountants through which AI systems become economically real — the precondition for economic existence that precedes regulation, valuation, and every market transaction involving an intelligent system. Distinguished from the *Invisible Constitution*. (Part II)

Invisible Constitution The operational rulebook that governs global AI regardless of what formal law provides — produced by compliance engineers, cloud platforms, and the propagation logic of convergent design rather than by legislatures alone. (Part IV)

Legitimacy Crisis The widening gap between authority written in statute and authority realised through the ability to direct and control AI systems in practice — arising when the pace of AI advancement outstrips the adaptation cycles of existing institutions. (Part V)

Legitimacy Debt The structural stress that accumulates when formal responsibility and real decision authority drift apart — hidden until the system is under strain. An original analytic term in this manuscript, modelled by analogy to technical debt. (Part V)

Licensed Operator An entity authorised to run frontier-scale AI systems under constitutional constraints — required to comply with the *Sovereign API*, maintain audit and interruption mechanisms, and operate under transparent public oversight. Licensing thresholds are tied to capability levels rather than fixed technical specifications. (Part V)

Measurement Gap The rupture that forms when industrial-era accounting frameworks confront probabilistic, self-updating capital — limiting the State's capacity to price AI's contribution even when it can regulate AI's conduct. (Part II)

Misleading Dashboard Effect The condition in which monitoring tools continue to signal compliance and control even as actual system behaviour moves beyond institutional reach — surface indicators remain calm while underlying misalignment deepens. (Part V)

Obsolescence Trap The dynamic in which a rule designed for one model generation arrives just in time to regulate the one that no longer exists — by the time a rule is drafted, the practice has evolved; by the time it is enforced, the system design has changed; by the time it is revised, the field has moved again. (Part I, Ch. 2)

Parametric Democracy A governance mechanism proposed in this manuscript that preserves democratic legitimacy by shifting regulation from rule authorship to constraint calibration — citizens vote on the trade-offs that executable systems must respect rather than on the executable logic itself. (Part V)

Permission Framework The grammar through which the State decides what may operate, on what terms, and at whose risk — composed of the *Five Functions of AI Legal Language*. (Part II)

Procedural Truth A valuation or measurement output that is methodologically justified, auditor-defensible, and regulator-acceptable — binding not because it is metaphysically accurate but because it satisfies the evidentiary and methodological requirements of the standards. The mechanism through which probabilistic systems acquire economic identity. (Part II)

Protocol of Federation An institutional mechanism proposed in this manuscript for managing legal pluralism at the machine interface — enabling interaction between incompatible normative systems without demanding alignment of their internal values. Nations need not agree on the meaning of liberty, privacy, order, or rights; they need only maintain protocol-level interoperability. (Part V)

Ratchet Effect The dynamic in which the most demanding rule in a *Blended Template* becomes the baseline applied worldwide — because maintaining a single global standard is cheaper than engineering regional variants. (Part I, Ch. 2)

Root Access / Admin Access / User Access The three tiers of the *Topology of Power*. Root-access states write the defaults that propagate globally. Admin-access actors shape how those defaults are implemented within specific domains. User-access countries retain full sovereignty but lack the incorporation weight to project their rules beyond their borders. (Part IV)

Rule-Execution Gap The structural separation between the declaration of a rule and its realisation inside the systems meant to execute it — the distance a rule must travel from legislative intent to technical implementation. (Part I)

Sovereign API An institutional interface specification defining the channels through which democratically legitimate authority shapes machine-executed outcomes. On the input side, it receives statutes, regulations, and obligations; on the output side, it renders them as constraint schemas, parameter ranges, prohibited and required behaviours, and documentation requirements that technical systems can consume. Not a technical instrument but a constitutional one. (Part V)

Sovereignty Paradox The condition in which the political map recognises nearly two hundred sovereign states while the AI economy organises itself around a far smaller number of regulatory gravity centres whose rules propagate unevenly through global systems. (Part IV)

Stack Dependence The condition in which a state regulates models built on chips it did not fabricate, trained on compute it does not own, and deployed on platforms whose core designs are external — its authority resting primarily at the regulatory and market interface rather than in shaping the technical base. (Part I)

Third Architecture The constitutional design for the AI era proposed in this manuscript, comprising the *Sovereign API*, *Parametric Democracy*, and the *Protocol of Federation*. Distinguished from the Westphalian territorial model (first architecture) and the current default governance through private intermediaries described as the *Invisible Constitution* (second architecture). (Part V)

Topology of Power The permission hierarchy derived from the *Incorporation Heuristic*: a gradient of operational influence determining how far a state's rules travel beyond its borders — resolving into *Root Access*, *Admin Access*, and *User Access* tiers. (Part IV)

Translation Imperative The foundational necessity of an intermediate layer capable of converting legal text into machine-readable constraints — arising because norms do not execute themselves and legislatures cannot perform this conversion at the required level of technical specificity. (Part II)

Translation Layer The distributed ecosystem of legal, accounting, technical, and assurance institutions — principally Big Law and the Big Four accounting networks — that operates between statutes and code, converting governance principles into executable protocols. As applied to AI governance, this is an original concept of this manuscript. (Part I, Part II, Part IV)

Tripartite Failure The combined effect of the *Visibility Gap*, *Rule-Execution Gap*, and *Categorical Gap* — defining the structural limits of governing AI through legal categories built for a different class of objects. The *Measurement Gap*, identified subsequently, extends the diagnosis to accounting frameworks. (Part II)

Visibility Gap The structural asymmetry in which regulators cannot observe the systems they claim to govern — composed of three dimensions: technical opacity, jurisdictional fragmentation, and temporal acceleration. (Part I)

Visibility Monopoly The structural condition in which the observational capacity required to govern AI concentrates among actors whose position cannot be replicated or bypassed by public authority — requiring three conditions: privileged access to system-level observability, non-substitutability, and functional dependence. (Part I)
